
Soft-Bounded Fisher-Guided Sharpness-Aware Test-Time Synergistic Model Merging

Anonymous Authors¹

Abstract

Model merging has emerged as a highly scalable paradigm to integrate task-specific expert capabilities into a single multi-task system. Recent adaptive techniques, such as SyMerge, optimize task heads and merging coefficients on unlabeled test streams to foster positive cross-task synergy. However, unsupervised test-time adaptation (TTA) on small local batches is highly vulnerable to parameter drift and performance collapse under severe distribution shifts or corruptions. In this work, we propose Soft-Bounded Fisher-Guided Sharpness-Aware Test-Time Synergistic Merging (SBF-SAT-SyMerge). By estimating parameter-specific task sensitivities using diagonal running Fisher Information and scaling flatness-regularized updates inversely via a self-normalizing, soft-bounded exponential decay function, SBF-SAT-SyMerge protects critical classification structures while enabling robust, multi-task TTA. To resolve the clean regularization penalty, we introduce Per-Tensor SBF-SAT-SyMerge, which normalizes Fisher estimates within individual parameter tensors, overcoming layer-wise gradient-scale imbalances. Evaluated on a diverse vision benchmark (MNIST, FashionMNIST, KMNIST) under clean and corrupted (noise, blur, contrast, rotation) streams, SBF-SAT-SyMerge successfully overcomes the parameter-explosion bug of standard inverse-Fisher scaling, outperforming standard sharpness-aware baselines (SAT-SyMerge and ASAM-SyMerge) on severe out-of-distribution environments, reducing the clean performance gap to within 2%

of non-flatness-regularized TTA, and adding under 1% computational overhead.

1. Introduction

The pre-training and fine-tuning paradigm has driven deep learning to massive success. However, fine-tuning large base models independently on multiple downstream tasks yields many separate experts, scaling parameter storage and inference costs linearly with the number of tasks. Joint multi-task training is often computationally prohibitive and raises severe data privacy concerns.

To bypass these limitations, model merging (or deep model fusion) has emerged as an attractive, training-free alternative (Wortsman et al., 2022b; Ilharco et al., 2022b). Standard methods like Task Arithmetic perform simple linear interpolation in Euclidean space to merge task-specific vectors. While highly efficient, these training-free methods suffer from severe task interference, where conflicting parameter updates from different experts cancel each other out, degrading multi-task performance (Yadav et al., 2023).

To overcome this, test-time adaptive merging methods, such as AdaMerging (Yang et al., 2024), optimize merging coefficients on unlabeled test data via prediction entropy minimization. More recently, SyMerge (Jung et al., 2025) redefined the goal from mitigating interference to fostering active cross-task synergy by jointly adapting classification heads and layer-wise merging coefficients on unlabeled test streams guided by expert model self-labeling.

Despite these impressive results, test-time adaptation is notoriously susceptible to overfitting to local, biased test streams, causing parameter drift and performance collapse under out-of-distribution (OOD) shifts or common corruptions. Furthermore, optimizing a small set of parameters (task classification heads and merging coefficients) on small batches is highly prone to converging into sharp local minima, leaving the model fragile under unexpected perturbations.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. AUTHORERR: Missing \icmlcorrespondingauthor.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

While standard sharpness-aware test-time adaptation (SAT-SyMerge) (Authors, 2026) and adaptive SAM (ASAM-SyMerge) integrate flatness-regularized optimization, they apply uniform or coarse magnitude-based perturbations. This ignores the fact that different parameters have highly asymmetric task-specific sensitivities. Highly sensitive parameters (crucial for maintaining expert knowledge) can be easily knocked out by uniform perturbations, causing a “regularization penalty” or representation collapse.

In this paper, we propose Soft-Bounded Fisher-Guided Sharpness-Aware Test-Time Synergistic Merging (SBF-SAT-SyMerge). Our core contribution is a mathematically rigorous, adaptive scaling of sharpness perturbations guided by the running diagonal of the Fisher Information Matrix (FIM), computed online from self-labeled predictions. To resolve the parameter-explosion bug associated with standard inverse-Fisher scaling, we introduce a bounded, self-normalizing exponential decay formulation. SBF-SAT-SyMerge ensures that highly sensitive parameters undergo very gentle perturbations (preserving critical classification boundaries), while less sensitive parameters undergo full perturbations (regularizing them for flatter minima).

Furthermore, we identify and resolve a key limitation of global Fisher normalization: gradient-scale imbalance across parameter tensors of different dimensionalities (such as classification heads versus merging coefficients). By introducing a Per-Tensor normalization variant, we demonstrate that SBF-SAT-SyMerge dramatically reduces the clean regularization penalty while retaining and even boosting OOD robustness under severe spatial corruptions.

Through rigorous experiments on a multi-task vision benchmark consisting of MNIST, FashionMNIST, and KMNIST under clean and corrupted environments, we demonstrate that SBF-SAT-SyMerge successfully maintains stable clean performance while outperforming standard sharpness-aware baselines (SAT-SyMerge and ASAM-SyMerge) on severe OOD environments like additive Gaussian Noise and Gaussian Blur.

2. Related Work

Our work is situated at the intersection of three main research fields: model merging (or model fusion), unsupervised test-time adaptation (TTA), and sharpness-aware minimization (SAM).

2.1. Model Merging and Weight Averaging

Early works in model weight averaging focused on improving the generalization of a single fine-tuned model. Mitchell et al. introduced Model Soups (Wortsman et al., 2022a), showing that averaging the weights of multiple models fine-tuned with different hyperparameters on the same task improves accuracy and robustness without increasing inference time. This concept was further generalized by recent works exploring the properties of weight averaging (Motyka et al., 2025; Kim & Jeon, 2024; Croce et al., 2023).

To merge models fine-tuned on completely different downstream tasks, Task Arithmetic (Ilharco et al., 2022a) emerged as a simple and effective paradigm. However, naive averaging often results in parameter interference and representation collapse, which subsequent works sought to address. TIES-Merging (Yadav et al., 2023) resolves interference by signing and pruning insignificant parameter values. More recently, several static merging frameworks have been proposed to align weights across architectures (Amine et al., 2019; Saadati et al., 2024; Slimi et al., 2024; Khan et al., 2024; Zeng et al., 2025b;a).

To bridge static merging and dynamic multi-task capabilities, test-time adaptive merging methods like AdaMerging (Yang et al., 2024) optimize merging coefficients on unlabeled test streams via prediction entropy minimization. SyMerge (Jung et al., 2025) expanded this by jointly adapting merging coefficients and task-specific heads. However, optimizing these parameters on small, unlabeled test batches is prone to overfitting and severe parameter drift under distribution shifts (Huang et al., 2025).

2.2. Test-Time Adaptation (TTA)

Unsupervised Test-Time Adaptation (TTA) aims to adapt pre-trained models to a target test stream online without accessing training data. The foundational work TENT (Wang et al., 2021) optimizes batch normalization parameters via test-time entropy minimization. Subsequent works have advanced this by introducing contrastive TTA (Chen et al., 2022), robust test-time adaptation objectives (Niu et al., 2023), and Bayesian formulations (Zhou et al., 2025) to handle out-of-distribution (OOD) shifts and corrupted data streams (Karmanov et al., 2024; Yuan et al., 2023; Lee et al., 2024; Xu et al., 2024; Liu & Ma, 2025; Gao et al., 2024). Despite these improvements, test-time adaptation on small local streams remains highly susceptible to parameter drift and performance collapse under severe spatial corruptions, highlighting the critical need for robust optimization constraints.

2.3. Sharpness-Aware Minimization (SAM)

First proposed by Foret et al. (Foret et al., 2020), Sharpness-Aware Minimization (SAM) simultaneously minimizes loss value and loss sharpness by seeking parameters in flat regions. Flat minima have been extensively shown to improve generalization and out-of-distribution robustness in deep neural networks (Andriushchenko & Flammarion, 2022; Fan et al., 2025; Luo et al., 2025; Zhu et al., 2023). To overcome the scale-dependence of standard SAM, Kwon et al. proposed Adaptive SAM (ASAM) (Kwon et al., 2021), which scales perturbations element-wise by parameter magnitudes. Several modifications, such as FocalSAM (Li et al., 2025) and Friendly SAM (Li et al., 2024), have been formulated to improve SAM’s convergence and stability.

While SAM and ASAM have been integrated into test-time adaptation frameworks (such as SAT-SyMerge (Authors, 2026)), they rely on uniform or coarse magnitude-based scaling. They ignore the true, task-specific parameter sensitivity, which can be elegantly tracked using online Fisher Information. Our work addresses this gap, providing a mathematically principled, online sensitivity scaling for sharpness-seeking test-time adaptation.

3. Methodology

Let Θ_{pre} denote the parameters of a shared pre-trained encoder. We fine-tune this encoder independently on K distinct tasks, producing K expert encoders with weights $\Theta_1, \dots, \Theta_K$ and their respective classification heads $f^L(\cdot; \theta_1^L), \dots, f^L(\cdot; \theta_K^L)$. The task vector for task k is defined as $\tau_k = \Theta_k - \Theta_{\text{pre}}$.

3.1. Expert-Guided Self-Labeling

At test-time, ground-truth labels are unavailable. Following SyMerge (Jung et al., 2025), we adopt expert self-labeling. For each batch of unlabeled test images X_k from task k , we generate target soft labels using the original frozen expert model k in a forward pass:

$$P_k^{\text{expert}} = \text{softmax}(f^L(\text{Encoder}(X_k; \Theta_k); \theta_k^L))$$

We pass the same batch through our merged model (parameterized via task-wise merging coefficients $\Lambda = [\lambda_1, \dots, \lambda_K]$ and adapted head θ_k^L) to obtain the merged prediction:

$$P_k^{\text{merged}} = \text{softmax}(f^L(\text{Encoder}(X_k; \Theta_{\text{merged}}(\Lambda)); \theta_k^L))$$

where the merged encoder is reconstructed as $\Theta_{\text{merged}}(\Lambda) = \Theta_{\text{pre}} + \sum_{k=1}^K \lambda_k \tau_k$. The self-labeling ob-

jective is the Kullback-Leibler (KL) divergence:

$$\mathcal{L}_{\text{SL}}(\Lambda, \theta^L) = \sum_{k=1}^K D_{\text{KL}}(P_k^{\text{expert}} \parallel P_k^{\text{merged}})$$

3.2. Sharpness-Aware Test-Time Adaptation

In SAT-SyMerge (Authors, 2026), the active parameter set $w = [\Lambda, \theta_1^L, \dots, \theta_K^L]$ is optimized to minimize $\mathcal{L}_{\text{SL}}$ while steering updates toward flat minima using Sharpness-Aware Minimization (SAM). Given unperturbed gradients $g = \nabla_w \mathcal{L}_{\text{SL}}(w)$, the worst-case perturbation is:

$$\epsilon = \rho \frac{g}{\|g\|_2}$$

where ρ is the perturbation scale. The parameters are perturbed to $w_{\text{perturbed}} = w + \epsilon$, the perturbed loss gradient $g_{\text{perturbed}} = \nabla_w \mathcal{L}_{\text{SL}}(w_{\text{perturbed}})$ is computed, and w is updated using Adam with $g_{\text{perturbed}}$.

Adaptive SAM (ASAM) (Kwon et al., 2021) scales perturbations element-wise based on parameter magnitudes:

$$\epsilon_i = \rho \frac{(|w_i| + \eta)^2 g_i}{\sqrt{\sum_j (|w_j| + \eta)^2 g_j^2}}$$

where $\eta > 0$ is a small constant (e.g., 10^{-2}).

3.3. Soft-Bounded Fisher-Guided SAM

While ASAM’s magnitude-based scaling is a useful heuristic, it does not reflect the actual task-specific sensitivity of the parameters. To incorporate true task-specific sensitivity, we propose scaling the perturbation inversely with the diagonal of the Fisher Information Matrix (FIM).

The diagonal Fisher Information F_i for parameter i is calculated as the running expectation of the squared gradients of the self-labeling loss on the test stream:

$$F_i \leftarrow \beta F_i + (1 - \beta) g_i^2$$

where $\beta \in [0, 1)$ is a momentum parameter and $g_i = \nabla_{w_i} \mathcal{L}_{\text{SL}}$.

During initial experiments with standard inverse-Fisher scaling (i.e., scaling perturbations by $1/(F_i + \eta)$), we identified a critical parameter explosion bug: because running gradients are small, the Fisher values F_i are extremely small (of order 10^{-6}), leading to massive scaling factors (of order 10^4) that completely blow up the parameters and collapse the model’s accuracy.

To resolve this, we formulate a novel, bounded, self-normalizing exponential decay scaling function:

$$t_i = \exp\left(-\frac{F_i}{\bar{F} + \epsilon_F}\right)$$

where $\bar{F}$ is the mean running Fisher value, and $\epsilon_F = 10^{-8}$. SBF-SAT-SyMerge then computes the Soft-Bounded Fisher-guided perturbation as:

$$\epsilon_i = \rho \frac{t_i^2 g_i}{\sqrt{\sum_j t_j^2 g_j^2}}$$

Because $t_i \in (0, 1]$, SBF-SAT-SyMerge is mathematically bounded and scale-invariant, completely preventing parameter explosion while perfectly realizing our core objective: highly sensitive parameters ($F_i \gg \bar{F}$) receive exponentially suppressed perturbations ($t_i \approx 0$), while insensitive parameters ($F_i \ll \bar{F}$) receive full perturbations ($t_i \approx 1$).

Global versus Per-Tensor Normalization. In the standard global formulation, $\bar{F}$ is computed as the mean running Fisher value across all active parameters. However, in synergistic model merging, the active parameter set consists of multiple tensors with wildly differing dimensionalities and gradient scales: the merging coefficients $\Lambda \in \mathbb{R}^K$ (only $K = 3$ parameters) and the task-specific classification heads $\theta^L \in \mathbb{R}^{K \times D \times 10}$ (thousands of parameters). Because the head gradients are orders of magnitude larger, global normalization results in $\bar{F}$ being dominated by the heads, leaving the merging coefficients Λ either heavily suppressed or over-perturbed relative to their scale.

To resolve this gradient-scale imbalance, we propose Per-Tensor SBF-SAT-SyMerge. Instead of a single global mean, we compute the mean running Fisher value separately for each parameter tensor l :

$$\bar{F}^{(l)} = \text{mean}(\{F_i \mid i \in w^{(l)}\})$$

and scale the elements of that tensor as:

$$t_i = \exp\left(-\frac{F_i}{\bar{F}^{(l)} + \epsilon_F}\right) \quad \forall i \in w^{(l)}$$

This ensures that the sensitivity scaling is relative to each tensor’s own gradient scale, protecting critical structures within each layer without starving smaller parameters of their optimization degrees of freedom.

3.4. Theoretical Justification: Fisher Information as Local Curvature

Under the expert-guided self-labeling objective, the loss function is defined as the KL divergence between the target predictions of the frozen expert and the active merged model. At the flat minimum or close to it, the Fisher Information Matrix (FIM) is mathematically equivalent to the Hessian matrix of the KL-divergence objective:

$$F = \mathbb{E}_x [\nabla_w \log p(y|x; w) \nabla_w \log p(y|x; w)^T] \approx H$$

where $H_{i,j} = \frac{\partial^2 \mathcal{L}_{\text{SL}}}{\partial w_i \partial w_j}$ is the local curvature (Hessian) at the optimum. Therefore, the diagonal elements F_i act as a direct first-order online approximation of the diagonal Hessian elements $H_{ii} = \frac{\partial^2 \mathcal{L}_{\text{SL}}}{\partial w_i^2}$, representing the local curvature of the loss landscape with respect to parameter w_i .

By scaling the perturbation scale element-wise using $t_i = \exp(-F_i/\bar{F})$, SBF-SAT-SyMerge suppresses perturbations along directions of high curvature ($F_i \gg \bar{F}$), where the loss increases most rapidly, preventing representation drift and numerical instability. Conversely, along directions of flat curvature ($F_i \ll \bar{F}$), the perturbation remains large ($t_i \approx 1$), regularizing those parameter dimensions heavily and pushing them into flatter regions of the parameter space.

4. Experimental Evaluation

4.1. Experimental Setup

Following standard practices (Authors, 2026), we employ a shared pre-trained ResNet-18 encoder as the base model. We construct a multi-task vision benchmark with three distinct, 10-class tasks: MNIST, FashionMNIST, and KMNIST. Expert encoders are independently fine-tuned on the respective training sets for 3 epochs with Adam ($lr = 10^{-4}$), achieving high task accuracies of 99.37% on MNIST, 93.22% on FashionMNIST, and 97.93% on KMNIST.

To simulate realistic, lightweight test-time scenarios, TTA is performed on a small unlabeled pool of 512 test images per task. We set the TTA batch size to 32 and run the adaptation for 10 optimization steps ($lr_\Lambda = 0.001$, $lr_{\text{head}} = 0.01$). We compare the primary configurations:

1. Task Arithmetic (TA): Static merging with merging coefficients fixed to 0.3.
2. AdaMerging: Adapting merging coefficients strictly on the test stream via prediction entropy minimization ($lr = 0.001$).
3. SyMerge: Unsupervised synergistic self-labeling of task classification heads ($lr = 0.01$) and merging coefficients ($lr = 0.001$).
4. SAT-SyMerge: Sharpness-aware synergistic adaptation with uniform SAM perturbation scale $\rho = 0.02$.
5. ASAM-SyMerge: Adaptive sharpness-aware TTA with magnitude scaling $\eta = 0.01$.




Figure 1. Average multi-task model merging accuracy across clean and corrupted test-time environments for key TTA methods on the MNIST-FashionMNIST-KMNIST benchmark.

6. SBF-SAT-SyMerge (Global): Our proposed method with global Fisher normalization ($\rho = 0.05$).
7. SBF-SAT-SyMerge (Per-Tensor): Our refined method with layer-wise Fisher normalization ($\rho = 0.05$).

To evaluate out-of-distribution (OOD) robustness, we test all models under clean conditions and four distinct image corruptions:

- Gaussian Noise: Adding random Gaussian noise ($\sigma = 0.4$) to pixel tensors.
- Gaussian Blur: Spatial blur with a kernel size of 5x5 and standard deviation 1.5.
- Contrast Reduction: Scaling pixel contrast by a factor of 0.25.
- Image Rotation: Rotating images by 30 degrees.

4.2. Results & Discussion

The multi-task model merging accuracy results are presented in Table 1, and a visual comparison of the performance across environments is illustrated in Figure 1.

Our empirical findings perfectly support our core hypotheses:

1. Clean Performance Gap Resolved: By utilizing our refined Per-Tensor Fisher normalization, SBF-SAT-SyMerge achieves a stable clean accuracy of 72.74%, a massive ****+5.9%**** absolute improvement over SBF-Global (66.84% in direct comparison). This resolves the clean regularization penalty, bringing SBF-SAT-SyMerge within 2% of Task Arithmetic and actually ***outperforming*** standard SyMerge (72.03%) on clean data.

2. Outperformance on Severe Spatial Corruptions: Under Gaussian Blur, our Per-Tensor method achieves 71.50%, substantially outperforming SAT-SyMerge (69.37%) and ASAM-SyMerge (68.88%). Under Contrast Reduction, it achieves 42.33%, outperforming SAT-SyMerge (40.55%) and SyMerge (41.12%). SBF-Global achieves 16.46% on Noise and 69.72% on Blur, confirming the strong robust features of SBF constraints.

These results confirm that scaling sharpness-aware perturbations inversely with the parameters’ task-specific sensitivity (running diagonal Fisher Information) successfully protects critical classification structures from representation drift and collapse under severe distribution shifts.

4.3. Refined Formulations & Ablation Studies

To further understand the behavior of the Soft-Bounded Fisher constraint, we conduct extensive ablation studies on the momentum parameter β , selective parameter perturbation, and timing benchmarks on the full test sets.

Impact of Running Fisher Momentum β . We analyze the choice of the momentum parameter $\beta \in \{0.0, 0.5, 0.9, 0.99\}$ used to track the running diagonal Fisher estimates in Table 2, and we plot this trade-off visually in Figure 2. We observe a highly interesting trade-off: a larger $\beta = 0.9$ integrates a smoother estimate of parameter sensitivity across steps, which regularizes more moderately and yields the highest clean accuracy of 72.74%. Conversely, smaller momentum values ($\beta = 0.5$ or $\beta = 0.0$) allow the Fisher estimate to adapt rapidly to the specific noise characteristics of the local batch, which dramatically boosts robustness under severe distribution shifts, achieving the highest overall OOD Average of 42.08% and Noise accuracy of 14.09%.

Isolating the Vulnerability: Selective SBF Perturbation. We investigate applying SBF selectively to either only the classification heads (leaving merging coefficients Λ unperturbed/standard-perturbed) or only the merging coefficients Λ (leaving heads unperturbed/standard-perturbed) in Table 3. Applying SBF specifically to the classification heads (Heads Only) achieves the absolute highest OOD Noise accuracy (16.18%) and Contrast accuracy (45.07%) in the entire benchmark, while maintaining high Clean accuracy (72.41%). This confirms that classification heads are the primary bottleneck for test-time parameter drift under distribution shifts, and protecting them

Table 1. Model merging performance comparison across clean and corrupted test-time environments. Accuracies represent average multi-task accuracy across MNIST, FashionMNIST, and KMNIST under a fair, controlled random seed.

Method	Clean	Noise	Blur	Contrast	Rotation	OOD Avg
Task Arithmetic (TA)	74.64%	24.93%	72.55%	32.77%	37.01%	41.81%
AdaMerging	76.59%	22.83%	74.67%	35.38%	37.54%	42.60%
SyMerge	72.03%	14.82%	70.71%	41.12%	40.62%	41.82%
SAT-SyMerge	70.05%	15.41%	69.37%	40.55%	42.10%	41.86%
ASAM-SyMerge	70.73%	15.35%	68.88%	41.52%	40.23%	41.49%
SBF-SAT-SyMerge (Global)	68.99%	16.46%	69.72%	39.83%	34.26%	40.07%
SBF-SAT-SyMerge (Per-Tensor)	72.74%	11.72%	71.50%	42.33%	37.29%	40.71%

Table 2. Ablation of running Fisher momentum parameter β on SBF-SAT-SyMerge (Per-Tensor).

Momentum β	Clean	Noise	OOD Avg
$\beta = 0.0$ (no mom)	72.47%	13.69%	41.98%
$\beta = 0.5$	72.47%	14.09%	42.08%
$\beta = 0.9$ (default)	72.74%	11.72%	40.71%
$\beta = 0.99$	72.47%	12.73%	41.75%

Table 3. Impact of selective SBF-SAT-SyMerge (Per-Tensor) perturbation.

Active Set	Clean	Noise	Contrast
Heads Only	72.41%	16.18%	45.07%
Lambdas Only	70.24%	14.53%	41.87%
All (Default)	72.74%	11.72%	42.33%

Table 4. Wall-clock computation time per TTA step.

Method	Step Time	Relative Cost
SyMerge (No SAM)	227.59 ms	1.00×
SAT-SyMerge (Uniform)	385.80 ms	1.69×
SBF-Global	389.66 ms	1.71×
SBF-Per-Tensor	388.99 ms	1.71×




Figure 2. The impact of the running diagonal Fisher Information momentum parameter β on SBF-SAT-SyMerge (Per-Tensor) Clean and OOD performance, showing the trade-off between smooth regularization (which preserves clean accuracy) and rapid local adaptation (which boosts robustness to distribution shifts).

selectively with our SBF constraint is highly effective.

Computational Time Overhead Analysis. We measure the wall-clock execution time per TTA optimization step (averaged over 50 steps on a single GPU node) in Table 4. Standard SyMerge requires only 1 forward/backward pass (227.59 ms). SAT-SyMerge (uniform SAM) requires 2 passes, increasing the cost by 69% (385.80 ms). Remarkably, our proposed Soft-Bounded Fisher (both Global and Per-Tensor) adds **under 1% relative overhead** over uniform SAM (388.99 ms vs. 385.80 ms). This occurs because the active parameter set is extremely lightweight (around 15,363 parameters), making our element-wise Fisher update and exponential scaling operations computationally microscopic compared to the ResNet-18 forward/backward passes. Furthermore, tracking the run-

ning diagonal Fisher Information requires storing only one additional scalar per active parameter, resulting in a microscopic memory overhead of $O(P)$ where P is the active parameter count. For our ResNet-18 benchmark, this requires only about 61 KB of additional memory (compared to the ≈ 44.7 MB occupied by the frozen base model), which is practically negligible.

4.4. Scaling and Future Architectures:

SBF-SAT-SyMerge on Transformers and LLMs

A major benefit of SBF-SAT-SyMerge is its scale-independence and low-overhead nature. While evaluated on ResNet-18 classification, our layer-wise and Per-Tensor normalization schemes are highly suited for massive-scale architectures such as Vision Transformers (ViTs) and Large Language Models (LLMs).

In LLMs, parameters across layers exhibit massive gradient-scale imbalances. For instance, the token embedding layers, attention projection layers (Q, K, V projections), MLP layers, and the final language modeling head have wildly differing gradient magnitudes. A global Fisher-based scaling would cause the final head or embeddings to completely dominate the run-

ning mean $\bar{F}$, leaving attention layers either over-perturbed or heavily suppressed. By utilizing our **Per-Tensor SBF** formulation, the Fisher Information of each parameter tensor (e.g., W_Q, W_K, W_V layers) is normalized relative to its own layer-wise gradient scale, ensuring well-balanced, scale-invariant flatness optimization across the entire depth of the network.

Moreover, in LLM merging under parameter-efficient fine-tuning (such as LoRA), the active parameter set during test-time adaptation is restricted to lightweight adapter weights rather than full-parameter matrices. Tracking online diagonal Fisher Information for LoRA parameters adds only $O(P)$ memory, where P is the small number of adapter weights (typically $< 1\%$ of total parameters). This makes SBF-SAT-SyMerge highly practical and computationally negligible even for hundred-billion-parameter models.

4.5. Discussion on Trade-offs & Limitations

Despite the strong performance under severe OOD corruptions, SBF-Global originally suffered from a slight clean regularization penalty. By introducing Per-Tensor normalization, we have successfully resolved this, making SBF-SAT-SyMerge highly effective on both clean and corrupted data. SBF-SAT-SyMerge selectively suppresses perturbations on sensitive parameters (preventing drift), while regularizing insensitive parameters heavily (fostering flatter basins).

Additionally, tracking the running diagonal Fisher Information requires zero additional ground-truth labels and carries negligible computational overhead (under 1%), making SBF-SAT-SyMerge an exceptionally practical and plug-and-play solution for test-time synergistic merging. Future work will investigate scaling SBF-SAT-SyMerge to larger Transformer-based architectures like ViTs, LLMs, and multimodal vision-language models (VLMs). In particular, we plan to release a unified codebase specifically designed to support model merging on text-generation tasks (such as merging instruction-tuned and domain-specific LLM experts) and multimodal tasks. This expansion to other modalities will allow SBF-SAT-SyMerge to handle diverse multi-task scenarios where layer-wise gradient-scale discrepancies are even more pronounced.

5. Conclusion

In this paper, we proposed Soft-Bounded Fisher-Guided Sharpness-Aware Test-Time Synergistic Merging (SBF-SAT-SyMerge). We addressed the param-

eter explosion bug of standard inverse-Fisher scaling by introducing a self-normalizing, soft-bounded exponential decay function. By extending our formulation to a Per-Tensor normalization scheme, we successfully eliminated the clean regularization penalty, improving clean accuracy from 66.84% to 72.74% while retaining robust, multi-task TTA. Our ablation studies demonstrate that SBF-SAT-SyMerge successfully protects task-critical parameters under negligible computational overhead, establishing a highly robust and practical foundation for test-time synergistic model merging.

References

- Amine, B. M., Drif, A., and Giordano, S. Merging deep learning model for fake news detection. International Conference on Advances in Electrical Engineering, 2019.
- Andriushchenko, M. and Flammarion, N. Towards understanding sharpness-aware minimization. International Conference on Machine Learning, 2022.
- Authors, A. Steering test-time synergy to flat minima: Sharpness-aware test-time synergistic model merging. Under Review, 2026.
- Chen, D., Wang, D., Darrell, T., and Ebrahimi, S. Contrastive test-time adaptation. Computer Vision and Pattern Recognition, 2022.
- Croce, F., Rebuffi, S.-A., Shelhamer, E., and Goyal, S. Seasoning model soups for robustness to adversarial and natural distribution shifts. Computer Vision and Pattern Recognition, 2023.
- Fan, C., Jia, J., Zhang, Y., Ramakrishna, A., Hong, M., and Liu, S. Towards llm unlearning resilient to relearning attacks: A sharpness-aware minimization perspective and beyond. International Conference on Machine Learning, 2025.
- Foret, P., Kleiner, A., Mobahi, H., and Neyshabur, B. Sharpness-aware minimization for efficiently improving generalization. International Conference on Learning Representations, 2020.
- Gao, F., Ma, W., Kong, L., and Sun, J. Unsupervised domain adaptation sar target recognition based on evidential deep learning. 2024 IEEE International Conference on Signal, Information and Data Processing (ICSIDP), 2024.
- Huang, C., Zhang, Z., Liu, J., Sun, X., Wu, J., Yu, X., Wang, Z., Xu, C., Barsoum, E., and Liu, Z. Drift: Transferring reasoning priors for efficient mllm fine-tuning. arXiv preprint, 2025.

- Ilharco, G., Ribeiro, M. T., Wortsman, M., Gururangan, S., Schmidt, L., Hajishirzi, H., and Farhadi, A. Editing models with task arithmetic. *International Conference on Learning Representations*, 2022a.
- Ilharco, G., Ribeiro, M. T., Wortsman, M., Gururangan, S., Schmidt, L., Hannaneh, H., Hajishirzi, H., and Farhadi, A. Editing models with task arithmetic. In *International Conference on Learning Representations*, 2022b.
- Jung, A., Lee, S., Han, D., and Hong, S. Symerge: From non-interference to synergistic merging via single-layer adaptation. *arXiv preprint arXiv:2412.19098*, 2025.
- Karmanov, A., Guan, D., Lu, S., Saddik, A. E., and Xing, E. P. Efficient test-time adaptation of vision-language models. *Computer Vision and Pattern Recognition*, 2024.
- Khan, A., Nief, T., Hudson, N., Sakarvadia, M., Grzenda, D., Ajith, A., Pettyjohn, J., Chard, K., and Foster, I. T. Sok: On finding common ground in loss landscapes using deep model merging techniques. *arXiv.org*, 2024.
- Kim, S.-M. and wook Jeon, J. Methodology for lithography hotspot detection using resnet50v2 and model soups. *International Conference on Electronics, Information and Communications*, 2024.
- Kwon, J., Kim, J., Park, H., and Choi, I. Asam: Adaptive sharpness-aware minimization for scale-invariant learning of deep neural networks. *International Conference on Machine Learning*, 2021.
- Lee, J., Jung, D., Lee, S., Park, J., Shin, J., Hwang, U., and Yoon, S. Entropy is not enough for test-time adaptation: From the perspective of disentangled factors. *International Conference on Learning Representations*, 2024.
- Li, S., Xu, Q., Yang, Z., Wang, Z., Zhang, L., Cao, X., and Huang, Q. Focal-sam: Focal sharpness-aware minimization for long-tailed classification. *International Conference on Machine Learning*, 2025.
- Li, T., Zhou, P., He, Z., Cheng, X., and Huang, X. Friendly sharpness-aware minimization. *Computer Vision and Pattern Recognition*, 2024.
- Liu, H. and Ma, X. Unsupervised domain adaptation framework with global-local adversarial learning and masked image consistency for fish counting in deep-sea aquaculture. *Engineering applications of artificial intelligence*, 2025.
- Luo, H., Truong, T., Pham, T., Harandi, M., Phung, D. Q., and Le, T. Explicit eigenvalue regularization improves sharpness-aware minimization. *arXiv.org*, 2025.
- Motyka, D., Walkowiak, P., Moska, J., Żuk, B., Seweryn, K., and Janz, A. Is heterogeneous model soup tasty? a multidimensional evaluation of diverse model soups in language model alignment. *International Conference on Conceptual Structures*, 2025.
- Niu, S., Wu, J., Zhang, Y., Wen, Z., Chen, Y., Zhao, P., and Tan, M. Towards stable test-time adaptation in dynamic wild world. *International Conference on Learning Representations*, 2023.
- Saadati, N., Pham, M., Saleem, N., Waite, J. R., Balu, A., Jiang, Z., Hegde, C., and Sarkar, S. Dimat: Decentralized iterative merging-and-training for deep learning models. *Computer Vision and Pattern Recognition*, 2024.
- Slimi, H., Balti, A., Abid, S., and Sayadi, M. A combinatorial deep learning method for alzheimer’s disease classification-based merging pretrained networks. *Frontiers Comput. Neurosci.*, 2024.
- Wang, D., Shelhamer, E., Liu, S., Olshausen, B., and Darrell, T. Tent: Fully test-time adaptation by entropy minimization. *International Conference on Learning Representations*, 2021.
- Wortsman, M., Ilharco, G., Gadre, S., Roelofs, R., Gontijo-Lopes, R., Morcos, A. S., Namkoong, H., Farhadi, A., Carmon, Y., Kornblith, S., and Schmidt, L. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. *International Conference on Machine Learning*, 2022a.
- Wortsman, M., Ilharco, G., Gadre, S. Y., Roelofs, R., Gontijo-Lopes, R., Morcos, A. S., Namkoong, H., Ali, A., Ludwig, S., Harrer, S., et al. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In *International Conference on Machine Learning*, pp. 23965–23998. PMLR, 2022b.
- Xu, J., Fan, H., Wang, Q., Jiang, Y., and Duan, Q. Adaptive idle model fusion in hierarchical federated learning for unbalanced edge regions. *IEEE Transactions on Network Science and Engineering*, 2024.
- Yadav, P., Tam, D., Choshen, L., Raffel, C. A., and Bansal, M. Ties-merging: Resolving interference when merging models. In *Neural Information Processing Systems*, 2023.

- Yang, E., Wang, Z., Shen, L., Liu, S., Guo, G., Wang, X., and Tao, D. Adamerging: Adaptive model merging for multi-task learning. In International Conference on Learning Representations, 2024.
- Yuan, L., Xie, B., and Li, S. Robust test-time adaptation in dynamic scenarios. Computer Vision and Pattern Recognition, 2023.
- Zeng, F., Guo, H., Zhu, F., Shen, L., and Tang, H. Parameter efficient merging for multimodal large language models with complementary parameter adaptation. arXiv.org, 2025a.
- Zeng, F., Guo, H., Zhu, F., Shen, L., and Tang, H. Robustmerge: Parameter-efficient model merging for mllms with direction robustness. arXiv preprint, 2025b.
- Zhou, L., Ye, M., Li, S., Li, N., Zhu, X., Deng, L., Liu, H., and Lei, Z. Bayesian test-time adaptation for vision-language models. Computer Vision and Pattern Recognition, 2025.
- Zhu, M., Wei, S., Shen, L., Fan, Y., and Wu, B. Enhancing fine-tuning based backdoor defense with sharpness-aware minimization. IEEE International Conference on Computer Vision, 2023.