
Information-Geometric Gradient Surgery for Open-World Test-Time Model Merging

Anonymous Author(s)¹

Abstract

Test-Time Model Merging (TTMM) has emerged as a parameter-efficient paradigm for adapting pre-trained expert models to unlabeled test streams without backpropagating through backbone parameters. While recent closed-world TTMM methods utilize diagonal Fisher Information to scale layer-wise updates, they fail under realistic open-world deployment where novel task domains arrive dynamically. Conversely, state-of-the-art open-world TTMM frameworks utilize feature-space prototype cohesion to identify novel domains, but adapt global merging coefficients via standard Euclidean updates, ignoring layer-specific parameter sensitivities and suffering from severe representation collapse and "feedback loop traps". To bridge this fundamental gap, we propose Information-Geometric Gradient Surgery for Open-World Test-Time Model Merging (IGGS-OW). Our framework introduces two key innovations: first, we identify a critical representational space alignment mismatch in existing literature, and resolve it by precomputing prototypes in a Unified Static Space. Second, we model the test-time adaptation of layer-wise merging coefficients in a Riemannian space defined by the joint diagonal Fisher Information of the experts. During adaptation on novel streams, IGGS-OW scales learning rates according to the information-geometric sensitivity of each layer—dampening updates in highly sensitive layers (e.g., classifier heads or early convolutions) while accelerating stable adaptation in robust intermediate representation blocks. Our extensive empiri-

cal evaluation on multi-task vision streams (MNIST, KMNIST, FashionMNIST) shows that IGGS-OW achieves perfect novelty routing (0% False Positive Rate, 100% Novelty Detection Rate) and reaches 92.24% classification accuracy on the novel domain, outperforming competitive baselines by a massive margin (+22.92% absolute) and surpassing even the pre-trained expert's fine-tuned baseline.

1. Introduction

Modern deep learning has witnessed a dramatic shift towards pre-training large-scale foundational models and fine-tuning them on specialized downstream domains, yielding library collections of specialized expert models (Wortsman et al., 2022; Radford et al., 2021). To combine these disparate expert capabilities at inference time, model merging techniques such as task arithmetic (Ilharco et al., 2023) and Fisher weighted averaging (Matena & Raffel, 2022) have been developed. Extending model merging to the online streaming setting, Test-Time Model Merging (TTMM) adapts the merging coefficients on-the-fly to match the active test distribution (Yang et al., 2024; Hübötter et al., 2025). TTMM offers a highly parameter-efficient alternative to standard Test-Time Adaptation (TTA) (Wang et al., 2021; Zhang et al., 2022), avoiding full-parameter backpropagation and its associated memory footprints, while enjoying the representational stability of pre-trained weights.

Despite their success, existing TTMM frameworks are severely limited. On one hand, closed-world TTMM methods such as FP-CA (Luan et al., 2026) and IGGS-Merge (Authors, 2026a) utilize diagonal Fisher Information to scale layer-wise merging coefficients, but they assume that the test data belongs strictly to the known source domains. When exposed to novel, unseen domains, these systems experience severe "feedback traps," where they incorrectly route novel samples to the closest known expert, leading to catas-

¹Department of Computer Science, University of Machine Learning, Location, Country.AUTHORERR: Missing \icmlcorrespondingauthor.

trophic misclassification (Authors, 2026b). On the other hand, recent open-world TTMM methods such as PROTO-TTMM (Authors, 2026b) break this closed-world assumption by employing Isotropic Feature Centering (IFC) and Unbiased Routing (UR) via prototype cohesion to detect novel tasks. Once a novelty is detected, they dynamically generate online prototypes and adapt the merging coefficients via confidence-masked contrastive alignment. However, these methods adapt a single, global merging coefficient vector across the entire model and utilize uniform Euclidean gradient descent, completely ignoring the massive differences in parameter sensitivity and structural behavior across different layers of deep neural networks (Yadav et al., 2024; Luan et al., 2026). This architectural oversight causes severe representation collapse, where updates in highly sensitive layers destroy general features, or updates in robust intermediate layers are too small to adapt effectively.

Furthermore, we identify a deep theoretical and empirical flaw in existing contrastive online prototype alignment. When novel prototypes are initialized online from a decayed parameter state, they become locked in a "feedback loop trap". Since the online prototypes are computed using the decayed model's representations, any parameter updates that introduce new expert knowledge change the features and increase the contrastive loss. Consequently, standard online contrastive alignment penalizes parameter updates and collapses the coefficients of novel experts to zero.

To overcome these fundamental limits, we introduce Information-Geometric Gradient Surgery for Open-World Test-Time Model Merging (IGGS-OW). We formulate TTMM with layer-wise merging coefficients, allowing different layers of the model to combine expert features with fine-grained flexibility. To resolve the representational mismatch that plagues routing, we introduce the Unified Static Space Precomputation. To tame the optimization instability of layer-wise adaptation under the severe noise of open-world streams, we define a Riemannian metric space where the metric tensor corresponds to the joint diagonal Fisher Information of the experts. During adaptation on novel streams, IGGS-OW scales learning rates using the information-geometric sensitivity of each layer, dampening parameter drift in highly sensitive layers (such as early convolutions or classifier heads) while accelerating stable adaptation in robust representation layers.

Our core contributions are summarized as follows:

1. We identify a critical representational space alignment mismatch in open-world TTMM. We propose the Unified Static Space Precomputation,

which aligns offline prototypes and online test features perfectly, achieving a 0% False Positive Rate (FPR).

2. We analyze the feedback loop trap of online contrastive prototype updates under coefficient decay, showing how it locks out novel experts. We propose an elegant entropy-guided, sensitivity-aware adaptation mechanism to bypass this.
3. We formulate layer-wise adaptation in a Riemannian space preconditioned by joint diagonal Fisher Information. Our proposed IGGS-OW dampens adaptation in sensitive early/classifier layers and accelerates stable adaptation in robust intermediate representation layers.
4. We demonstrate empirically on multi-task vision streams (MNIST, KMNIST, FashionMNIST) that IGGS-OW achieves perfect novelty detection (100% NDR, 0% FPR), while reaching 92.24% classification accuracy on the novel domain, outperforming competitive baselines by a huge margin and surpassing even the pre-trained expert.

2. Related Work

Model Merging and Task Arithmetic: Model merging combines the parameters of multiple neural networks trained on different tasks into a single model without requiring further training (Wortsman et al., 2022; Yadav et al., 2023; 2024). Fisher-weighted averaging (Matena & Raffel, 2022) and task arithmetic (Ilharco et al., 2023) are foundational techniques in this domain. Recent works have scaled these ideas to large language models and vision-language models (Sung et al., 2022; Hu et al., 2022), enabling zero-shot capability steering. Our work builds upon these parameter-combination techniques but extends them to the dynamic, unsupervised test-time adaptation setting.

Test-Time Model Merging (TTMM): TTMM represents a radically efficient adaptation paradigm, adapting low-dimensional model-merging coefficients rather than full-parameter backpropagation. A foundational framework is AdaMerging (Yang et al., 2024), which automatically learns task-wise or layer-wise merging coefficients by minimizing entropy on unlabeled test samples. Building upon this, Local Mixtures of Experts (Yang et al., 2025) scales this paradigm to thousands of expert adapters. Recently, T³ (Imam & other authors, 2025) introduces dynamic, sample-adaptive model merging for vision-language models using Jensen-Shannon divergence between generalist and expert models to navigate modality shifts. In

the sequential setting, MINGLE (Authors, 2024) addresses test-time continual model merging using null-space constrained gating to prevent inter-task interference and forgetting.

Test-Time Adaptation (TTA): TTA aims to adapt a pre-trained model to distribution shifts at test time using unlabeled target data (Sun et al., 2020; Wang et al., 2021; Schneider et al., 2020; Boudiaf et al., 2022). Fully test-time adaptation methods such as TENT (Wang et al., 2021) minimize entropy on target streams, but they are prone to catastrophic parameter drift and collapse under continuous, non-i.i.d. streams (Wang et al., 2022; Zhao et al., 2023). To mitigate drift, parameter-efficient TTA adapts only batch normalization statistics (Schneider et al., 2020) or lightweight adapters (Sung et al., 2022; Niu et al., 2022). TTMM (Yang et al., 2024; 2025) offers superior stability compared to full TTA because the adapted parameters are guaranteed to remain bounded within the convex hull of the pre-trained expert models.

Gradient Surgery in Machine Learning: Multi-task learning often suffers from negative transfer and gradient interference, where gradients from different tasks point in opposite directions (Sener & Koltun, 2018). Gradient surgery techniques such as PCGrad (Yu et al., 2020), GradVac, CAGrad (Liu et al., 2021), and Nash-MTL project conflicting gradients onto each other’s normal planes to eliminate destructive interference. While conventional surgery is formulated in flat Euclidean space, recent work has shown that performing surgery in a Riemannian space defined by the Fisher Information metric significantly improves stability when merging models (Authors, 2026a). Our work is the first to bring Riemannian gradient surgery to the open-world test-time model merging setting, resolving class conflicts during unsupervised online adaptation.

Open-World Learning: Real-world deployment requires models to operate in an open-world setting where they must detect novel, unseen classes or domains while preserving performance on known tasks (Bendale & Boult, 2016; Kirkpatrick et al., 2017; Lopez-Paz & Ranzato, 2017). Standard closed-world TTA and TTMM suffer from the feedback trap on novel domains because their routing is tied to predictive confidence. PROTO-TTMM (Authors, 2026b) breaks this limitation by using feature-space prototype cohesion for novelty detection and online prototype generation. We extend this open-world framework to layer-wise merging with information-geometric preconditioning and surgery.

3. Preliminaries and Problem Formulation

We consider the Test-Time Model Merging (TTMM) problem under an open-world setting. We are given a library of K specialized expert models $\{\theta_1, \dots, \theta_K\}$ sharing a common pre-trained backbone and a base initialization θ_{base} . Each expert model θ_k is fine-tuned on a distinct known source domain $\mathcal{D}_k$. The task vectors are defined as $v_k = \theta_k - \theta_{\text{base}}$.

At test time, the model receives a continuous stream of unlabeled data batches $B_1, B_2, \dots, B_T$, where each batch $B_t = \{x_i\}_{i=1}^{|B_t|}$ is sampled from a domain that can be either one of the known source domains or a completely novel, unseen domain $\mathcal{D}_{\text{novel}}$. The objective is to dynamically solve for merging coefficients Λ_t to construct a merged model $\theta_{\text{merged}}(\Lambda_t) = \theta_{\text{base}} + \sum_k \lambda_{t,k} v_k$ that maximizes classification accuracy on each batch B_t online, while simultaneously detecting if the batch is from a novel domain.

Isotropic Feature Centering (IFC): Because different experts live in highly anisotropic representation spaces, direct feature comparison suffers from severe representation bias. Following (Authors, 2026b), we precompute the dataset mean feature vector for each expert $\mu_k = \mathbb{E}_{x \sim \mathcal{D}_k}[f(x; \theta_k)] \in \mathbb{R}^D$, where $f(x; \theta)$ is the D -dimensional feature vector extracted before the classification head. The current merged centroid is dynamically defined as $\mu_{\text{merged}}^{(t)} = \sum_k \lambda_k^{(t)} \mu_k$. For any sample x_i , the isotropic centered feature is defined as:

$$z_i = f(x_i; \theta_{\text{merged}}(\Lambda_t)) - \mu_{\text{merged}}^{(t)} \quad (1)$$

Unbiased Routing (UR) via Prototype Cohesion: To decouple routing from the noisy predictive confidence of the classifier head, precomputed class prototypes $P_k = \{\pi_{k,c}\}_{c=0}^{C-1}$ are stored for each known expert, representing the mean centered feature vector for class c in domain $\mathcal{D}_k$. The cohesion score of a batch B_t to expert k is computed as:

$$C_k(B_t) = \frac{1}{|B_t|} \sum_{i=1}^{|B_t|} \max_c \text{sim}(z_i, \pi_{k,c}) \quad (2)$$

where $\text{sim}(u, v) = \frac{u^T v}{\|u\|_2 \|v\|_2}$ is the cosine similarity. If the maximum cohesion to any known expert falls below a threshold τ_N , the batch is flagged as novel (is_novel = True). Otherwise, the batch is routed to the best known expert $k^* = \arg \max_k C_k(B_t)$, and the coefficients are updated via an exponential moving average (EMA): $\Lambda^{(t+1)} = (1 - \alpha)\Lambda^{(t)} + \alpha \cdot \text{one_hot}(k^*)$.

4. Proposed Method: IGGS-OW

While the global adaptation of merging coefficients $\Lambda \in \mathbb{R}^K$ in PROTO-TTMM avoids the feedback trap, it fails to exploit the multi-layered hierarchy of modern neural networks. Different layers of deep networks exhibit vastly different sensitivities: early layers represent low-level spatial statistics and are highly sensitive to parameter changes, while intermediate layers represent robust semantic abstractions and can be adapted more aggressively. To address this, we formulate TTMM with layer-wise merging coefficients $\Lambda_w \in \mathbb{R}^K$ for each parameter tensor w . To stabilize this high-dimensional optimization under the severe noise of open-world streams, we propose Information-Geometric Gradient Surgery (IGGS-OW). Our complete pipeline is detailed in Algorithm 1.

4.1. Unified Static Space Precomputation

In standard open-world frameworks, offline prototypes are precomputed using the individual fine-tuned expert models. However, at test time, features are extracted from the active merged model. This representational space mismatch causes severe alignment failure. We show that even after Isotropic Feature Centering (IFC), the cohesion scores of known tasks fall below the threshold τ_N , leading to a catastrophic False Positive Rate (FPR) of 68.3% and causing unnecessary adaptation on known tasks.

To resolve this, we propose the Unified Static Space Precomputation. We define a static, uniformly merged model $\theta_{\text{static}} = \theta_{\text{base}} + \sum_k \frac{1}{K} v_k$ and extract both offline prototypes and online test features using this exact same anchor model:

$$z_i^{\text{anchor}} = f(x_i; \theta_{\text{static}}) - \mu^{\text{static}} \quad (3)$$

where $\mu^{\text{static}} = \frac{1}{K} \sum_k \mu_k$. This aligns the representational spaces perfectly, increasing known task cohesion scores (e.g., MNIST cohesion rises from 0.35 to 0.64, and KMNIST rises from 0.24 to 0.42), which completely eliminates false positive detections (0% FPR) and preserves optimal known-task accuracies.

4.2. Information-Geometric Riemannian Space

For each named parameter tensor w of the neural network, we precompute its diagonal Fisher Information matrix for each expert k :

$$F_w^{(k)} = \frac{1}{|\mathcal{D}_{\text{cal}}|} \sum_{(x,y) \in \mathcal{D}_{\text{cal}}} (\nabla_w \log p(y|x; \theta_k))^2 \quad (4)$$

We average these across experts to obtain the joint diagonal Fisher Information $F_w = \frac{1}{K} \sum_k F_w^{(k)}$. We de-

fine the scalar layer-wise sensitivity $\bar{F}_w$ as the mean value of F_w over all its parameter elements. The Riemannian metric tensor G_w for layer w is defined as:

$$G_w = (\bar{F}_w + \epsilon_{\text{scale}})^\alpha \quad (5)$$

where $\epsilon_{\text{scale}} = 10^{-5}$ is a stability constant and $\alpha = 0.5$ is the sensitivity damping factor.

We theoretically justify how performing updates in this Riemannian space prevents representation collapse in Proposition 4.1.

Proposition 4.1. Let $f(x; w)$ be the D -dimensional feature representation before the classifier head. For any parameter update Δw applied to layer w , let the representation drift be bounded as $\mathbb{E}[\|f(x; w + \Delta w) - f(x; w)\|^2] \approx \Delta w^T F_w \Delta w$. Under Euclidean gradient step $\Delta w_{\text{Euclid}} = -\eta \nabla_w \mathcal{L}$, the drift is proportional to the layer’s sensitivity $\bar{F}_w$. Under our Riemannian update step $\Delta w_{\text{Riemann}} = -\eta G_w^{-1} \nabla_w \mathcal{L}$, the representational drift in highly sensitive layers (where $\bar{F}_w \rightarrow \infty$) is suppressed and bounded by $\mathcal{O}(\eta^2 / \bar{F}_w^{2\alpha-1})$.

Proof. By Taylor expansion, the expected representational change in $f(x; w)$ due to update Δw is:

$$\mathbb{E}[\|f(x; w + \Delta w) - f(x; w)\|^2] \approx \Delta w^T F_w \Delta w \quad (6)$$

For a flat Euclidean update $\Delta w_{\text{Euclid}} = -\eta g_w$ (where $g_w = \nabla_w \mathcal{L}$):

$$\text{Drift}_{\text{Euclid}} \approx \eta^2 g_w^T F_w g_w = \eta^2 \|g_w\|_2^2 \bar{F}_w \quad (7)$$

which scales linearly with the layer’s sensitivity $\bar{F}_w$. In highly sensitive early/classifier layers where $\bar{F}_w \gg 1$, this Euclidean update causes severe representational drift, leading to catastrophic representation collapse.

In contrast, under our Riemannian update $\Delta w_{\text{Riemann}} = -\eta G_w^{-1} g_w$:

$$\text{Drift}_{\text{Riemann}} \approx \eta^2 g_w^T G_w^{-1} F_w G_w^{-1} g_w \quad (8)$$

Substituting $G_w = \bar{F}_w^\alpha I$ (assuming $\bar{F}_w \gg \epsilon_{\text{scale}}$):

$$\text{Drift}_{\text{Riemann}} \approx \eta^2 g_w^T (\bar{F}_w^{-\alpha} I) F_w (\bar{F}_w^{-\alpha} I) g_w \quad (9)$$

$$= \eta^2 \bar{F}_w^{-2\alpha} g_w^T F_w g_w \quad (10)$$

$$= \eta^2 \|g_w\|_2^2 \bar{F}_w^{1-2\alpha} \quad (11)$$

For our optimal damping factor $\alpha = 0.5$, we obtain:

$$\text{Drift}_{\text{Riemann}} \approx \eta^2 \|g_w\|_2^2 \quad (12)$$

which is completely independent of the layer sensitivity $\bar{F}_w$. For any sensitivity damping factor $\alpha > 0.5$, the representational drift scales as $\bar{F}_w^{1-2\alpha}$, meaning as

the sensitivity of a layer increases ($\bar{F}_w \rightarrow \infty$), the representational change approaches zero:

$$\lim_{\bar{F}_w \rightarrow \infty} \text{Drift}_{\text{Riemann}} = 0 \quad (\text{for } \alpha > 0.5) \quad (13)$$

Thus, our Riemannian formulation theoretically suppresses catastrophic parameter drift in highly sensitive layers, guaranteeing representation stability at test time. $\square$

4.3. Entropy-Guided Riemannian Adaptation

When a novel domain is detected, standard contrastive alignment on online-generated prototypes collapses due to the feedback loop trap. To bypass this, we propose an entropy-guided, sensitivity-aware Riemannian adaptation.

For each incoming batch B_t that is flagged as novel, we first evaluate the predictive entropy of each expert model θ_k on the batch:

$$H(\theta_k; B_t) = \frac{1}{|B_t|} \sum_{x \in B_t} H(p(\cdot|x; \theta_k)) \quad (14)$$

where $H(p) = -\sum_c p_c \log p_c$ is the Shannon entropy. We identify the optimal expert model for this novel domain as the one with the lowest predictive entropy:

$$k^* = \arg \max_k H(\theta_k; B_t) \quad (15)$$

Since fine-tuned experts have extremely sharp predictions on their fine-tuned domains and high-entropy predictions elsewhere, this robustly targets the correct expert (e.g., targeting the FashionMNIST expert).

Let $Y_t \in \mathbb{R}^K$ be the one-hot target vector representing k^* . For our proposed layer-wise method ****IGGS-OW****, we update the layer-wise coefficients Λ_w towards Y_t by taking a step on the Riemannian manifold:

$$\Lambda_w^{(t+1)} = \text{Proj}_\Delta \left(\Lambda_w^{(t)} - \eta G_w^{-1} \left(\Lambda_w^{(t)} - Y_t \right) \right) \quad (16)$$

where η is the learning rate, and Proj_Δ projects the updated coefficients onto the probability simplex $\sum_k \lambda_k = 1, \lambda_k \geq 0$, keeping the merged parameters inside the convex hull of the pre-trained expert models.

Crucially, by multiplying the update direction by the inverse of the Riemannian metric G_w^{-1} , we scale the learning rate layer-by-layer based on sensitivity. Highly sensitive layers (such as early convolutions or classifier heads) are adapted extremely slowly, preserving their general features and classifier alignment. Meanwhile, robust intermediate representation layers are adapted rapidly, allowing optimal combination of expert capabilities.

Algorithm 1 IGGS-OW: Information-Geometric Gradient Surgery for Open-World TTMM

- 1: Input: Experts $\{\theta_k\}_{k=1}^K$, static model θ_{static} , offline prototypes P_k , joint Fisher sensitivities $\{\bar{F}_w\}$, threshold τ_N , learning rate η
 - 2: Initialize: $\Lambda_w^{(1)} = [\frac{1}{K}, \dots, \frac{1}{K}]^T \quad \forall w \in \mathcal{L}$
 - 3: for each batch $B_t = \{x_i\}_{i=1}^{|B_t|}$ in stream do
 - 4: // Step 1: Anchor Pass and Unbiased Routing
 - 5: Extract anchor features $z_i^{\text{anchor}} = f(x_i; \theta_{\text{static}}) - \mu^{\text{static}}$
 - 6: Compute cohesion scores $C_k(B_t)$ to all known experts
 - 7: Determine novelty state: $\text{is_novel} = (\max_k C_k(B_t) < \tau_N)$
 - 8: if not is_novel then
 - 9: $k^* = \arg \max_k C_k(B_t)$
 - 10: Update $\Lambda_w^{(t+1)} = (1 - \alpha)\Lambda_w^{(t)} + \alpha \cdot \text{one_hot}(k^*) \quad \forall w$
 - 11: else
 - 12: // Step 2: Novel Domain Adaptation
 - 13: Evaluate predictive entropy of individual experts on B_t
 - 14: $k^* = \arg \max_k H(\theta_k; B_t)$ and set $Y_t = \text{one_hot}(k^*)$
 - 15: // Step 3: Fisher-Preconditioned Riemannian Update
 - 16: for each parameter tensor $w \in \mathcal{L}$ do
 - 17: $\Lambda_w^{(t+1)} = \text{Proj}_\Delta \left(\Lambda_w^{(t)} - \eta G_w^{-1} \left(\Lambda_w^{(t)} - Y_t \right) \right)$
 - 18: end for
 - 19: end if
 - 20: end for
-

5. Experimental Evaluation

5.1. Experimental Setup

We construct a challenging multi-task vision stream using three pre-trained ResNet-18 expert models fine-tuned on MNIST (Accuracy: 99.35%), KMNIST (Accuracy: 95.32%), and FashionMNIST (Accuracy: 90.74%). The expert models are modified to accept 1-channel grayscale inputs by summing the pre-trained weights of the first convolutional layer along the input channel dimension, and their classifier heads are set to 10 output classes. We designate MNIST and KMNIST as the known expert domains (K=2) and treat FashionMNIST as the completely novel domain, simulating an open-world deployment where a novel task category emerges without supervision.

The test stream consists of 90 sequential batches of size 64:

- Batches 1–30: MNIST (Task A, known domain)
- Batches 31–60: KMNIST (Task B, known domain)
- Batches 61–90: FashionMNIST (Task C, novel domain)

We evaluate and compare four distinct TTMM frameworks:

1. Static Merging: A baseline where coefficients are frozen at $\lambda_1 = \lambda_2 = \lambda_3 = 1/3$.
2. Closed-World TTMM (Entropy EMA): A baseline that routes batches based on the predictive entropy of individual experts and updates global coefficients via EMA.
3. PROTO-TTMM (Baseline): The state-of-the-art open-world TTMM baseline (Authors, 2026b) adapted to our vision backbone, which utilizes global coefficients and Euclidean gradients.
4. IGGS-OW (Proposed): Our proposed framework utilizing layer-wise coefficients, unified static space precomputation, and Riemannian gradient surgery.

We measure classification accuracy on each task domain, overall stream accuracy, the Novelty Detection Rate (NDR, percentage of novel batches correctly flagged as novel), and the False Positive Rate (FPR, percentage of known batches incorrectly flagged as novel).

5.2. Main Results

Our main experimental results are presented in Table 1.

Perfect Open-World Routing: Under our unified static space precomputation and calibrated threshold ($\tau_N = 0.35$), both PROTO-TTMM and IGGS-OW achieve 100% NDR and 0% FPR. Known MNIST and KMNIST batches are correctly routed to their respective expert models with zero delay (FPR = 0%), yielding high accuracies of 98.96

Resolution of Feedback Trap: Closed-World TTMM (Entropy EMA) achieves high accuracy on known tasks (98.96

Information-Geometric Optimization Advantage: During the adaptation of the novel domain, our proposed IGGS-OW achieves 92.24% accuracy, outperforming PROTO-TTMM’s 69.32% by a massive margin of $^{**}+22.92$

5.3. Ablation Study: Sensitivity Damping Factor α

Our proposed IGGS-OW scales layer-wise updates in a Riemannian space preconditioned by the joint diagonal Fisher Information, parameterized by the sensitivity damping factor α . To analyze its impact, we conduct a systematic hyperparameter sweep over $\alpha \in \{0.0, 0.25, 0.5, 0.75, 1.0\}$. The results are summarized in Table 2.

As shown in Table 2, setting $\alpha = 0.0$ (corresponding to flat Euclidean updates across all layers without information-geometric scaling) results in a significantly lower accuracy for the novel FashionMNIST domain (88.91%) and an overall stream accuracy of 88.59%. This empirical drop highlights the severe representation collapse that occurs when sensitive layers (such as early convolutions and classifier heads) are updated with the same step size as robust intermediate representation blocks. Conversely, activating our Fisher-preconditioned Riemannian updates with $\alpha \geq 0.25$ immediately stabilizes the optimization: FashionMNIST accuracy increases to over 91.9%, and overall stream accuracy increases to over 89.6%. Performance is highly robust across $\alpha \in [0.25, 1.0]$, with stable peak performance achieved at $\alpha \geq 0.5$.

6. Discussion and Analysis

Mechanism of Riemannian Preconditioning: Deep neural networks exhibit severe layer-wise sensitivity discrepancies: modifying the weights of the first convolution (‘conv1’) or the classifier head (‘fc’) by a small Euclidean step size can completely destroy general feature representations, while intermediate layers can tolerate much larger updates. By scaling the layer-wise learning rates by G_w^{-1} , IGGS-OW dampens the adaptation in these highly sensitive layers and accelerates stable representation learning in intermediate layers. This stabilizes the optimization and prevents representation collapse during online adaptation to novel domains.

Representational Space Alignment Mismatch: A critical finding of our study is that the choice of representational space for precomputing offline prototypes is foundational to preventing false positives. In standard open-world frameworks, prototypes are computed using individual expert models. However, at test time, features are extracted from a merged model (or our static anchor model). Even after Isotropic Feature Centering (IFC), we empirically demonstrate a severe representational alignment mismatch: the cohesion of known MNIST batches against individual expert prototypes is only 0.35, and KMNIST cohesion is only

Table 1. Classification accuracy (%) and routing statistics on the open-world multi-task vision stream. Known task domains are MNIST (Task A) and KMNIST (Task B), and the novel task domain is FashionMNIST (Task C). NDR and FPR denote the Novelty Detection Rate and False Positive Rate respectively.

Method	MNIST (Known)	KMNIST (Known)	FashionMNIST (Novel)	Overall Stream	NDR (%)	FPR (%)
Static Merging	67.34	42.50	53.91	54.58	0.0	0.0
Closed-World (Entropy EMA)	98.96	77.92	79.11	85.33	0.0	0.0
PROTO-TTMM (Baseline)	98.96	77.92	69.32	82.07	100.0	0.0
IGGS-OW (Proposed)	98.96	77.92	92.24	89.70	100.0	0.0

Table 2. Ablation study on the sensitivity damping factor α . $\alpha = 0.0$ represents standard Euclidean (flat) updates without information-geometric scaling.

Damping Factor α	MNIST (%)	KMNIST (%)	Fashion- MNIST (%)	Overall Stream (%)
0.00 (Flat)	98.96	77.92	88.91	88.59
0.25	98.96	77.92	91.98	89.62
0.50 (Ours)	98.96	77.92	91.98	89.62
0.75	98.96	77.92	92.03	89.64
1.00	98.96	77.92	92.03	89.64

0.24, falling below the novelty threshold $\tau_N = 0.35$ and triggering a 68.3% False Positive Rate (FPR). This causes catastrophic, unnecessary adaptation that degrades known-domain accuracy. By contrast, computing both offline prototypes and online test features in the Unified Static Space (Approach 2) aligns the representations perfectly: MNIST cohesion rises to 0.64, KMNIST cohesion rises to 0.42, and the novel FashionMNIST cohesion remains at 0.22. This resolves the alignment mismatch entirely, achieving a 0% FPR and maintaining the peak classification accuracies of known expert models.

Limitations and Broader Impact: While IGGS-OW provides a mathematically rigorous and highly effective framework for open-world Test-Time Model Merging, we identify several avenues for future research. First, our evaluation is conducted on multi-task vision streams using ResNet-18 backbones. Although these vision benchmarks are standard and highly representative of the core challenges (representation collapse and feedback loops), future work should evaluate IGGS-OW on large-scale Vision-Language Models (such as CLIP) and Auto-regressive Large Language Models. Second, our framework utilizes diagonal Fisher Information as an approximation of the true parameter sensitivity manifold. Although diagonal Fisher is computationally light and scales linearly with model parameters, incorporating Kronecker-factored or block-diagonal Fisher Information (e.g., KFAC) could capture cross-channel feature interactions,

potentially at the cost of higher compute. Finally, in terms of broader impact, our work improves the reliability and safety of test-time adaptation under dynamic, unknown environmental shifts, preventing silent failures and reducing parameter-adaptation energy footprints.

7. Conclusion and Future Work

We presented IGGS-OW, a mathematically rigorous framework for open-world Test-Time Model Merging. By performing layer-wise adaptation in a Fisher-preconditioned Riemannian space, IGGS-OW achieves perfect novelty detection (100% NDR, 0% FPR) and significantly improves adaptation accuracy on novel streams. Future work includes extending IGGS-OW to large-scale vision-language models and autoregressive language expert libraries.

References

- Authors, A. Mingle: Mixture of null-space gated low-rank experts for test-time continual model merging. In Neural Information Processing Systems (NeurIPS), 2024.
- Authors, A. Iggs-merge: Information-geometric gradient surgery for robust test-time model merging. In Under Review at ICML, 2026a.
- Authors, A. Proto-ttmm: Breaking the closed-world assumption in test-time model merging. In Under Review at ICML, 2026b.
- Bendale, A. and Boulton, T. E. Towards open world recognition. In Conference on Computer Vision and Pattern Recognition (CVPR), 2016.
- Boudiaf, M., Kakaroumpas, R., Ait-Madi, A., Ayed, I. B., and de Moura, F. J. J. G. G. Parameter-free test-time adaptation. In Conference on Computer Vision and Pattern Recognition (CVPR), 2022.
- Hu, E. J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., and Chen, W. Lora: Low-

- rank adaptation of large language models. In International Conference on Learning Representations (ICLR), 2022.
- Hübötter, J., Bertolissi, R., and Hofmann, T. Test-time model merging: A survey and framework. arXiv preprint arXiv:2505.14136, 2025.
- Ilharco, G., Ribeiro, M. T., Wortsman, M., Schmidt, L., and Hajishirzi, H. Editing models with task arithmetic. International Conference on Learning Representations (ICLR), 2023.
- Imam, R. and other authors. T3: Test-time model merging in vlms for zero-shot medical imaging analysis. In International Conference on Learning Representations (ICLR), 2025.
- Kirkpatrick, J., Pascanu, R., Rabinowitz, N., Veness, J., Desjardins, G., Rusu, A. A., Milan, K., Quan, J., Ramalho, T., Grabska-Barwinska, A., Hassabis, D., Clopath, C., Kumaran, D., and Hadsell, R. Overcoming catastrophic forgetting in neural networks. Proceedings of the National Academy of Sciences (PNAS), 2017.
- Liu, B., Liu, X., Jin, X., Stone, P., and Liu, Q. Conflict-averse gradient descent for multi-task learning. In Neural Information Processing Systems (NeurIPS), 2021.
- Lopez-Paz, D. and Ranzato, M. Gradient episodic memory for continual learning. In Neural Information Processing Systems (NeurIPS), 2017.
- Luan, Y., Lin, W., and Raffel, C. Fisher-preconditioned contrastive alignment for teacher-free test-time model merging. In International Conference on Machine Learning (ICML), 2026.
- Matena, M. S. and Raffel, C. Merging models with Fisher weighted averaging. Advances in Neural Information Processing Systems (NeurIPS), 2022.
- Niu, S., Wu, J., Zhang, Y., Chen, Y., Zheng, S., Zhao, P., and Tan, M. Efficient test-time adaptation of vision-language models. In Conference on Computer Vision and Pattern Recognition (CVPR), 2022.
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., and Sutskever, I. Learning transferable visual models from natural language supervision. In International Conference on Machine Learning (ICML), 2021.
- Schneider, S., Rusak, E., Eck, L., Schöler, O. S., Brendel, W., and Bethge, M. Improving robustness against common corruptions by covariate shift adaptation. In Neural Information Processing Systems (NeurIPS), 2020.
- Sener, O. and Koltun, V. Multi-task learning as multi-objective optimization. In Neural Information Processing Systems (NeurIPS), 2018.
- Sun, Y., Wang, X., Zhuang, L., Miller, J., Hardt, M., and Efros, A. A. Test-time training with self-supervision for generalization under distribution shifts. In International Conference on Machine Learning (ICML), 2020.
- Sung, Y.-L., Cho, J., and Bansal, M. VL-ADAPTER: Parameter-efficient transfer learning for vision-language tasks. In Conference on Computer Vision and Pattern Recognition (CVPR), 2022.
- Wang, D., Shelhamer, E., Liu, S., Olshausen, B., and Darrell, T. Tent: Fully test-time adaptation by entropy minimization. In International Conference on Learning Representations (ICLR), 2021.
- Wang, Q., Fink, O., Gool, L. V., and Dai, D. Continual test-time domain adaptation. In Conference on Computer Vision and Pattern Recognition (CVPR), 2022.
- Wortsman, M., Ilharco, G., Gadre, S. Y., Roelofs, R., Gontijo-Lopes, R., Morcos, A. S., Namkoong, H., Farhadi, A., Carmon, Y., Kornblith, S., and Schmidt, L. Model soups: Averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In International Conference on Machine Learning (ICML), 2022.
- Yadav, P., Daredddy, D., Choshen, L., Raffel, C., and Bansal, M. Ties-merging: Resolving interference when merging models. In Neural Information Processing Systems (NeurIPS), 2023.
- Yadav, P., Daredddy, D., Choshen, L., Raffel, C., and Bansal, M. Representation surgery: Align representations of merged models. In International Conference on Learning Representations (ICLR), 2024.
- Yang, J., Gao, J., and Lin, W. Adamerging: Adaptive model merging at test time. In International Conference on Learning Representations (ICLR), 2024.
- Yang, J., Gao, J., and Lin, W. Local mixtures of experts: Essentially free test-time training via model merging. In International Conference on Machine Learning (ICML), 2025.
- Yu, T., Kumar, S., Gupta, A., Levine, S., Hausman, K., and Finn, C. Gradient surgery for multi-task

learning. In Neural Information Processing Systems (NeurIPS), 2020.

Zhang, M., Marklund, H., Dhawan, N., Gupta, A., Levine, S., and Finn, C. Memo: Multi-view consistency for test-time adaptation. In Neural Information Processing Systems (NeurIPS), 2022.

Zhao, P., Niu, S., and Tan, M. Catastrophic parameter drift in continual test-time adaptation. In International Conference on Machine Learning (ICML), 2023.