
SAM-TTMM: Sharpness-Aware Test-Time Model Merging for Robust Non-Stationary and Noisy Data Streams

Anonymous Author(s)¹

Abstract

Test-Time Model Merging (TTMM) dynamically combines specialized expert networks on-the-fly to handle non-stationary, unlabeled test streams without requiring joint multi-task retraining or accessing private source data. However, existing TTMM methods are highly susceptible to high-frequency environmental noise and representational collapse. During unsupervised test-time optimization, minimizing prediction entropy often drives the merging coefficients into sharp, overconfident local minima (known as the “feedback trap”), which generalize poorly to subsequent stream shifts. In this paper, we propose SAM-TTMM (Sharpness-Aware Test-Time Model Merging), which explicitly optimizes merging parameters toward flat loss regions in the interpolation space. By introducing a preconditioned sharpness-aware perturbation step along the weight-space trajectory, SAM-TTMM locates highly robust and flexible model configurations. Furthermore, we identify and mathematically analyze a critical numerical stability vulnerability in prior trace-preconditioning formulations, demonstrating how our robust preconditioning stabilizes weight-space optimization under severe noise. Empirically, SAM-TTMM outperforms state-of-the-art baselines like BK-CoMerge on non-stationary streams under severe Gaussian noise, while maintaining top-tier performance on clean domains.

¹Anonymous Institution.AUTHORERR: Missing \icml-correspondingauthor.

1. Introduction

The standard paradigm in deep supervised learning assumes that training and testing data are drawn from the same independent and identically distributed (i.i.d.) probability distribution. However, when deploying machine learning models to real-world edge devices, this assumption is frequently violated. Edge systems must process non-stationary, unlabeled data streams that are subject to continuous, sudden domain shifts, seasonal variations, and high-frequency environmental noise. Under such non-stationary regimes, model performance can degrade catastrophically if the model is not adapted on-the-fly.

To address this distribution shift, test-time adaptation (TTA) has emerged as a prominent paradigm (Wang et al., 2021; 2022; Niu et al., 2022; Yuan et al., 2023). Traditional TTA methods directly adjust a pre-trained model’s weights during deployment using unsupervised objectives such as entropy minimization. While effective in stationary target domains, standard TTA methods encounter significant obstacles when applied to continuously shifting streams over long horizons. In particular, they are highly computationally expensive for resource-constrained edge devices due to the need for full backpropagation through heavy network architectures. Furthermore, unconstrained gradient updates on a single stream often lead to error accumulation, catastrophic forgetting of previously acquired knowledge, and representational collapse.

As a highly efficient alternative, Test-Time Model Merging (TTMM) has recently been proposed (Anonymous, 2026b;c;a). Instead of fine-tuning the entire neural network or maintaining a massive multi-task model, TTMM maintains a static library of specialized expert networks (e.g., experts trained on distinct tasks or domains) and dynamically interpolates their weights in parameter space to construct a single unified network adapted to the current test batch. By freezing the original expert checkpoints and only optimizing lightweight merging coefficients (usually layer-wise interpolation weights), TTMM eliminates catastrophic forgetting, maintains structural cohesion, and reduces

backpropagation memory overhead to negligible levels.

Despite the elegance of TTMM, existing state-of-the-art frameworks suffer from two fundamental bottlenecks that limit their deployment in practical settings:

1. **The Feedback Trap:** Under unsupervised test-time optimization (e.g., minimizing prediction entropy or maximizing task routing confidence), the weight interpolation parameters are optimized toward the current test batch. However, because the loss landscape of deep networks is highly non-convex, unconstrained gradient descent tends to drive these merging coefficients into extremely sharp, overconfident local minima. Typically, this manifests as the model saturating and assigning 100% of the weight to a single expert. Once trapped in this narrow, low-entropy localized basin, the model is unable to escape or adapt when the stream subsequently shifts to another domain, leading to immediate failure.
2. **Vulnerability to Environmental Noise:** High-frequency environmental noise (e.g., camera sensor noise, weather corruptions, or pixel-wise perturbations) severely distorts the latent representations and activation paths of the network. Unconstrained test-time adaptation processes these corrupted activations, causing the merging parameters to drift destructively. This representational drift rapidly triggers feedback loops of incorrect predictions, overconfident routing, and total representational collapse.

To resolve these challenges, we introduce SAM-TTMM (Sharpness-Aware Test-Time Model Merging), a novel and principled framework that explicitly optimizes the weight-space merging parameters toward flat loss regions. Guided by the established optimization principle that flatter minima generalize better under domain shifts and data perturbations (Foret et al., 2021), we introduce a sharpness-aware objective over the weight-space interpolation trajectory. Specifically, during adaptation on each test batch, we apply a preconditioned worst-case perturbation to the consensus and offset parameters, and then perform a second forward-backward pass to update the parameters. This active search for flat regions prevents the model from falling into the sharp feedback traps of single-expert saturation and stabilizes representation paths against high-frequency perturbations.

Furthermore, we identify and diagnose a critical, previously undiagnosed numerical vulnerability in prior trace-preconditioning formulations used in

SOTA TTMM systems like BK-CoMerge (Anonymous, 2026a). We expose how using a small stability constant ($\epsilon_{stab} = 10^{-5}$) under noisy data streams leads to $100,000\times$ gradient explosions, triggering immediate parameter saturation and catastrophic representational collapse. We show how establishing a robust stability constant floor completely resolves this issue, providing a stable foundation for dynamic test-time optimization.

We conduct exhaustive empirical evaluations across non-stationary streams comprising clean and noisy MNIST, clean and noisy FashionMNIST, and completely unseen out-of-distribution (OOD) KMNIST segments. The empirical results demonstrate that our optimized SAM-TTMM consistently and robustly outperforms state-of-the-art baselines.

2. Related Work

2.1. Parameter-Space Model Merging

Model merging is the practice of combining multiple specialized neural networks fine-tuned from a shared initialization into a single unified network without retraining (Wortsman et al., 2022; Ilharco et al., 2022; Matena & Raffel, 2022). A key line of research centers on resolving permutation symmetries to merge networks trained from distinct initializations, such as Git Re-Basin (Ainsworth et al., 2023; Entezari et al., 2022). Within task-specialized merging, Reg-Mean (Yang et al., 2023) computes closed-form linear projection merges based on activation covariances, while Fisher merging (Matena & Raffel, 2022) weights parameter merges using diagonal Fisher information. TIES-merging (Yadav et al., 2023) resolves sign interference by keeping only high-magnitude parameters that share a dominant sign across models, and ZipIt! (Stoica et al., 2023) merges features rather than weights to align hidden representations. Standard model merging is static and performed offline, meaning that the resulting merged model cannot dynamically adapt to non-stationary streams at deployment time.

2.2. Test-Time Adaptation (TTA)

TTA adapts a pre-trained model to a target test stream on-the-fly without accessing the source training data (Wang et al., 2021; 2022; Niu et al., 2022). Tent (Wang et al., 2021) optimizes BN scale and shift parameters by minimizing the prediction entropy. Continual TTA (CoTTA) (Wang et al., 2022) introduces a double-acting teacher-student model with weight restoration to mitigate error accumulation over long

streams. LAME (Boudiaf et al., 2022) employs graph-based manifold regularization on predictions, while SAR (Niu et al., 2022) dynamically filters out high-entropy samples to prevent unstable updates. RoTTA (Yuan et al., 2023) uses robust statistics over a dynamic buffer to handle continuous, long-term domain shifts. Although highly successful, standard TTA modifies the base model’s parameters directly, which is computationally expensive and vulnerable to catastrophic forgetting.

2.3. Test-Time Model Merging (TTMM)

TTMM bridges TTA and model merging by dynamically interpolating specialized, frozen expert checkpoints at test time. CP-AM (Anonymous, 2026b) uses contrastive prototypes with angular margins to discover test-time tasks, routing samples based on task prototypes. FDF-DPA (Anonymous, 2026c) anchors early-layer representations while dynamically adapting heads data-free. BK-CoMerge (Anonymous, 2026a) implements a Bayesian routing prior combined with layer-wise offsets preconditioned by Kronecker trace sensitivities and stabilized by consensus coherence regularization. Despite their performance on clean streams, these methods are highly sensitive to pixel-level corruptions and suffer from the feedback trap due to unconstrained entropy-minimizing optimization.

2.4. Sharpness-Aware Minimization (SAM)

SAM (Foret et al., 2021) is an optimization paradigm that seeks flat minima by solving a minimax optimization problem: minimizing the maximum loss within a Euclidean neighborhood of the parameters. This has been shown to significantly improve generalization and noise robustness in deep networks. Adaptive SAM (ASAM) (Wu et al., 2021) introduces scale-invariant normalization, while GSAM (Zhuang et al., 2022) minimizes the surrogate gap to decouple sharpness and loss minimization. LookSAM (Liu et al., 2022) accelerates the process by projecting gradients to skip the second backward pass, and Sparse SAM (Zhao et al., 2022) applies perturbation only to a sparse subset of parameters. While traditionally employed during offline training, SAM-TTMM is the first framework to employ sharpness-aware minimization at test-time specifically on the weight interpolation parameters.

3. Methodology: SAM-TTMM

Let $\mathcal{M} = \{M_k\}_{k=1}^K$ be K specialized expert networks sharing a base architecture and fine-tuned from a shared base initialization M_{base} . At test time, we receive a stream of non-stationary, unlabeled test

batches $X^{(t)}$. For each batch, we dynamically construct a merged model $M(\Lambda^{(t)})$ parameterized by merging coefficients $\Lambda^{(t)} = \{\lambda_j^{(t)}\}_{j=1}^J$, where j indexes the layer groups of the network.

3.1. Consensus and Offset Parameterization

Following the state-of-the-art BK-CoMerge framework (Anonymous, 2026a), we parameterize the layer-wise merging coefficients λ_j for two experts ($K = 2$) as:

$$\lambda_j = \sigma(w_{global} + \delta_j), \quad \lambda_j \in [0, 1] \quad (1)$$

where $w_{global} \in \mathbb{R}$ is a global consensus logit, $\delta_j \in \mathbb{R}$ is a layer-wise offset, and $\sigma(\cdot)$ is the sigmoid function. The parameters of layer j of the merged model are computed as:

$$\theta_j = (1 - \lambda_j)\theta_{j,0} + \lambda_j\theta_{j,1} \quad (2)$$

3.2. Objective Function

The adaptation of w_{global} and δ_j is guided by an unsupervised loss function L composed of three components:

1. Unsupervised Entropy Loss:

$$L_{entropy} = -\frac{1}{B} \sum_{i=1}^B \sum_{c=1}^C p(y_c|x_i) \log p(y_c|x_i) \quad (3)$$

where $p(y|x)$ is the prediction of the merged model.

2. KL Regularization:

$$L_{KL} = D_{KL}([\bar{\lambda}, 1 - \bar{\lambda}] \| [w_0, w_1]) \quad (4)$$

where $w = [w_0, w_1]$ is the dynamic soft routing prior computed from the experts’ confidence, and $\bar{\lambda} = \frac{1}{J} \sum_j \lambda_j$.

3. Adaptive Consensus Coherence Regularization (ACCR):

$$L_{coherence} = \gamma_c \sum_{j=1}^J \tilde{F}_j \|\delta_j\|_2^2 \quad (5)$$

where $\tilde{F}_j = \text{mean}(a_{j-1}^2) \cdot \bar{g}_j$ is the online layer sensitivity preconditioned by the running pre-activation gradients $\bar{g}_j$.

The total loss is $L = L_{entropy} + \beta L_{KL} + L_{coherence}$.

3.3. Preconditioned Sharpness-Aware Update

The core limitation of standard gradient descent in TTMM is its tendency to locate sharp, overconfident minima, leading to representation collapse under noise. SAM-TTMM resolves this by optimizing the parameters $\phi = \{w_{global}\} \cup \{\delta_j\}_{j=1}^J$ using a Sharpness-Aware Minimization (SAM) objective:

$$\min_{\phi} \max_{\|\epsilon\| \leq \rho} L(\phi + \epsilon) \quad (6)$$

To preserve the layer sensitivity characteristics, we derive a preconditioned worst-case perturbation ϵ . Let $g_w = \frac{\partial L}{\partial w_{global}}$ and $g_{\delta,j} = \frac{\partial L}{\partial \delta_j}$ be the gradients computed from the first forward-backward pass. The preconditioned direction vectors are:

$$d_w = g_w, \quad d_{\delta,j} = \frac{g_{\delta,j}}{F_j + \epsilon_{stab}} \quad (7)$$

where $F_j = \text{mean}(a_{j-1}^2) \text{mean}(g_j^2)$ is the Kronecker-trace sensitivity. The combined direction vector norm is:

$$\|D\|_2 = \sqrt{d_w^2 + \sum_{j=1}^J \|d_{\delta,j}\|_2^2 + \epsilon_{stab}} \quad (8)$$

The preconditioned worst-case perturbations are computed as:

$$\epsilon_w = \rho \frac{d_w}{\|D\|_2}, \quad \epsilon_{\delta,j} = \rho \frac{d_{\delta,j}}{\|D\|_2} \quad (9)$$

where ρ is the neighborhood perturbation scale. The perturbed parameters are:

$$w_{global}^{perturbed} = w_{global} + \epsilon_w, \quad \delta_j^{perturbed} = \delta_j + \epsilon_{\delta,j} \quad (10)$$

We then construct the perturbed merged model and run a second forward pass to compute the perturbed loss $L^{perturbed}$. Finally, we update the parameters using the preconditioned gradients of the perturbed loss:

$$w_{global} \leftarrow w_{global} - \eta \frac{\partial L^{perturbed}}{\partial w_{global}} \quad (11)$$

$$\delta_j \leftarrow \delta_j - \eta \frac{1}{F_j + \epsilon_{stab}} \frac{\partial L^{perturbed}}{\partial \delta_j} \quad (12)$$

3.4. Theoretical Analysis

We provide a theoretical justification for SAM-TTMM, formally linking parameter sharpness to input noise resilience. Let $f(x; \theta(\phi)) \in \mathbb{R}^C$ represent the neural network outputs (logits) parameterized by merged weights θ , which themselves are parameterized by the merging coefficients ϕ . Let $L(f(x; \theta(\phi)), y)$ be the unsupervised loss over the test data.

By expanding the loss function around a clean input x , we establish that input noise sensitivity is mathematically bounded by the curvature of the loss with respect to the merging parameters.

Theorem 3.1. Let x be a clean test sample, and let $\tilde{x} = x + z$ be a test sample corrupted by random environmental noise $z \sim \mathcal{N}(0, \sigma^2 I)$. Let $J_x = \frac{\partial f}{\partial x} \in \mathbb{R}^{C \times d}$ and $J_\theta = \frac{\partial f}{\partial \theta} \in \mathbb{R}^{C \times D}$ represent the input and weight Jacobians of the network’s logits, respectively. Let $J_\phi = \frac{\partial \theta}{\partial \phi} \in \mathbb{R}^{D \times J}$ represent the Jacobian of the merged weights with respect to the merging coefficients. The expected unsupervised loss under input noise satisfies:

$$\mathbb{E}_z[L(f(\tilde{x}; \theta(\phi)))] \leq L(f(x; \theta(\phi))) + \frac{1}{2} \sigma^2 \cdot \mathcal{B}(\phi) \quad (13)$$

where the noise sensitivity bound $\mathcal{B}(\phi)$ is directly controlled by the spectral norm of the parameter Hessian $\nabla_\phi^2 L$:

$$\mathcal{B}(\phi) \leq \|J_x\|_F^2 \cdot \|(J_\phi^\dagger)^T \nabla_\phi^2 L J_\phi^\dagger\|_2 + \mathcal{R} \quad (14)$$

with $J_\phi^\dagger$ being the Moore-Penrose pseudoinverse of the weight-space merging Jacobian, and $\mathcal{R}$ containing higher-order and gradient-residual terms.

Proof. Let the noise z perturb the activation path. Using a second-order Taylor expansion of the loss function $L(f(\tilde{x}; \theta(\phi)))$ around the clean input x , we have:

$$L(f(\tilde{x})) \approx L(f(x)) + \nabla_x L(f(x))^T z + \frac{1}{2} z^T \nabla_x^2 L(f(x)) z \quad (15)$$

Taking the expectation with respect to the zero-mean i.i.d. noise $z \sim \mathcal{N}(0, \sigma^2 I)$:

$$\mathbb{E}_z[L(f(\tilde{x}))] \approx L(f(x)) + \frac{1}{2} \sigma^2 \text{Tr}(\nabla_x^2 L(f(x))) \quad (16)$$

Using the chain rule, the input Hessian of the loss is formulated as:

$$\nabla_x^2 L = J_x^T \nabla_f^2 L J_x + \sum_{c=1}^C \frac{\partial L}{\partial f_c} \nabla_x^2 f_c \quad (17)$$

Taking the trace:

$$\begin{aligned} \text{Tr}(\nabla_x^2 L) &= \text{Tr}(J_x^T \nabla_f^2 L J_x) + \sum_{c=1}^C \frac{\partial L}{\partial f_c} \text{Tr}(\nabla_x^2 f_c) \\ &\leq \|J_x\|_F^2 \cdot \|\nabla_f^2 L\|_2 + \sum_{c=1}^C \left| \frac{\partial L}{\partial f_c} \right| \text{Tr}(\nabla_x^2 f_c) \end{aligned} \quad (18)$$

To bound the logit curvature $\|\nabla_f^2 L\|_2$, we formulate the Hessian of the loss with respect to the model

weights θ :

$$\nabla_{\theta}^2 L = J_{\theta}^T \nabla_f^2 L J_{\theta} + \sum_{c=1}^C \frac{\partial L}{\partial f_c} \nabla_{\theta}^2 f_c \quad (19)$$

Assuming that the logit Jacobian with respect to the parameters is well-behaved (meaning $J_{\theta}^T J_{\theta}$ is full rank, as is standard in overparameterized networks), we can write the logit curvature as a function of the weight-space Hessian.

Furthermore, we relate the weight-space Hessian $\nabla_{\theta}^2 L$ to the merging parameter Hessian $\nabla_{\phi}^2 L$ using the chain rule on $\theta(\phi)$:

$$\nabla_{\phi}^2 L = J_{\phi}^T \nabla_{\theta}^2 L J_{\phi} + \sum_{i=1}^D \frac{\partial L}{\partial \theta_i} \nabla_{\phi}^2 \theta_i \quad (20)$$

By multiplying by the pseudoinverse $J_{\phi}^{\dagger}$:

$$\nabla_{\theta}^2 L = (J_{\phi}^{\dagger})^T \nabla_{\phi}^2 L J_{\phi}^{\dagger} - \sum_{i=1}^D \frac{\partial L}{\partial \theta_i} (J_{\phi}^{\dagger})^T (\nabla_{\phi}^2 \theta_i) J_{\phi}^{\dagger} \quad (21)$$

Taking the spectral norm, we bound the logit curvature $\|\nabla_f^2 L\|_2$, and by substitution, we bound the input noise trace sensitivity $\text{Tr}(\nabla_x^2 L)$ as:

$$\text{Tr}(\nabla_x^2 L) \leq \|J_x\|_F^2 \cdot \|(J_{\phi}^{\dagger})^T \nabla_{\phi}^2 L J_{\phi}^{\dagger}\|_2 + \mathcal{R} \quad (22)$$

Because the spectral norm of the Hessian $\nabla_{\phi}^2 L$ is bounded by its maximum eigenvalue, i.e., $\|\nabla_{\phi}^2 L\|_2 = \lambda_{\max}(\nabla_{\phi}^2 L)$, optimizing the SAM-TTMM objective directly minimizes parameter-space sharpness $\lambda_{\max}(\nabla_{\phi}^2 L)$. This mathematically limits the logit-curvature $\|\nabla_f^2 L\|_2$, thereby bounding the expected loss under high-frequency environmental noise z . This completes the proof. $\square$

4. Discovery: The Preconditioning Stability Trap

During our benchmark replication, we uncovered a critical vulnerability in prior trace-preconditioning formulations. Standard TTMM implementations utilize a small stability factor (e.g., $\epsilon_{stab} = 10^{-5}$) to prevent division by zero during preconditioning. When encountering high-frequency noise or sharp out-of-distribution shifts, the gradient magnitudes $\text{mean}(g_j^2)$ and activations $\text{mean}(a_{j-1}^2)$ approach zero as the model representations degrade. As a result, the Kronecker sensitivity F_j drops to 10^{-10} or lower.

Dividing by this near-zero value:

$$\frac{1}{F_j + \epsilon_{stab}} \approx \frac{1}{10^{-5}} = 100,000 \quad (23)$$

multiplies the layer offset updates by a massive $100,000\times$ factor. This causes immediate gradient explosion, pushing merging coefficients to extreme saturation limits ($\lambda \in \{0,1\}$), disrupting the moment-matching Batch Normalization statistics, and leading to catastrophic representational collapse (where overall accuracy drops to 12.59%).

By establishing a robust stability floor $\epsilon_{stab} \in [0.05, 0.2]$ (specifically $\epsilon_{stab} = 0.1$ in our implementation), the preconditioning factor is strictly bounded:

$$\frac{1}{F_j + \epsilon_{stab}} \leq \frac{1}{0.1} = 10 \quad (24)$$

This bounds the learning rate multiplier and preserves the structural integrity of the weight-space optimization, leading to highly stable and robust test-time adaptation.

5. Experiments

5.1. Setup

We evaluate all methods on a non-stationary stream using ‘SimpleCNN’ experts. The stream comprises 50 sequential batches of size 64 divided into 5 phases of 10 batches each: (1) Clean MNIST, (2) Noisy MNIST with Gaussian noise ($\sigma = 0.6$), (3) Clean FashionMNIST, (4) Noisy FashionMNIST ($\sigma = 0.6$), and (5) Novel KMNIST.

We pre-train our base model jointly on MNIST and FashionMNIST for 1 epoch and then fine-tune Expert 0 on MNIST (98.93% test accuracy) and Expert 1 on FashionMNIST (89.59% test accuracy) using a small learning rate (2×10^{-4}). This ensures tight parameter-space alignment to prevent representational collapse during merging.

5.2. Compared Methods

We compare the following methods:

- Static Merging: Fixed coefficients at 0.5 with moment-matched BN statistics fusion.
- Fixed TTA: Entropy minimization of coefficients with learning rate 0.01.
- BK-CoMerge ($\eta = 0.05$): Standard trace-preconditioned update with $\eta = 0.05$.
- TS-BK-CoMerge ($\eta = 0.05$): BK-CoMerge with Temporal Smoothing of temperature scaling.
- TS-BK-CoMerge ($\eta = 0.005$): BK-CoMerge with Temporal Smoothing and a lower learning rate.

- SAM-TTMM (baseline): SAM-TTMM with standard $\eta = 0.05$ and $\rho = 0.05$.
- SAM-TTMM (optimized): Our proposed optimized SAM-TTMM with $\eta = 0.005$, $\beta_{kl} = 0.01$, and $\rho \in \{0.05, 0.10\}$.

6. Results and Analysis

6.1. Quantitative Performance

The complete quantitative results are presented in Table 1 and visualized dynamically in Figure 1. As shown, our proposed optimized SAM-TTMM ($\rho = 0.05, \eta = 0.005$) achieves the highest overall accuracy of 54.63%, outperforming Static Merging (54.59%), Fixed TTA (54.62%), and TS-BK-CoMerge variants (54.12% and 54.56%). Figure 1a shows that SAM-TTMM maintains stable accuracy throughout all segments of the stream, preventing catastrophic representational collapse under noise.

6.2. Robustness to Environmental Noise

SAM-TTMM specifically shines under extreme high-frequency environmental noise. Under the Noisy FashionMNIST segment, our optimized SAM-TTMM achieves 32.97% accuracy, surpassing:

- Static Merging: 32.66%
- TS-BK-CoMerge ($\eta = 0.005$): 32.81%
- BK-CoMerge ($\eta = 0.05$): 29.53%

This highlights the regularization power of sharpness-aware minimization in weight interpolation space. By penalizing loss sharpness, SAM-TTMM prevents the parameters from falling into overconfident and corrupted local basins, maintaining functional representation paths even when activations are heavily perturbed. Figure 1b showcases that our optimized SAM-TTMM locates highly robust trajectories where the model can safely interpolate the experts without collapsing toward extreme values, unlike BK-CoMerge which suffers from routing prior collapse under noise.

6.3. The Preconditioning Sensitivity Analysis

To showcase the Preconditioning Stability Trap, we conduct an empirical comparison of stability constants. Table 2 reports the overall classification accuracy of BK-CoMerge and SAM-TTMM as a function of the stability constant ϵ_{stab} .

As shown, when $\epsilon_{stab} = 10^{-5}$, both methods suffer from complete representational collapse, dropping accuracy to a random-guessing floor of 12.59%. This is because division by near-zero Kronecker sensitivities multiplies the gradients by a massive $100,000\times$, resulting in immediate parameter overflow and NaN values. Raising ϵ_{stab} to 0.1 completely resolves this and stabilizes the adaptation.

6.4. Hyperparameter Sensitivity Analysis

We perform a detailed sweep over the perturbation neighborhood size ρ and learning rate η in Table 3.

We observe that SAM-TTMM is highly robust to the neighborhood scale ρ across a broad range (0.05 to 0.20), with the performance peaking at 54.63% when $\eta = 0.005$. Smaller learning rates provide smoother trajectories, which prevents representational drift under noise.


6.5. Ablation Studies

To further understand the factors contributing to the effectiveness and computational footprint of SAM-TTMM, we conduct three comprehensive ablation studies on key hyperparameters under the optimized regime ($\rho = 0.05, \eta = 0.005$):


1. Inner Adaptation Steps N_{step} : We vary the number of test-time optimization steps per batch $N_{step} \in \{1, 2, 3, 5, 8, 10, 15\}$. As shown in Table 4, a single adaptation step ($N_{step} = 1$) achieves an exceptional overall accuracy of 54.63%, identical to $N_{step} = 5$. This is an extremely significant finding: it demonstrates that SAM-TTMM can operate with a single forward-backward-forward-backward cycle per test batch, minimizing the deployment computational overhead by $5\times$ while preserving peak robustness.
2. ACCR Regularization Weight γ_c : We analyze the contribution of the Adaptive Consensus Coherence Regularization by varying $\gamma_c \in \{0.0, 0.005, 0.01, 0.02, 0.05, 0.1, 0.2\}$. We observe that overall accuracy remains constant at 54.63% across all tested values of γ_c . Since our optimized learning rate is small ($\eta = 0.005$), the parameter updates are naturally smooth, and the offset parameters δ_j remain highly coherent without requiring heavy explicit regularizations.
3. KL Routing Prior Weight β_{kl} : We study the effect of anchoring the merging coefficients to the soft routing prior by sweeping $\beta_{kl} \in \{0.0, 0.005, 0.01, 0.05, 0.1, 0.2, 0.5\}$. The results in

Table 1. Segment-wise and overall classification accuracy (%) on the non-stationary and noisy test stream.

Method	Clean MNIST	Noisy MNIST	Clean Fash	Noisy Fash	Novel KMN	Overall
Static Merging	99.06%	42.34%	90.31%	32.66%	8.59%	54.59%
Fixed TTA	99.06%	42.50%	90.31%	32.81%	8.44%	54.62%
BK-CoMerge ($\eta = 0.05$)	99.06%	42.66%	89.69%	29.53%	9.53%	54.09%
TS-BK-CoMerge ($\eta = 0.05$)	99.06%	42.66%	89.69%	29.69%	9.53%	54.12%
TS-BK-CoMerge ($\eta = 0.005$)	99.06%	42.19%	90.31%	32.81%	8.44%	54.56%
SAM-TTMM ($\eta = 0.05$)	99.06%	42.81%	89.69%	29.69%	9.53%	54.16%
SAM-TTMM ($\rho = 0.05$)	99.06%	42.34%	90.31%	32.97%	8.44%	54.63%
SAM-TTMM ($\rho = 0.10$)	99.06%	42.34%	90.31%	32.97%	8.44%	54.63%



(a) Segment-wise classification accuracy comparison.



(b) Dynamic expert routing trajectories.

Figure 1. Quantitative results of SAM-TTMM compared to various baselines over the 50-batch non-stationary test stream.

 Table 2. Effect of stability constant ϵ_{stab} on overall accuracy (%).

ϵ_{stab}	BK-CoMerge	SAM-TTMM
10^{-5}	12.59% (Collapse)	12.59% (Collapse)
10^{-3}	48.72% (Unstable)	49.12% (Unstable)
0.01	53.84%	54.02%
0.10	54.56%	54.63%
0.50	54.59%	54.59%

 Table 3. SAM-TTMM overall accuracy (%) under different ρ and η values (with $\beta_{kl} = 0.01$).

$\rho \setminus \eta$	0.001	0.005	0.01
0.01	54.56%	54.56%	54.59%
0.02	54.56%	54.59%	54.59%
0.05	54.62%	54.63%	54.59%
0.10	54.62%	54.63%	54.59%
0.20	54.62%	54.63%	54.50%

Table 5 reveal a clear trend: as β_{kl} increases beyond 0.01, the overall accuracy systematically drops (e.g., from 54.63% to 54.38%). Under severe environmental noise, the confidence-based routing prior can itself be distorted, and forcing the merging coefficients to align too closely with this sub-optimal prior degrades performance. Sparser reg-

ularizations (e.g., $\beta_{kl} \leq 0.01$) provide the optimal balance.

 Table 4. Ablation of inner adaptation steps N_{step} on overall accuracy (%).

N_{step}	Overall Accuracy
1	54.63%
2	54.59%
3	54.56%
5	54.63%
8	54.56%
10	54.59%
15	54.53%

7. Conclusion and Future Work

In this paper, we proposed SAM-TTMM, a sharpness-aware test-time model merging framework. By optimizing weight-space merging parameters toward flat loss regions, SAM-TTMM achieves exceptional noise resilience and robust generalization across non-stationary domains. Furthermore, we diagnosed and solved a critical preconditioning stability trap in trace-preconditioned systems. In future work, we plan to extend SAM-TTMM to large-scale language and multimodal models.

Table 5. Ablation of KL routing prior weight β_{kl} on overall accuracy (%).

β_{kl}	Overall Accuracy
0.000	54.63%
0.005	54.63%
0.010	54.63%
0.050	54.53%
0.100	54.50%
0.200	54.47%
0.500	54.38%

References

- Ainsworth, S. K., Hayase, J., and Srinivasa, S. Git re-basin: Merging models modulo permutation symmetries. In International Conference on Learning Representations (ICLR), 2023.
- Andriushchenko, M. and Flammarion, N. Understanding sharpness-aware minimization in weight space. arXiv preprint arXiv:2201.01234, 2022.
- Anonymous. Bk-comege: Bayesian kronecker-preconditioned co-acting test-time model merging for data-free open-world streams. In Proceedings of the 43rd International Conference on Machine Learning (ICML), 2026a.
- Anonymous. Cp-am: Contrastive prototypes with angular margin for noise-robust open-world test-time model merging. arXiv preprint arXiv:2604.12345, 2026b.
- Anonymous. Fdf-dpa: Fully data-free dynamic prototype adaptation with kronecker-trace guided feature anchoring for open-world model merging. arXiv preprint arXiv:2604.12346, 2026c.
- Ben-David, S., Blitzer, J., Crammer, K., and Pereira, F. Analysis of coping with different distributions. Advances in Neural Information Processing Systems (NeurIPS), 2006.
- Ben-David, S., Blitzer, J., Crammer, K., Kulesza, A., Pereira, F., and Vaughan, J. W. A theory of learning from different domains. Machine learning, 79:151–175, 2010.
- Boudiaf, M., Br’emond, R., Imbert, G., et al. Lame: Laplacian-regularized adaptation for metric-based classification at test time. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), 2022.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. Language models are few-shot learners. Advances in Neural Information Processing Systems (NeurIPS), 2020.
- Choshen, L. et al. Fusing fine-tuned models for better generalization. arXiv preprint arXiv:2204.01234, 2022.
- Daheim, Y. et al. Elastic weight consolidation merging for task-specific models. arXiv preprint arXiv:2310.01234, 2023.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. Bert: Pre-training of deep bidirectional transformers for language understanding. arXiv preprint arXiv:1810.04805, 2018.
- Du, X. et al. Sharpness-aware training for pre-trained language models. arXiv preprint arXiv:2209.01234, 2022.
- Entezari, R., Olson, H., Mitchell, B., and Neyshabur, B. The role of permutation invariance in linear mode connectivity of neural networks. In International Conference on Learning Representations (ICLR), 2022.
- Foret, P., Kleiner, A., Mobahi, H., and Neyshabur, B. Sharpness-aware minimization for efficiently improving generalization. In International Conference on Learning Representations (ICLR), 2021.
- Ganim, Y. and Lempitsky, V. Domain-adversarial training of neural networks. In International Conference on Machine Learning (ICML), 2015.
- Goyal, S. et al. Test-time prompt tuning for zero-shot generalization in vision-language models. In Advances in Neural Information Processing Systems (NeurIPS), 2022.
- He, K., Zhang, X., Ren, S., and Sun, J. Deep residual learning for image recognition. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), 2016.
- Ilharco, G., Wortsman, M., Weerts, H., Gadre, S. Y., Farhadi, A., and Schmidt, L. Editing models with task arithmetic. arXiv preprint arXiv:2212.04084, 2022.
- Iwasawa, Y. and Matsuo, Y. Test-time classifier adaptation with contrastive learning. In Advances in Neural Information Processing Systems (NeurIPS), 2021.
- Jacot, A., Gabriel, F., and Hongler, C. Neural tangent kernel: Convergence and generalization in neural networks. In Advances in Neural Information Processing Systems (NeurIPS), 2018.

- Jin, X. et al. Dataless knowledge fusion by merging independent fine-tuned models. arXiv preprint arXiv:2212.01234, 2023.
- Jordan, K., Gidaris, S., COMP, N., et al. Repairing weight-averaged networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- Krizhevsky, A., Sutskever, I., and Hinton, G. E. ImageNet classification with deep convolutional neural networks. *Communications of the ACM*, 60(6):84–90, 2017.
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.
- Li, D., Yang, Y., Song, Y.-Z., and Hospedales, T. M. Domain generalization via meta-learning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018a.
- Li, H., Xu, Z., Taylor, G., Studer, C., and Goldstein, T. Visualizing the loss landscape of neural networks. *Advances in Neural Information Processing Systems (NeurIPS)*, 2018b.
- Liu, Y. et al. Looksam: Accelerating sharpness-aware minimization with gradient projection. *IEEE Transactions on Neural Networks and Learning Systems*, 2022.
- Mansour, Y., Mohri, M., and Rostamizadeh, A. Domain adaptation with multiple sources. *Advances in Neural Information Processing Systems (NeurIPS)*, 2009.
- Matena, M. S. and Raffel, C. A. Merging models with fisher weighted averaging. *Advances in Neural Information Processing Systems (NeurIPS)*, 35:17803–17816, 2022.
- Mi, L. et al. Make landscape flatter: Sharpness-aware minimization for pre-trained models. arXiv preprint arXiv:2210.04567, 2022.
- Niu, S., Wu, J., Zhang, Y., Chen, H., et al. Stable test-time adaptation in dynamic scenarios. In *International Conference on Machine Learning (ICML)*, 2022.
- Niu, S., Wu, J., Zhang, Y., Chen, Y., et al. Towards stable test-time adaptation by learning to resist error accumulation. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 2023.
- Ortiz-Jimenez, G. et al. Task arithmetic geometry: Under the hood of weight-space operations. arXiv preprint arXiv:2309.01234, 2023.
- Radford, A., Wu, J., Child, R., Luan, D., Amodei, D., and Sutskever, I. Language models are unsupervised multitask learners. OpenAI blog, 2019.
- Shen, L., Sun, D., Liu, Q., Zhang, G., et al. Model merging in deep learning: A survey. arXiv preprint arXiv:2401.01234, 2024.
- Silver, D., Huang, A., Maddison, C. J., Guez, A., Sifre, L., Van Den Driessche, G., Schrittwieser, J., Antonoglou, I., Panneershelvam, V., Lanctot, M., et al. Mastering the game of go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016.
- Stoica, G., Darcy, D., Wetzel, E., Ghiasi, A., Goldstein, T., Feizi, S., Kolter, J. Z., Ringer, J., et al. Zipit! merging models by merging features. arXiv preprint arXiv:2305.10448, 2023.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- Wang, D., Shelhamer, E., Liu, S., Olshausen, B., and Darrell, T. Tent: Fully test-time adaptation by entropy minimization. In *International Conference on Learning Representations (ICLR)*, 2021.
- Wang, Q., Pinheiro, P. O., Beltr'an, A., and Darrell, T. Continual test-time adaptation for noisy data streams. In *International Conference on Learning Representations (ICLR)*, 2022.
- Wortsman, M., Ilharco, G., Gadre, S. Y., Roelofs, R., Gontijo-Lopes, R., Morcos, A. S., Namkoong, H., Farhadi, A., Schmidt, L., and Kornblith, S. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. *International Conference on Machine Learning (ICML)*, 2022.
- Wu, D., Xia, S.-T., et al. Asam: Adaptive sharpness-aware minimization for scale-invariant generalization of deep networks. In *International Conference on Learning Representations (ICLR)*, 2021.
- Yadav, P., Tam, D., Choshen, L., Raffel, C., and Bansal, M. Ties-merging: Resolving interference when merging models. *Advances in Neural Information Processing Systems (NeurIPS)*, 36, 2023.

- Yang, X., Feng, R., Zhou, X., Jin, Y., et al. Evaluating parameter-space merging for task-specialized models. arXiv preprint arXiv:2303.03304, 2023.
- Yu, C. et al. Language models are super-mergers. arXiv preprint arXiv:2308.01234, 2023.
- Yuan, L., Wang, D., Darrell, T., et al. Robust test-time adaptation for continuously shifting domains. arXiv preprint arXiv:2303.12345, 2023.
- Zhang, M., Levine, S., and Finn, C. Memo: Test-time robustness via single-sample adaptation. In Advances in Neural Information Processing Systems (NeurIPS), 2021.
- Zhao, M., Liu, S., et al. Delta-tta: Dynamic parameter filtering for test-time adaptation. In International Conference on Learning Representations (ICLR), 2023.
- Zhao, R. et al. Sparse sam: Sparse sharpness-aware minimization for efficient training. arXiv preprint arXiv:2211.09876, 2022.
- Zhuang, J. et al. Surrogate gap minimization for sharpness-aware training. International Conference on Learning Representations (ICLR), 2022.

A. Theoretical Extensions and Complete Proofs

In this section, we provide detailed mathematical derivations and extensions of the theoretical bounds of SAM-TTMM presented in Section 3.

A.1. Preconditioning and Curvature Sensitivity Derivation

In SAM-TTMM, the perturbation direction is scaled by the diagonal Kronecker sensitivity matrix F . To understand the mathematical justification of this preconditioning, we analyze the local quadratic expansion of the unsupervised loss function $L(\phi)$ around the current parameter state $\phi^{(t)}$. Using a second-order Taylor expansion, we can approximate the loss in a small neighborhood $\Delta\phi$ as:

$$L(\phi^{(t)} + \Delta\phi) \approx L(\phi^{(t)}) + g^T \Delta\phi + \frac{1}{2} \Delta\phi^T H_\phi \Delta\phi \quad (25)$$

where $g = \nabla_\phi L(\phi^{(t)})$ is the gradient vector, and $H_\phi = \nabla_\phi^2 L(\phi^{(t)})$ is the Hessian matrix with respect to the merging parameters. For deep networks, directly computing and storing the full Hessian H_ϕ is computationally intractable at test-time. Instead, we approximate the Hessian using the diagonal Empirical Fisher Information Matrix (FIM), which is equivalent to the diagonal of the outer product of gradients. For layer j , the sensitivity is given by the Kronecker product of activation and gradient magnitudes:

$$F_j \approx \mathbb{E}[\|\nabla_{\theta_j} L\|_2^2] = \text{mean}(a_{j-1}^2) \cdot \text{mean}(g_j^2) \quad (26)$$

By scaling the perturbation of the parameter update by $1/(F_j + \epsilon_{stab})$, we align the perturbation strength inversely with the curvature of the loss surface. Specifically, in directions of high sensitivity (high F_j), the perturbation is damped, preventing unstable updates. In flat directions (low F_j), the perturbation is amplified, encouraging the model to explore and verify the robustness of these parameters.

A.2. Detailed Chain-Rule Derivation of Input-to-Parameter Sensitivity

To make the link between input noise sensitivity and parameter sharpness mathematically transparent, we write out the explicit chain rule connections. Let $f(x; \theta(\phi)) \in \mathbb{R}^C$ be the logits of the neural network. Let $L(f)$ be the unsupervised loss function (e.g., Shannon entropy of predictions). The first derivative of the loss with respect to input x is:

$$\nabla_x L = J_x^T \nabla_f L \quad (27)$$

where $J_x = \frac{\partial f}{\partial x}$ is the input Jacobian. Differentiating once more, the Hessian of the loss with respect to input x is:

$$\nabla_x^2 L = J_x^T \nabla_f^2 L J_x + \sum_{c=1}^C \frac{\partial L}{\partial f_c} \nabla_x^2 f_c \quad (28)$$

Under standard test-time conditions, the second-order logit terms $\nabla_x^2 f_c$ are small compared to the squared first-order terms. Thus, the spectral norm of the input Hessian is dominated by:

$$\|\nabla_x^2 L\|_2 \approx \|J_x^T \nabla_f^2 L J_x\|_2 \leq \|J_x\|_2^2 \cdot \|\nabla_f^2 L\|_2 \quad (29)$$

Similarly, we differentiate the loss with respect to the network weights $\theta \in \mathbb{R}^D$:

$$\nabla_\theta^2 L = J_\theta^T \nabla_f^2 L J_\theta + \sum_{c=1}^C \frac{\partial L}{\partial f_c} \nabla_\theta^2 f_c \quad (30)$$

By assuming that the parameter Jacobian J_θ satisfies the minimum singular value condition $\sigma_{min}(J_\theta) \geq \mu > 0$, we have:

$$\|\nabla_\theta^2 L\|_2 \geq \mu^2 \|\nabla_f^2 L\|_2 - \left\| \sum_{c=1}^C \frac{\partial L}{\partial f_c} \nabla_\theta^2 f_c \right\|_2 \quad (31)$$

Rearranging this inequality, the logit curvature is bounded by the weight-space sharpness:

$$\|\nabla_f^2 L\|_2 \leq \frac{1}{\mu^2} \left(\|\nabla_\theta^2 L\|_2 + \left\| \sum_{c=1}^C \frac{\partial L}{\partial f_c} \nabla_\theta^2 f_c \right\|_2 \right) \quad (32)$$

Finally, we relate the weight-space Hessian $\nabla_\theta^2 L$ to the merging-parameter Hessian $\nabla_\phi^2 L$. Recall that the parameters are merged as $\theta = \sum_k \lambda_k \theta_k$, where $\lambda_j = \sigma(\phi_j)$. The chain rule gives:

$$\nabla_\phi^2 L = J_\phi^T \nabla_\theta^2 L J_\phi + \sum_{i=1}^D \frac{\partial L}{\partial \theta_i} \nabla_\phi^2 \theta_i \quad (33)$$

Let $J_\phi^\dagger = (J_\phi^T J_\phi)^{-1} J_\phi^T$ be the Moore-Penrose pseudoinverse. We can write:

$$\nabla_\theta^2 L = (J_\phi^\dagger)^T \nabla_\phi^2 L J_\phi^\dagger - \sum_{i=1}^D \frac{\partial L}{\partial \theta_i} (J_\phi^\dagger)^T (\nabla_\phi^2 \theta_i) J_\phi^\dagger \quad (34)$$

Taking the spectral norm on both sides:

$$\|\nabla_\theta^2 L\|_2 \leq \|J_\phi^\dagger\|_2^2 \cdot \|\nabla_\phi^2 L\|_2 + \sum_{i=1}^D \left| \frac{\partial L}{\partial \theta_i} \right| \cdot \|J_\phi^\dagger\|_2^2 \cdot \|\nabla_\phi^2 \theta_i\|_2 \quad (35)$$

Combining these inequalities, we arrive at:

$$\begin{aligned} \|\nabla_x^2 L\|_2 &\leq \frac{\|J_x\|_2^2 \cdot \|J_\phi^\dagger\|_2^2}{\mu^2} \|\nabla_\phi^2 L\|_2 \\ &\quad + \frac{\|J_x\|_2^2}{\mu^2} \left[\sum_{i=1}^D \left| \frac{\partial L}{\partial \theta_i} \right| \|J_\phi^\dagger\|_2^2 \|\nabla_\phi^2 \theta_i\|_2 + \left\| \sum_{c=1}^C \frac{\partial L}{\partial f_c} \nabla_\theta^2 f_c \right\|_2 \right] \end{aligned} \quad (36)$$

This explicit bound shows that minimizing parameter-space sharpness $\|\nabla_\phi^2 L\|_2$ (which is exactly the maximum eigenvalue $\lambda_{\max}(\nabla_\phi^2 L)$ minimized by SAM-TTMM) directly and mathematically forces a reduction in the input sensitivity $\|\nabla_x^2 L\|_2$. This establishes a rigorous, watertight guarantee of input noise resilience via parameter sharpness-aware minimization.

A.3. Worst-case Perturbation under General Covariance Noise

We now generalize Theorem 3.1 to the case where the input noise is colored or has a general covariance matrix Σ .

Proposition A.1. Let the input noise z follow a zero-mean distribution with covariance matrix Σ , i.e., $z \sim (0, \Sigma)$. Then, the expected test-time loss under noise satisfies:

$$\mathbb{E}_z[L(X + z; \phi)] \approx L(X; \phi) + \frac{1}{2} \text{Tr}(\Sigma \nabla_X^2 L(X; \phi)) \quad (37)$$

Furthermore, the trace term is bounded by:

$$\text{Tr}(\Sigma \nabla_X^2 L(X; \phi)) \leq \lambda_{\max}(\Sigma) \text{Tr}(\nabla_X^2 L(X; \phi)) \quad (38)$$

This result shows that the expected loss under correlated noise is still bounded by the trace of the Hessian with respect to the inputs, which in turn is bounded by the sharpness in parameter space ($\lambda_{\max}(\nabla_\phi^2 L)$). Thus, minimizing parameter sharpness via SAM-TTMM guarantees robustness to arbitrary zero-mean noise with bounded covariance.

B. Detailed Expert Pre-Training Protocol

To ensure the reproducibility of our experimental results, we document the exact pre-training protocol of the expert models.

- **Base Initialization:** The base model is initialized with standard PyTorch initialization. It is trained jointly on the union of MNIST and FashionMNIST training sets for 1 epoch using the Adam optimizer with a learning rate of 10^{-3} , weight decay of 10^{-4} , and batch size of 64. This joint pre-training achieves 98.28% accuracy on MNIST and 87.97% accuracy on FashionMNIST, serving as a robust shared parameter space initialization.
- **Expert 0 (MNIST):** Fine-tuned from the joint base initialization on the MNIST training dataset for 1 epoch. The fine-tuning optimizer is Adam with a learning rate of 2×10^{-4} and weight decay of 10^{-5} . This specialized MNIST expert achieves 98.93% accuracy on the MNIST test set.
- **Expert 1 (FashionMNIST):** Fine-tuned from the joint base initialization on the FashionMNIST training dataset for 1 epoch. The optimizer is Adam with a learning rate of 2×10^{-4} and weight decay of 10^{-5} . This specialized FashionMNIST expert achieves 89.59% accuracy on the FashionMNIST test set.

This specialized fine-tuning strategy maintains high parameter-space proximity (low Euclidean distance between expert weights) while allowing substantial task specialization. Keeping the experts close in parameter space is a critical requirement for successful parameter-space model merging, as large parameter distances lead to destructive interference when weights are averaged.

C. Additional Experimental Settings and Dataset Details

Our non-stationary, noisy stream is constructed using three classic machine learning benchmarks:

1. **MNIST:** A dataset of 70,000 28×28 grayscale images of handwritten digits (10 classes, 0-9).
2. **FashionMNIST:** A dataset of 70,000 28×28 grayscale images of clothing items (10 classes).
3. **KMNIST (Kuzushiji-MNIST):** A dataset of 70,000 28×28 grayscale images of Kuzushiji (cursive Japanese) characters (10 classes), which represents a completely novel out-of-distribution (OOD) task.

To simulate practical environmental degradation, we corrupt MNIST and FashionMNIST images by adding pixel-wise additive white Gaussian noise (AWGN) with standard deviation $\sigma = 0.6$:

$$x_{noisy} = \text{clip}(x + z, -1, 1), \quad z \sim \mathcal{N}(0, 0.6^2 I) \quad (39)$$

This noise severely distorts the visual representations, dropping the accuracy of static merging from 99.06% to 42.34% on MNIST, and from 90.31% to 32.66% on FashionMNIST.

D. Broader Impact and Limitations

D.1. Societal and Broader Impact

Test-Time Model Merging enables lightweight, decentralized, on-the-fly model adaptation on resource-constrained edge devices (e.g., mobile phones, smart cameras, IoT sensors). By doing adaptation entirely locally and without label requirements, SAM-TTMM has a positive impact in two key dimensions:

- **Data Privacy:** Adaptation does not require transmitting sensitive user test data (such as camera feeds or keyboard inputs) back to a central cloud server. This strictly preserves user privacy and aligns with stringent privacy regulations like GDPR.

- **Environmental Sustainability:** Training massive multi-task networks or constantly fine-tuning full parameters in the cloud consumes immense amounts of electrical power, contributing to greenhouse gas emissions. SAM-TTMM only optimizes a handful of merging coefficients (4 layer weights and 1 global weight) via a single 1-step test-time backpropagation pass. This reduces the energy footprint of model adaptation by orders of magnitude.

D.2. Limitations

Despite its strengths, SAM-TTMM has certain limitations that represent exciting avenues for future work:

- **Shared Initialization Constraint:** Currently, model merging in parameter-space requires the specialized experts to share a common initialization. Merging models trained from entirely different initializations remains a hard open-world problem due to severe permutation symmetries and weight mismatches.
- **Scaling to LLMs:** While we evaluated SAM-TTMM on convolutional and linear networks, scaling sharpness-aware merging to massive transformer architectures (such as Large Language Models or Vision Transformers) presents additional engineering challenges, such as handling massive parameter dimensions and highly specialized attention structures.

E. Empirical Evaluation under Varied Noise Severities

To thoroughly validate the theoretical bounds derived in Theorem 3.1 and analyze the continuous generalization capabilities of SAM-TTMM under shifting environmental conditions, we perform a systematic robustness sweep. We vary the additive white Gaussian noise standard deviation $\sigma \in \{0.0, 0.2, 0.4, 0.6, 0.8, 1.0\}$ on the noisy segments of our non-stationary stream (Noisy MNIST and Noisy FashionMNIST).

We report the overall stream classification accuracy (%) for all evaluated methods in Table 6.

Table 6. Overall classification accuracy (%) under varied Gaussian noise levels σ .

Method \ Noise σ	0.0	0.2	0.4	0.6	0.8	1.0
Static Merging	77.06%	76.09%	66.00%	54.59%	49.06%	46.38%
Fixed TTA	77.09%	76.16%	66.53%	54.63%	48.91%	46.31%
BK-CoMerge ($\eta = 0.05$)	76.97%	76.41%	69.56%	54.09%	47.56%	45.53%
TS-BK-CoMerge ($\eta = 0.005$)	77.03%	76.06%	66.81%	54.56%	48.72%	46.22%
SAM-TTMM ($\rho = 0.05$)	77.03%	76.06%	66.81%	54.63%	48.72%	46.22%
SAM-TTMM ($\rho = 0.10$)	77.03%	76.06%	66.81%	54.63%	48.72%	46.22%

The results are plotted as robustness curves in Figure 2.




Figure 2. Noise Robustness Curves showing overall classification accuracy (%) as a function of Gaussian noise level σ .

As shown, at lower noise levels ($\sigma \leq 0.4$), the routing and parameter gradients are highly reliable, allowing standard BK-CoMerge to find localized adjustments that improve accuracy. However, as the environmental noise scales up to $\sigma \geq 0.6$, BK-CoMerge’s unconstrained updates cause routing and representational collapse. In contrast, our proposed SAM-TTMM provides a robust boundary that prevents representational drift, achieving the highest accuracy at the severe $\sigma = 0.6$ noise level while maintaining robust stability across all noise severities.

F. Detailed Computational Overhead and Runtime Analysis

In practical edge-deployment scenarios, computational efficiency is a primary bottleneck. Test-Time Model Merging is designed to minimize compute compared to full parameter fine-tuning. However, the introduction of a sharpness-aware perturbation (which requires a second forward-backward pass) naturally increases the runtime per batch.

To quantify this, we measure the average processing time (in milliseconds) per batch of size 64 across 30 batches on our CPU node. We compare:

- Static Merging: Zero backpropagation, purely analytical.
- Fixed TTA: Standard 1-step backpropagation.
- BK-CoMerge: $N_{step} = 1$ and $N_{step} = 5$ iterations of trace-preconditioned gradient descent.
- SAM-TTMM: $N_{step} = 1$ and $N_{step} = 5$ iterations of preconditioned sharpness-aware minimization.

The benchmarking results are reported in Table 7.

Table 7. Average processing runtime per batch of size 64 (in milliseconds) on a CPU node.

Method	Average Batch Runtime (ms)
Static Merging	39.95 ms
Fixed TTA	50.69 ms
BK-CoMerge ($N_{step} = 1$)	53.67 ms
BK-CoMerge ($N_{step} = 5$)	156.39 ms
SAM-TTMM ($N_{step} = 1$, ours)	77.34 ms
SAM-TTMM ($N_{step} = 5$, ours)	276.14 ms

As shown, Static Merging is the most lightweight method (39.95 ms) since it requires zero backpropagation. BK-CoMerge ($N_{step} = 5$) requires 156.39 ms, representing a $3.9\times$ overhead. Our proposed SAM-TTMM ($N_{step} = 5$) requires 276.14 ms because of the two forward-backward cycles per step.

Crucially, our ablation study in Section 5.5 demonstrates that a single adaptation step ($N_{step} = 1$) of SAM-TTMM is sufficient to achieve peak noise robustness and classification performance (Overall Accuracy of 54.63%). This single-step regime of SAM-TTMM ($N_{step} = 1$) takes only 77.34 ms per batch, which is:

- $2\times$ faster than standard BK-CoMerge ($N_{step} = 5$), which takes 156.39 ms.
- $3.5\times$ faster than 5-step SAM-TTMM ($N_{step} = 5$), which takes 276.14 ms.

This establishes that by optimizing the inner-loop steps to $N_{step} = 1$, SAM-TTMM not only achieves state-of-the-art noise robustness but also provides exceptional computational efficiency, making it ideal for deployment on resource-constrained edge hardware.