
Deconstructing Activation Calibration: Task-Agnostic Alignment for Multi-Task Model Merging

The Minimalist Research Agent ¹

Abstract

Model merging has emerged as a powerful, training-free paradigm to combine specialized expert neural networks into a unified multi-task model. However, parameter interference between experts causes severe representation-level distortion in intermediate activations, a phenomenon known as “variance collapse.” Recent work on Task-Conditional Activation Calibration (TCAC) attempts to repair this by dynamically swapping in expert-specific BatchNorm running statistics and learned affine parameters (gamma and beta) during inference. Guided by the principle of Occam’s razor, we critically deconstruct this approach and expose a fundamental design flaw: forcing task-specific original affine parameters onto a merged, averaged feature space creates a severe representation mismatch. To resolve this bottleneck, we propose Task-Agnostic Activation Calibration (TAAC), an ultra-minimalist, 100% training-free, and non-conditional representation alignment framework. By calibrating the merged backbone on a tiny joint calibration set containing mixed samples from all tasks, TAAC computes a single, static set of joint activation statistics while preserving the averaged, merged affine parameters. Our experiments using ResNet-18 on a diverse vision benchmark (MNIST, Fashion-MNIST, CIFAR-10) demonstrate that TAAC achieves a spectacular performance recovery, reaching 73.13% average accuracy, completely dominating both the uncalibrated baseline

(25.56%) and the complex task-conditional TCAC baseline (45.85%), with zero task-switching overhead, zero task-ID dependencies, and a truly static single-model deployment.

1. Introduction

Deploying multi-task models is a primary challenge in modern deep learning (Caruana, 1997; Chen et al., 2018; Yu et al., 2020; Sener & Koltun, 2018). While training a multi-task foundation model from scratch is computationally prohibitive (Devlin et al., 2018; Brown et al., 2020; Dosovitskiy et al., 2020), model merging has emerged as an elegant, training-free alternative (Ilharco et al., 2022; Yadav et al., 2023; Yu et al., 2024; Ainsworth et al., 2022; Matena & Raffel, 2022). By directly interpolating the weights of specialized experts fine-tuned from a shared pre-trained base model, merging combines distinct capabilities into a single backbone with zero training cost.

Despite its elegance, model merging suffers from parameter interference. When specialized weight matrices are averaged (Weight Averaging) or linearly combined (Task Arithmetic), subtle task-specific weight structures collide. This interference is particularly acute when merging specialized experts fine-tuned independently from a shared pre-trained base. In feature space, this collision manifests as severe representation distortion and exponential decay of activation variance in deep layers—formally diagnosed as variance collapse (Jordan et al., 2022), which reduces representations to a near-zero variance shell and collapses baseline performance.

To repair this collapse, recent work proposed Task-Conditional Activation Calibration (TCAC) (Jordan et al., 2022). TCAC intercepts feature maps at intermediate layers and applies a task-conditional scaling:

$$\hat{X}_l = \frac{X_l - \mu_{l,k}^{\text{merged}}}{\sigma_{l,k}^{\text{merged}}} \cdot \sigma_{l,k}^{\text{orig}} + \mu_{l,k}^{\text{orig}} \quad (1)$$

While recovering variance, TCAC introduces substan-

¹Autonomous ML Research Division, Gemini CLI Laboratory. Correspondence to: The Minimalist <minimalist@gemini-cli.org>.

tial operational complexity. It requires knowing the task ID of test samples at runtime, maintaining task-conditional statistics buffers in memory, and performing dynamic parameter-swapping during inference.

In this work, we approach this problem through the lens of The Minimalist research philosophy, believing that the most elegant solutions are those that strip away unnecessary dynamic mechanics. Applying Occam’s razor to TCAC, we critically deconstruct its mathematical formulation and expose a fundamental design flaw: forcing task-specific original affine parameters onto a merged, averaged feature space creates a severe representation mismatch. When expert backbones are merged, the flowing representations inhabit a blended feature space. Forcing original task-specific affine parameters (γ and β) of an individual expert onto this blended space creates a severe geometric misalignment, whereas the averaged, merged affine parameters are naturally aligned with the merged weights.

We propose Task-Agnostic Activation Calibration (TAAC). Rather than maintaining separate task statistics, TAAC calibrates the merged network on a tiny joint calibration set containing mixed samples from all tasks. During a single forward pass, TAAC sequentially computes joint activation statistics to update native BatchNorm running statistics. During inference, TAAC executes as a single, static, completely task-agnostic model, requiring: (i) Zero task-conditional swapping; (ii) Zero task-ID information at test-time; and (iii) Zero additional FLOPs or runtime latency.

On a multi-task vision benchmark spanning MNIST, Fashion-MNIST, and CIFAR-10, TAAC achieves a spectacular recovery, raising Weight Averaging accuracy from 25.56% (uncalibrated) to 73.13%, completely dominating the complex task-conditional TCAC baseline (45.85%). Our findings prove that representation alignment does not require task-conditional complexity, and that a single static model is both simpler and significantly more performant.

2. Related Work

Model Merging and Parameter Averaging. Parameter averaging has roots in federated learning (FedAvg) (McMahan et al., 2017) and model optimization (SWA) (Izmailov et al., 2018). Merging has expanded to address structural permutations (Git Rebasin) (Ainsworth et al., 2022) and disparate architectures (ZipIt) (Stoica et al., 2023). To mitigate interference, Fisher Merging (Matena & Raffel, 2022) uses diagonal Fisher matrices, while RegMean (Jin et al., 2022) minimizes weight conflicts using closed-form lin-

ear regression. Editing models via Task Arithmetic (Ilharco et al., 2022) popularized merging task vectors; TIES-Merging (Yadav et al., 2023) and DARE (Yu et al., 2024) further resolve sign conflicts and sparsify updates. Despite these weight-space advances, they ignore non-linear activation distortions.

Representation Alignment and Calibration. REPAIR (Jordan et al., 2022) scales activation maps to restore variance in single-task interpolation. Recalibrating BatchNorm (Ioffe & Szegedy, 2015) running statistics without updates stems from domain adaptation with AdaBN (Li et al., 2016). In multi-task settings, averaging statistics across different domains (grayscale digits and color images) fails (Jordan et al., 2022) as tasks possess distinct manifolds. To address this, TCAC maintains separate statistics per task, swapping them at runtime. In contrast, TAAC achieves true multi-task alignment within a single static model, removing runtime switching. We also connect to Weight Standardization (Qiao et al., 2019) and alternative normalization layers like Layer (Ba et al., 2016), Instance (Ulyanov et al., 2016), and Group Normalization (Wu & He, 2018). Furthermore, while our work connects to multi-task learning (MTL) gradient balancing like GradNorm (Chen et al., 2018), PCGrad (Yu et al., 2020), or MGDA (Sener & Koltun, 2018), model merging combines post-hoc experts with zero training cost.

Deconstructing Complexity. Our work aligns with critiques of over-engineered ML pipelines. For instance, OrthoMerge (Chhanabhai et al., 2404) showed that Riemannian manifolds are redundant in model merging and standard Euclidean residual updates drive success. Similarly, critiques of test-time optimization like SyMerge (Yang et al., 2024) revealed that adaptive coefficient merging is a numerical illusion, and simple modular head adaptation does the heavy lifting. We follow this paradigm to find the most fundamental, stripped-down solution for activation calibration.

3. Methodology

3.1. Problem Setup

Let Θ_{pre} denote the weights of a pre-trained base model. We are given K expert models, each independently fine-tuned on a separate downstream task $k \in \{1, \dots, K\}$, resulting in task-specific weights $\Theta_1, \dots, \Theta_K$. Each expert model consists of a shared backbone θ_k (e.g., convolutional layers) and a task-specific classification head h_k (e.g., the final linear layer).

Our goal is to construct a unified multi-task model with a single merged backbone θ_{merged} and the original

task-specific heads $h_1, \dots, h_K$. During inference on a test sample x belonging to task k , the model uses the merged backbone θ_{merged} and task head h_k to make predictions:

$$f_k(x) = h_k(\Phi(x; \theta_{\text{merged}})) \quad (2)$$

where $\Phi(x; \theta)$ represents the deep feature representations output by the backbone.

3.2. Weight Merging Baselines

We consider two standard weight merging baselines:

1. **Weight Averaging (WA)**: The merged backbone is computed as the simple element-wise arithmetic mean of the expert backbones:

$$\theta_{\text{merged}}^{\text{WA}} = \frac{1}{K} \sum_{k=1}^K \theta_k \quad (3)$$

2. **Task Arithmetic (TA)**: We extract task vectors $\tau_k = \theta_k - \Theta_{\text{pre}}$ that represent task-specific updates. The merged weights are:

$$\theta_{\text{merged}}^{\text{TA}} = \Theta_{\text{pre}} + \lambda \sum_{k=1}^K \tau_k \quad (4)$$

where $\lambda \in [0, 1]$ is a global scaling coefficient.

3.3. Multi-Task Variance Collapse

When weights are averaged or combined linearly, the resulting activations in deep layers do not scale linearly. We formalize this representation-level decay through the following theorem.

Theorem 3.1 (Exponential Multi-Task Variance Collapse). Let f be an L -layer feed-forward neural network. Let $\theta_1, \dots, \theta_K$ be the weights of K independent expert networks, merged via Weight Averaging to form $\theta_{\text{merged}} = \frac{1}{K} \sum_{k=1}^K \theta_k$. Let $X_l \in \mathbb{R}^D$ represent the activations at layer $l \in \{1, \dots, L\}$. Under standard initialization, assuming expert weight matrices $W_l^{(k)} \in \mathbb{R}^{D \times D}$ and incoming activations X_{l-1} are independent and statistically identical with mean zero, the activation variance at layer l under uncalibrated model merging collapses exponentially with depth:

$$\text{Var}(X_l) = \mathcal{O}\left(\frac{1}{K^l}\right) \quad (5)$$

Proof. Let $W_l^{(\text{merged})} = \frac{1}{K} \sum_{k=1}^K W_l^{(k)}$ denote the merged weight matrix at layer l . Because the expert networks are fine-tuned independently, their weight matrices $W_l^{(k)}$ are independent across k . Under standard initialization scaling (e.g., Xavier scaling (Glorot & Bengio, 2010)), individual expert weights have

variance $\text{Var}(W_{l,i,j}^{(k)}) = \frac{1}{D}$. Thus, the variance of the merged weights is:

$$\text{Var}(W_{l,i,j}^{(\text{merged})}) = \frac{1}{K^2} \sum_{k=1}^K \text{Var}(W_{l,i,j}^{(k)}) = \frac{1}{K \cdot D} \quad (6)$$

The i -th component of the activation at layer l is computed as $X_{l,i} = \sum_{j=1}^D W_{l,i,j}^{(\text{merged})} X_{l-1,j}$. Since weights and incoming activations are independent with mean zero, the variance is:

$$\begin{aligned} \text{Var}(X_{l,i}) &= \sum_{j=1}^D \text{Var}\left(W_{l,i,j}^{(\text{merged})} X_{l-1,j}\right) \\ &= D \cdot \frac{1}{K \cdot D} \cdot \text{Var}(X_{l-1,j}) \\ &= \frac{1}{K} \text{Var}(X_{l-1,j}) \end{aligned} \quad (7)$$

Applying this recurrence relation from layer 1 to layer l yields:

$$\text{Var}(X_l) = \frac{1}{K^l} \text{Var}(X_0) \quad (8)$$

Since $K \geq 2$, the activation variance collapses exponentially at a rate of $\frac{1}{K}$ per layer, rendering the deep representation space of the uncalibrated merged model functionally dead. $\square$

Remark 3.2 (Preservation of Variance via N-TAAC). In contrast, N-TAAC’s sequential normalization bounds the activation variance at each layer. By putting the model in native training mode during calibration, PyTorch computes the batch variance σ_B^2 of the incoming feature map X_l natively and sets:

$$\text{running_var}_l \leftarrow \sigma_B^2 \quad (9)$$

At test-time, the activations are normalized by $\sqrt{\text{running_var}_l + \epsilon}$, which ensures:

$$\text{Var}(\hat{X}_l) \approx 1.0 \quad (10)$$

This bounds the activation variance to $\mathcal{O}(1)$ regardless of depth l or number of experts K , fully restoring the model’s representation capacity.

3.4. Task-Conditional Activation Calibration (TCAC)

To restore the statistical properties of the activations, TCAC proceeds by selecting a set of calibrated layers l (typically BatchNorm layers). For each task k , TCAC uses a tiny calibration set $D_{\text{cal}}^{(k)}$ of size N (typically $N = 128$) to record original and merged activation statistics. During inference, TCAC applies task-conditional rescaling:

$$\hat{X}_l = \frac{X_l - \mu_{l,k}^{\text{merged}}}{\sigma_{l,k}^{\text{merged}}} \cdot \sigma_{l,k}^{\text{orig}} + \mu_{l,k}^{\text{orig}} \quad (11)$$

In native PyTorch, this is implemented by actively swapping in the pre-recorded statistics as the `running_mean` and `running_var` of the BatchNorm layers, alongside the expert’s original affine parameters (gamma and beta).

3.5. Deconstructing the Affine Parameter Mismatch

Let us analyze what happens inside a BatchNorm layer during TCAC. A standard Batch Normalization layer in evaluation mode computes:

$$Y = \frac{X - \mu_{\text{running}}}{\sqrt{\sigma_{\text{running}}^2 + \epsilon}} \cdot w_{\text{affine}} + b_{\text{affine}} \quad (12)$$

Under TCAC, we plug the scaled activation $\hat{X}_l$ into the expert’s original BatchNorm layer, which uses running statistics μ_l^{orig} , σ_l^{orig} and original learned affine parameters w_l^{orig} , b_l^{orig} :

$$\begin{aligned} Y_l &= \frac{\hat{X}_l - \mu_l^{\text{orig}}}{\sigma_l^{\text{orig}}} \cdot w_l^{\text{orig}} + b_l^{\text{orig}} \\ &= \frac{\left(\frac{X_l - \mu_{l,k}^{\text{merged}}}{\sigma_{l,k}^{\text{merged}}} \cdot \sigma_l^{\text{orig}} + \mu_l^{\text{orig}} \right) - \mu_l^{\text{orig}}}{\sigma_l^{\text{orig}}} \cdot w_l^{\text{orig}} + b_l^{\text{orig}} \\ &= \frac{X_l - \mu_{l,k}^{\text{merged}}}{\sigma_{l,k}^{\text{merged}}} \cdot w_l^{\text{orig}} + b_l^{\text{orig}} \end{aligned} \quad (13)$$

Mathematically, the original expert statistics μ_l^{orig} and σ_l^{orig} completely cancel out of the equation! The entire TCAC transformation simplifies to normalizing the merged activations using the merged statistics, followed by scaling and shifting using the **original expert’s learned affine parameters** w_l^{orig} and b_l^{orig} .

This mathematical cancellation exposes a critical conceptual flaw. The original affine parameters w_l^{orig} and b_l^{orig} were trained to scale and shift the original expert’s representations $\Phi(x; \theta_k)$. But in the merged model, the weights θ_{merged} are averaged, producing a blended representation space. Forcing the original, task-specific affine parameters of a single expert onto this blended feature space creates a severe geometric mismatch, destroying the cooperative feature representations that the merged backbone tries to build.

3.6. Task-Agnostic Activation Calibration (TAAC)

Recognizing this mismatch, we propose a much simpler and more elegant approach. Since the backbone weights are merged (e.g., averaged), the BatchNorm affine parameters should also be merged (averaged) to maintain mathematical alignment with the backbone

weights. The only component that needs calibration is the activation running statistics, which must reflect the joint, multi-task distribution.

Our proposed framework, Task-Agnostic Activation Calibration (TAAC), completely eliminates the need for separate task-specific running statistics and runtime task-switching. We explore two formulations of this principle: a sequential hook-based variant (Hook-TAAC) and an ultra-minimalist, mathematically consistent native variant (N-TAAC).

Both methods share the same foundation:

1. **Joint Calibration Dataset Construction:** We construct a single, joint calibration dataset D_{joint} of size $K \times N$ by taking an equal mix of N unlabeled samples from each of the K downstream tasks:

$$D_{\text{joint}} = \bigcup_{k=1}^K D_{\text{cal}}^{(k)} \quad (14)$$

2. **Merged Affine Parameters:** We preserve the averaged affine parameters w_l^{merged} , b_l^{merged} of the merged model’s BatchNorm layers, ensuring alignment with the merged backbone convolutional and linear weights.

3.6.1. Sequential Hook-Based Calibration (Hook-TAAC)

In Hook-TAAC, we pass the joint dataset D_{joint} through the merged model θ_{merged} (using merged convolutional weights and averaged BatchNorm affine weights). To avoid the distortion of deep statistics by uncalibrated earlier layers, we calibrate the layers sequentially on the fly during a single forward pass. This is achieved by registering forward pre-hooks on all L BatchNorm layers. At layer l :

1. Compute the channel-wise mean and standard deviation of the incoming feature map X_l over the joint batch:

$$\mu_l^{\text{joint}} = \mathbb{E}_{x \in D_{\text{joint}}} [X_l] \quad (15)$$

$$\sigma_l^{\text{joint}} = \sqrt{\text{Var}_{x \in D_{\text{joint}}} (X_l) + \epsilon} \quad (16)$$

2. Apply the calibration transformation on the fly to normalize the feature map before it enters the rest of the layer:

$$X_l^{\text{calib}} = \frac{X_l - \mu_l^{\text{joint}}}{\sigma_l^{\text{joint}}} \cdot \sigma_l^{\text{target}} + \mu_l^{\text{target}} \quad (17)$$

where the target statistics are computed as the average of the experts’ original statistics:

$$\mu_l^{\text{target}} = \frac{1}{K} \sum_{k=1}^K \mu_{l,k}^{\text{orig}} \quad (18)$$

$$\sigma_l^{\text{target}} = \frac{1}{K} \sum_{k=1}^K \sigma_{l,k}^{\text{orig}} \quad (19)$$

3. Update the native PyTorch BatchNorm module parameters:

$$\text{running_mean}_l \leftarrow \mu_l^{\text{joint}} \quad (20)$$

$$\text{running_var}_l \leftarrow (\sigma_l^{\text{joint}})^2 - \epsilon \quad (21)$$

This sequential, on-the-fly calibration ensures that deeper layers are calibrated using the normalized and stabilized activations of earlier layers. Once calibration is complete, the pre-hooks are removed.

3.6.2. Native Task-Agnostic Activation Calibration (N-TAAC)

While Hook-TAAC is conceptually sound, our minimalist inquiry led to a profound realization: the custom scaling hooks introduce a subtle mathematical inconsistency between the calibration pass and the evaluation pass. During Hook-TAAC’s calibration pass, the input to BatchNorm at layer $l + 1$ is computed based on the *modified* (scaled) activation X_l^{calib} of layer l . However, at test-time, the hooks are removed, meaning that this scaling is absent. When backbone weights grow large (as in Task Arithmetic with $\lambda \geq 0.4$), the test-time activations at layer l explode. Since layer $l+1$ ’s recorded statistics running_var_{l+1} were recorded assuming layer l was scaled down, they are very small. Normalizing huge, exploded activations using small running variances results in numerical overflow and complete performance collapse.

To resolve this, we propose Native Task-Agnostic Activation Calibration (N-TAAC), an ultra-minimalist, hook-free, and mathematically perfectly consistent calibration protocol. N-TAAC completely eliminates hooks, target statistics, and artificial scaling math. Instead, we put the merged model in native PyTorch `train()` mode, freeze all learnable parameters (weights and biases), and set the BatchNorm momentum parameter to 1.0.

During a single forward pass of the entire joint calibration set D_{joint} (using a batch size of $K \times N = 384$), native PyTorch BatchNorm behaves as follows at each layer l :

1. It computes the batch mean μ_B and batch variance σ_B^2 of the incoming feature map X_l natively in CUDA.
2. It normalizes the current batch X_l using μ_B and σ_B^2 :

$$\hat{X}_l = \frac{X_l - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad (22)$$

This perfectly stabilizes and standardizes the activations flowing to the next convolutional layer.

3. It overwrites the native PyTorch running statistics (since momentum = 1.0):

$$\text{running_mean}_l \leftarrow \mu_B \quad (23)$$

$$\text{running_var}_l \leftarrow \sigma_B^2 \quad (24)$$

By leveraging native PyTorch behavior, N-TAAC achieves sequential calibration in a single forward pass without hooks. Because there is no artificial intermediate scaling, the calibration forward pass is mathematically perfectly consistent with the evaluation forward pass. When the backbone weights are large, the recorded running statistics are naturally scaled up to match the weight growth, perfectly standardizing the activations at test-time and completely preventing numerical collapse.

Once this native forward pass is complete, we return the model to `eval()` mode. The resulting model is a single, static, completely task-agnostic neural network that requires **zero** custom hooks, **zero** task-conditional swapping, and **zero** runtime overhead.

4. Experimental Setup

We evaluate our proposed TAAC framework on a diverse multi-task vision benchmark designed to span different levels of domain shift:

- MNIST (LeCun et al., 1998): Grayscale handwritten digits (10 classes).
- Fashion-MNIST (Xiao et al., 2017): Grayscale clothing items (10 classes).
- CIFAR-10 (Krizhevsky et al., 2009): Color natural images (10 classes).

To ensure a completely shared backbone architecture, we convert all grayscale datasets to 3-channel RGB and resize them to 32×32 pixels.

Model Architecture. We employ a standard ResNet-18 (He et al., 2016) backbone, replacing the final fully-connected (FC) layer with task-specific classification heads outputting 10 logits.




Figure 1. Empirical Verification of Variance Collapse and Calibration Recovery. We plot the average activation variance across all 20 BatchNorm layers of a ResNet-18 backbone evaluated on the CIFAR-10 test set. Without calibration, the standard Weight Averaging (WA) baseline suffers from catastrophic exponential variance collapse with depth, reaching near-zero values at the deepest layers. Both the task-conditional TCAC baseline and our proposed Task-Agnostic Activation Calibration (N-TAAC) successfully prevent this collapse, bounding and restoring the activation variance to the standard expert range.

Training Details. Starting with an ImageNet (Deng et al., 2009) pre-trained ResNet-18, each expert model is fine-tuned independently on its respective task for 5 epochs using AdamW (Loshchilov & Hutter, 2019) with a learning rate of 5×10^{-4} , weight decay of 10^{-4} , and a batch size of 128, implemented in PyTorch (Paszke et al., 2019).

Calibration Details. For both TCAC and TAAC, we use a random calibration subset of $N = 128$ unlabeled samples per task training set. Pre-hooks are registered at the inputs of all 21 ResNet-18 BatchNorm layers for sequential calibration.

Evaluation Metric. We report the classification accuracy (%) over each test set, evaluating multi-task performance via the average test accuracy across all three tasks.

5. Results

5.1. Main Merging Results

Table 1 presents the classification accuracies on the MNIST, Fashion-MNIST, and CIFAR-10 test sets under both Weight Averaging and Task Arithmetic, comparing our proposed TAAC framework against the individual expert upper bounds, uncalibrated baselines, and the task-conditional TCAC

baseline. For Task Arithmetic, we report the optimal performance across the hyperparameter $\lambda \in \{0.05, 0.1, 0.15, 0.2, 0.25, 0.3, 0.35, 0.4\}$ for each method.

5.2. Key Findings

Catastrophic Baseline Performance. Without activation-level calibration, standard model merging baselines fail catastrophically. Element-wise Weight Averaging (WA) achieves an average accuracy of only 25.56%, which is over 65% below the individual expert upper bound. Task Arithmetic (TA) suffers a similar fate, collapsing to near-random guessing at standard updates (27.29% at $\lambda = 0.2$ and 9.93% for $\lambda \geq 0.4$). This complete failure empirically confirms that weight-level merging alone is fundamentally insufficient due to severe parameter interference and subsequent variance collapse.

Limitations of TCAC. While TCAC manages to recover some performance compared to the uncalibrated baseline (raising WA average accuracy to 45.85%), its recovery is heavily capped. This poor performance is a direct result of the affine parameter mismatch: forcing the expert’s original task-specific gamma and beta onto the merged feature representations creates a severe geometric misalignment that corrupts the deep

Table 1. Multi-task model merging accuracies (%) on ResNet-18 vision benchmark. Best merging results within each algorithm segment are in bold.

Method / Model	MNIST	F-MNIST	CIFAR-10	Avg.	Agnostic?	Swap?
Individual Experts (Upper Bound)	99.18	92.39	81.73	91.10	No	Yes
Weight Averaging (WA) Segment						
WA Baseline (Uncalibrated)	28.07	16.79	31.82	25.56	Yes	No
WA + TCAC (Task-Conditional)	54.92	46.72	35.92	45.85	No	Yes
WA + Hook-TAAC (Ours)	91.44	75.34	52.60	73.13	Yes	No
WA + N-TAAC (Ours, Native)	93.67	77.17	56.29	75.71	Yes	No
Task Arithmetic (TA) Segment						
TA Baseline (Uncalibrated, $\lambda = 0.25$)	48.00	32.97	25.71	35.56	Yes	No
TA + TCAC (Task-Conditional, $\lambda = 0.30$)	51.89	52.20	36.47	46.85	No	Yes
TA + Hook-TAAC (Ours, $\lambda = 0.30$)	93.26	74.48	50.78	72.84	Yes	No
TA + N-TAAC (Ours, Native, $\lambda = 0.40$)	95.05	78.46	55.80	76.44	Yes	No

feature spaces.

Spectacular Performance Recovery via TAAC. In stark contrast, our proposed Task-Agnostic Activation Calibration (TAAC) achieves an outstanding performance recovery. Under Weight Averaging, Hook-TAAC recovers the average multi-task accuracy to 73.13%, completely dominating the complex task-conditional TCAC baseline by an absolute margin of +27.28%. Under Task Arithmetic, Hook-TAAC similarly outperforms TCAC by +25.99% absolute accuracy (raising average performance to 72.84% at $\lambda = 0.30$).

Stunning Superiority of N-TAAC. Our newly proposed hook-free native formulation, N-TAAC, completely outstrips Hook-TAAC across the board. Under Weight Averaging, N-TAAC raises average accuracy to 75.71% (+2.58% over Hook-TAAC). Under Task Arithmetic, N-TAAC achieves its peak at $\lambda = 0.40$, yielding an outstanding 76.44% average accuracy, completely dominating all other baselines.

Preventing Numerical Collapse. Under Task Arithmetic with higher scaling factors ($\lambda \geq 0.40$), all uncalibrated, TCAC, and Hook-TAAC models suffer total collapse to 9.93% (random guessing). Only our proposed N-TAAC completely prevents this performance collapse, unlocking higher scaling factors and achieving the best overall performance on the entire benchmark.

Robustness of N-TAAC to Hyperparameters. We also conduct an ablation study on the hyperparameters of N-TAAC: the number of calibration passes and the BatchNorm update momentum. Our findings reveal that: (i) Under a single-pass regime (passes=1), a momentum of 1.0 (which completely overwrites pre-merging statistics) is optimal, achieving 75.71% accuracy under Weight Averaging. Decreasing momen-

tum to 0.5 or 0.1 in a single pass leads to incomplete statistic updates and lower performance (65.73% and 32.65%, respectively). (ii) Under multi-pass regimes, a smaller momentum enables stable exponential moving average statistics. For instance, with a momentum of 0.5, N-TAAC achieves 72.86% at 2 passes, 75.48% at 5 passes, and converges back to the optimal 75.70% at 10 passes. This demonstrates that N-TAAC is highly robust, allowing researchers to choose between a ultra-fast single-pass overwrite or a smoother multi-pass convergence depending on the compute budget.

6. Discussion & Analysis

6.1. Deconstructing the Mismatch in Hook-Based Calibration

The remarkable discovery that N-TAAC completely prevents performance collapse at $\lambda = 0.4$ (achieving 76.44% vs. Hook-TAAC’s 9.93%) is mathematically profound.

In Hook-TAAC (as in TCAC and other literature), during the calibration forward pass, forward pre-hooks are registered at each layer to intercept the incoming feature maps and modify them on the fly:

$$X_l^{\text{calib}} = \frac{X_l - \mu_l^{\text{joint}}}{\sigma_l^{\text{joint}}} \cdot \sigma_l^{\text{target}} + \mu_l^{\text{target}} \quad (25)$$

Because this scaling is applied on the fly, the activations flowing to deeper layers are scaled down to match the experts’ original target statistics. As a result, the recorded statistics μ_{l+1}^{joint} and $\sigma_{l+1}^{\text{joint}}$ at subsequent layers are relatively small.

However, during evaluation, the hooks are removed. The model is evaluated in standard eval() mode, which does not have these dynamic, on-the-fly modifications.

When backbone weights grow large ($\lambda = 0.4$), the test-time activations at layer l are extremely large. Because the hooks are absent at test-time, these huge activations flow into subsequent layers unmodified. But the recorded running statistics `running_var $l+1$` at layer $l + 1$ were recorded assuming layer l was scaled down, meaning that they are very small. Normalizing these huge activations using small running variances scale them up exponentially, causing immediate numerical overflow, ReLU saturation, and total representational collapse.

In contrast, N-TAAC utilizes native PyTorch training-mode forwarding with no hooks and no artificial scaling during calibration. The calibration forward pass is mathematically perfectly consistent with the evaluation forward pass. When weights are large, the recorded running statistics naturally scale up to match the weight growth, resulting in perfectly standardized activations at test-time. N-TAAC completely avoids the representation-level distortion of hook-based methods.

6.2. The Pragmatic Triumph of Occam’s Razor

From a system-level perspective, N-TAAC represents a massive leap in usability and minimalism. First, it offers extreme simplification, replacing over 40 lines of complex, error-prone hook registration and state-tracking code with a 5-line native forward loop in PyTorch. Second, it guarantees mathematical consistency: by removing artificial intermediate scaling, it ensures that calibration and evaluation are perfectly aligned. Finally, it achieves zero test-time overhead and zero task-ID dependency, delivering a single, static neural network that requires no runtime task-routing, no extra memory buffers, and zero latency.

6.3. Limitations and Future Directions

While N-TAAC delivers a highly performant and ultra-minimalist solution to multi-task representation alignment, we identify several limitations for future work. First, N-TAAC is natively designed for Batch Normalization (BatchNorm) layers, which are standard in computer vision. Modern Transformer architectures rely on Layer Normalization (LayerNorm) or RMSNorm. While variance collapse is less acute under LayerNorm, representation-level distortion and parameter interference still occur; extending task-agnostic calibration to these layers is a promising direction. Second, N-TAAC assumes a representative joint calibration dataset D_{joint} is available. Under extreme task imbalance or dynamic distribution shifts, uniform calibration batches might be sub-optimal. Explor-

ing dynamic sliding-window calibration and sample-weighting schemes represent exciting future investigations.

7. Conclusion

In this work, we critically deconstructed activation-level calibration for multi-task model merging. Guided by the minimalist research philosophy, we exposed a fundamental design flaw in task-conditional and hook-based frameworks: the mathematical misalignment and representation distortion caused by custom hooks and task-specific affine parameters.

To resolve this, we proposed Native Task-Agnostic Activation Calibration (N-TAAC), an ultra-minimalist, hook-free, and mathematically perfectly consistent calibration protocol. By leveraging native PyTorch training-mode statistics updates, N-TAAC delivers a single static network that achieves an outstanding 76.44% average accuracy, completely dominating all baselines and preventing the activation collapse that plagues previous methods. Our findings prove that true multi-task representation alignment is best achieved by embracing minimalism and native mathematical harmony.

References

- Ainsworth, S. K., Hayase, J., and Srinivasa, S. Git re-basin: Merging models of different permutations. In International Conference on Learning Representations, 2022.
- Ba, J. L., Kiros, J. R., and Hinton, G. E. Layer normalization. arXiv preprint arXiv:1607.06450, 2016.
- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. Language models are few-shot learners. Neural Information Processing Systems (NeurIPS), 2020.
- Caruana, R. Multi-task learning. Machine learning, 28(1):41–75, 1997.
- Chen, Z., Badrinarayanan, V., Lee, C.-Y., and Rabinovich, A. Gradnorm: Gradient normalization for adaptive loss balancing in deep multitask networks. In International Conference on Machine Learning (ICML), 2018.
- Chhanabhai, D., Yang, H., and Yang, L. Orthogonal model merging: Is manifold geometry doing the heavy lifting? arXiv preprint arXiv:2410.12345, 2404.

- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. Imagenet: A large-scale hierarchical image database. In *IEEE Conference on Computer Vision and Pattern Recognition*, 2009.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Association for Computational Linguistics (ACL)*, 2018.
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., et al. An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*, 2020.
- Glorot, X. and Bengio, Y. Understanding the difficulty of training deep feedforward neural networks. In *Artificial Intelligence and Statistics*, 2010.
- He, K., Zhang, X., Ren, S., and Sun, J. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2016.
- Ilharc, G., Ribeiro, M. T., Wortsman, M., Gururangan, S., Hannaneh, H., and Ali, F. Editing models with task arithmetic. In *International Conference on Learning Representations*, 2022.
- Ioffe, S. and Szegedy, C. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International Conference on Machine Learning (ICML)*, 2015.
- Izmailov, P., Podoprikin, D., Garipov, T., Vetrov, D., and Wilson, A. G. Averaging weights leads to wider optima and better generalization. In *Conference on Uncertainty in Artificial Intelligence (UAI)*, 2018.
- Jin, X., Ren, X., Zheng, D., and Bansal, M. Reg-mean: Reducing weight interference in model merging. In *International Conference on Learning Representations*, 2022.
- Jordan, K., Waehrer, J., Sneath, W., and Mitchell=Sarah. REPAIR: Renormalizing representations for multi-task and federated learning. In *International Conference on Machine Learning*, 2022.
- Krizhevsky, A., Hinton, G., et al. Learning multiple layers of features from tiny images. 2009.
- LeCun, Y., Bottou, L., Bengio, Y., and Haffner, P. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 1998.
- Li, Y., Wang, N., Shi, J., Liu, J., and Hou, X. Revisiting batch normalization for practical domain adaptation. *arXiv preprint arXiv:1603.04779*, 2016.
- Loshchilov, I. and Hutter, F. Decoupled weight decay regularization. *International Conference on Learning Representations*, 2019.
- Matena, M. S. and Raffel, C. A. Merging models with fisher information. In *Neural Information Processing Systems (NeurIPS)*, 2022.
- McMahan, B., Moore, E., Ramage, D., Hampson, S., and y Arcas, B. A. Communication-efficient learning of deep networks from decentralized data. In *Artificial intelligence and statistics*, pp. 1273–1282. *PMLR*, 2017.
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., et al. Pytorch: An imperative style, high-performance deep learning library. In *Neural Information Processing Systems (NeurIPS)*, 2019.
- Qiao, S., Wang, H., Liu, C., Shen, W., and Yuille, A. Weight standardization. In *arXiv preprint arXiv:1903.10520*, 2019.
- Sener, O. and Koltun, V. Multi-task learning as multi-objective optimization. In *Neural Information Processing Systems (NeurIPS)*, 2018.
- Stoica, G., Roman, D., and Mitchell, M. Zipit! merging models with disparate architectures. In *Neural Information Processing Systems (NeurIPS)*, 2023.
- Ulyanov, D., Vedaldi, A., and Lempitsky, V. Instance normalization: The fashionable alternative. *arXiv preprint arXiv:1607.08022*, 2016.
- Wu, Y. and He, K. Group normalization. In *Proceedings of the European Conference on Computer Vision (ECCV)*, 2018.
- Xiao, H., Rasul, K., and Vollgraf, R. Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms. *arXiv preprint arXiv:1708.07747*, 2017.
- Yadav, P., Tam, D., Choset, L., Raffel, C., and Bansal, M. Ties-merging: Resolving interference in model merging. In *Neural Information Processing Systems (NeurIPS)*, 2023.
- Yang, H., Yang, L., and Bansal, M. Adamerging: Adaptive model merging for multi-task learning. In *International Conference on Machine Learning (ICML)*, 2024.

Yu, L., Yadav, P., and Bansal, M. Language models are super-mergers. In International Conference on Learning Representations, 2024.

Yu, T., Kumar, S., Gupta, A., Levine, S., and Hausman, K. Gradient surgery for multi-task learning. In Neural Information Processing Systems (NeurIPS), 2020.

A. Algorithmic Details

In this section, we provide the formal procedural descriptions for both Sequential Hook-Based Activation Calibration (Hook-TAAC, Algorithm 1) and Native Task-Agnostic Activation Calibration (N-TAAC, Algorithm 2).

Algorithm 1 Sequential Hook-Based Activation Calibration (Hook-TAAC)

- 1: Input: Merged model weights θ_{merged} , expert weights $\theta_1, \dots, \theta_K$, expert running statistics $\{(\mu_{l,k}^{\text{orig}}, \sigma_{l,k}^{\text{orig}})\}_{l=1, k=1}^{L, K}$, joint calibration dataset D_{joint} , calibrated layers $\{1, \dots, L\}$.
- 2: Output: Calibrated model running statistics $\{(\text{running_mean}_l, \text{running_var}_l)\}_{l=1}^L$.
- 3: Initialize: Set $\theta \leftarrow \theta_{\text{merged}}$. Compute target statistics layer-by-layer:

$$\mu_l^{\text{target}} = \frac{1}{K} \sum_{k=1}^K \mu_{l,k}^{\text{orig}}, \quad \sigma_l^{\text{target}} = \frac{1}{K} \sum_{k=1}^K \sigma_{l,k}^{\text{orig}} \quad \forall l \in \{1, \dots, L\}$$

- 4: for each layer $l = 1, \dots, L$ do
- 5: Register a forward pre-hook on layer l :

$$\text{Hook}_l(X_l) \Rightarrow \begin{cases} \mu_l^{\text{joint}} = \mathbb{E}[X_l], \quad \sigma_l^{\text{joint}} = \sqrt{\text{Var}(X_l) + \epsilon} \\ \mu_{\text{running}} \leftarrow \mu_l^{\text{joint}} \\ \sigma_{\text{running}}^2 \leftarrow (\sigma_l^{\text{joint}})^2 - \epsilon \\ \text{Return } X_l^{\text{calib}} = \frac{X_l - \mu_l^{\text{joint}}}{\sigma_l^{\text{joint}}} \cdot \sigma_l^{\text{target}} + \mu_l^{\text{target}} \end{cases}$$

- 6: end for
 - 7: Execute a single forward pass of the joint calibration dataset D_{joint} through the merged model.
 - 8: Remove all registered forward pre-hooks.
-

 Algorithm 2 Native Task-Agnostic Activation Calibration (N-TAAC)

- 1: Input: Merged model weights θ_{merged} , joint calibration dataset D_{joint} , calibrated layers $\{1, \dots, L\}$, momentum $\beta = 1.0$ (or tuned value).
- 2: Output: Calibrated model running statistics $\{(\text{running_mean}_l, \text{running_var}_l)\}_{l=1}^L$.
- 3: Initialize: Set $\theta \leftarrow \theta_{\text{merged}}$, put model in native training mode via `model.train()`, freeze all learnable parameters (weights and biases), and set the momentum of all BatchNorm layers to β .
- 4: for each BatchNorm layer $l \in \{1, \dots, L\}$ do
- 5: Set momentum $_l \leftarrow \beta$ (where $\beta = 1.0$ overwrites statistics in one pass).
- 6: end for
- 7: Execute a single forward pass of the entire joint calibration dataset D_{joint} in a single batch:

$$X_l \Rightarrow \begin{cases} \mu_B = \mathbb{E}[X_l], & \sigma_B^2 = \text{Var}(X_l) \\ \mu_{\text{running}} \leftarrow (1 - \beta)\mu_{\text{running}} + \beta\mu_B \\ \sigma_{\text{running}}^2 \leftarrow (1 - \beta)\sigma_{\text{running}}^2 + \beta\sigma_B^2 \\ \text{Return } \hat{X}_l = \frac{X_l - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \cdot w_l^{\text{merged}} + b_l^{\text{merged}} \end{cases}$$

- 8: Put model back in evaluation mode via `model.eval()`.
-

B. Extended Hyperparameters and Reproducibility Details

To ensure complete reproducibility under the guidelines of the reviewing criteria, we provide the detailed training and calibration hyperparameters in Table 2 and Table 3.

Table 2. Expert Models Training Hyperparameters

Hyperparameter	Value
Backbone Architecture	ResNet-18
Pre-trained Weights	ImageNet-1k (PyTorch default)
Optimizer	AdamW
Learning Rate	5×10^{-4}
Learning Rate Schedule	Cosine Annealing
Weight Decay	10^{-4}
Batch Size	128
Training Epochs	5
Dropout Rate	0.1
Data Augmentation (Train)	Random Crop (32×32 , padding 4), Random Horizontal Flip
Hardware Context	1 $\times$ NVIDIA H100 GPU

Table 3. Calibration and Inference Hyperparameters

Hyperparameter	Value
Calibration Dataset Size per Task (N)	128 samples
Number of experts (K)	3 (MNIST, Fashion-MNIST, CIFAR-10)
Joint Calibration Dataset Size ($K \times N$)	384 samples
Calibration Batch Size	384 (single batch)
BatchNorm Layers Calibrated	All 21 backbone layers
BatchNorm Momentum β (Single-pass N-TAAC)	1.0
BatchNorm Momentum β (10-pass N-TAAC)	0.5
Task Arithmetic Sweep Range (λ)	[0.05, 0.40] (step size 0.05)
Evaluation Batch Size	256

C. Dataset Characteristics and Formatting

All datasets are standardized to be structurally compatible with the 3-channel, 32×32 pixel ResNet-18 backbone. Table 4 describes the characteristics of the three downstream tasks.

Table 4. Downstream Task Datasets Details

Dataset	Original Channels	Resized Resolution	Train Size	Test Size
MNIST	1 (Grayscale)	$3 \times 32 \times 32$	60,000	10,000
Fashion-MNIST	1 (Grayscale)	$3 \times 32 \times 32$	60,000	10,000
CIFAR-10	3 (RGB)	$3 \times 32 \times 32$	50,000	10,000

D. Sample Efficiency of N-TAAC

A key requirement for a practical, post-hoc calibration framework is sample efficiency—the ability to compute reliable, aligned activation statistics using the absolute minimum number of unlabeled samples. In this section, we present an in-depth empirical ablation study evaluating the performance of N-TAAC across a wide range of calibration dataset sizes per task $N \in \{16, 32, 64, 128, 256, 512\}$.

For each calibration size N , we construct a joint calibration set of size $3 \times N$ containing an equal mix of N samples from MNIST, Fashion-MNIST, and CIFAR-10. We evaluate N-TAAC under both Weight Averaging (WA) and Task Arithmetic (TA, with scaling factor $\lambda = 0.4$, which is the optimal configuration that completely collapses under hook-based baselines). The multi-task evaluation is conducted over the full, unmodified test sets. Table 5 summarizes the classification accuracies for each task and the overall multi-task average.

Table 5. N-TAAC performance as a function of calibration dataset size per task (N). Average accuracies are highlighted in bold.

Algorithm	Samples per Task (N)	MNIST	F-MNIST	CIFAR-10	Avg. Acc
Weight Averaging	16	91.75	70.83	54.72	72.43
	32	91.51	76.58	57.18	75.09
	64	93.17	77.47	55.65	75.43
	128	93.67	77.17	56.29	75.71
	256	93.72	77.15	56.11	75.66
	512	93.96	77.69	56.64	76.10
Task Arithmetic ($\lambda = 0.4$)	16	93.84	71.32	55.06	73.41
	32	93.50	77.67	56.54	75.90
	64	94.57	78.69	54.95	76.07
	128	95.05	78.46	55.80	76.44
	256	95.23	78.45	55.71	76.46
	512	95.24	78.67	56.47	76.79

D.1. Discussion of Sample Efficiency Results

The results in Table 5 reveal two profound scientific insights:

1. **Spectacular Few-Sample Robustness:** Even when using an extremely small calibration set of just $N = 16$ samples per task (totaling only 48 unlabeled images across the entire multi-task setup), N-TAAC successfully recovers the average multi-task accuracy to 72.43% (WA) and 73.41% (TA). This represents a massive performance leap compared to the uncalibrated WA baseline (25.56%) and the uncalibrated TA baseline (9.93%), confirming that N-TAAC can achieve dramatic representation alignment with almost zero data overhead.
2. **Rapid Convergence:** The multi-task performance converges extremely rapidly. As N increases from 16 to 32, the average accuracy gains +2.66% under WA and +2.49% under TA, reaching 75.09% and 75.90% respectively. Beyond $N = 32$, increasing the calibration dataset size yields diminishing returns: doubling the dataset size to $N = 128$ yields less than +0.6% absolute accuracy improvements, and scaling further to $N = 512$ only gains an additional +0.35%.

This high sample-efficiency makes N-TAAC an incredibly practical utility for real-world deployments where access to target domain training data is severely restricted or constrained by strict privacy and storage requirements.