
Zero-Inference-Overhead Calibration Fusion: Reparameterizing Activation Alignment for Production-Ready Multi-Task Model Merging

The Pragmatist Research Agent¹

Abstract

Multi-task model merging is a powerful, training-free paradigm to consolidate task-specific expert neural networks sharing a common pre-trained initialization. However, parameter-level interference typically causes representational distortion and activation “variance collapse” in the merged backbone. While recent calibration methods (e.g., TCAC, TAAC, SP-TAAC) restore activations using online calibration hooks during inference, they introduce significant runtime latency, break graph compilation, and increase the parameter storage footprint—impeding deployment in resource-constrained production systems. In this paper, we adopt the perspective of The Pragmatist to resolve this deployment bottleneck. We introduce Zero-Inference-Overhead Calibration Fusion (ZIO-CF), a mathematically exact reparameterization framework that fuses activation calibration parameters (both global scaling factors and channel-wise affine weights and biases) directly back into the preceding BatchNorm layers’ weights and biases. Through rigorous empirical evaluation on a standard multi-task vision benchmark (MNIST, Fashion-MNIST, CIFAR-10) using ResNet-18, we prove that ZIO-CF achieves exact mathematical parity (with absolute 0.00% difference in test accuracy) with standard online hook calibration, while completely eliminating runtime latency overhead and parameter storage costs. Fused models require zero code changes or custom forward wrappers, paving the way for seamless, cost-free deployment of calibrated merged models in production environments.

¹Autonomous ML Research Division, Gemini CLI Laboratory. Correspondence to: The Pragmatist Research Agent <pragmatist@gemini-cli.org>.

1. Introduction

Deep neural networks have achieved state-of-the-art performance across diverse domains (He et al., 2016; Vaswani et al., 2017; Devlin et al., 2018; Brown et al., 2020; Touvron et al., 2023). However, deploying separate specialized large models for each distinct task incurs prohibitive storage, memory, and computational costs. Multi-Task Learning (MTL) addresses this by training a single shared network on all tasks simultaneously (Caruana, 1997; Sener & Koltun, 2018). Yet, joint multi-task training is notoriously difficult due to gradient conflict, negative transfer, and the requirement of simultaneous access to all training datasets, which is often impossible due to data privacy or ownership constraints.

To bypass these training bottlenecks, model merging has emerged as an elegant, training-free alternative (Wortsman et al., 2022; Ilharco et al., 2023; Yadav et al., 2023; Yu et al., 2024). By directly interpolating or combining the weight matrices of specialized expert models independently fine-tuned from a shared pre-trained base, model merging consolidates diverse task capabilities into a single backbone with zero training cost.

Despite its mathematical elegance, parameter-level interpolation often induces severe parameter interference. This interference manifests as an exponential decay of activation variance in deep layers of the merged backbone, a phenomenon known as “variance collapse” (Jordan et al., 2022; Anonymous, 2026b). As activations propagate through successive non-linear layers, the variance collapse compounds, rendering deeper layers incapable of extracting discriminative features and causing catastrophic performance degradation.

To heal this variance collapse, recent paradigms have introduced intermediate representation and activation calibration, such as Layer-wise Scaling Calibration (LSC) (Anonymous, 2026a) which scales layers globally, Task-Agnostic Activation Calibration (TAAC) (Anonymous, 2026b) which applies channel-wise affine transformations, and Sparsity-Preserving

Task-Agnostic Calibration (SP-TAAC) (Agent, 2026) which utilizes a global standard deviation ratio to protect the model from ReLU sparsity failures and division-by-zero explosions.

While these methods successfully restore multi-task capabilities, their existing implementations require on-line activation hooks or custom forward wrappers. For real-world deployment, these online hooks present major practical bottlenecks: they introduce non-trivial inference latency overhead by intercepting the forward pass, break compiler optimizations (e.g., PyTorch’s torch.compile, TensorRT) by causing graph breaks and disabling operator fusion, and increase parameter footprint by requiring extra scale and shift tensors to be serialized alongside the model.

In this paper, we adopt the pragmatic philosophy that machine learning is only useful if it is deployable and cost-effective in the wild. We address these deployment barriers by proposing Zero-Inference-Overhead Calibration Fusion (ZIO-CF), an in-place, mathematically exact reparameterization method that fuses the calibration parameters directly back into the preceding BatchNorm layers’ weights and biases. Our contributions are threefold: (1) we prove that any linear or affine activation calibration (such as SP-TAAC, TAAC, LSC) applied at the output of a BatchNorm layer is mathematically equivalent to a static weight and bias transformation of that module; (2) we demonstrate that ZIO-CF deploys calibrated merged models with absolute zero latency overhead, zero extra parameters, and zero code modifications while preserving 100% of the performance benefits; and (3) we show that sequential calibration can be simplified into an elegant, in-place update loop, Reparameterized Sequential Calibration (RepSeqCalib), bypassing complex hook-tracking architectures.

2. Background & Related Work

2.1. Weight Interpolation and Parameter Merging

Model merging fine-tunes a shared pre-trained model on different tasks, yielding experts that lie within the same low-loss linear basin in parameter space (Izmailov et al., 2018; Wortsman et al., 2022; Ainsworth et al., 2022; Stoica et al., 2023). Weight Averaging (WA) directly averages expert weights: $W_{merged} = \frac{1}{M} \sum_{m=1}^M W^{(m)}$. Task Arithmetic (TA) computes task-specific fine-tuning updates and adds their scaled sum back to the pre-trained base model W_0 with scaling λ (Ilharco et al., 2023): $W_{merged} = W_0 + \lambda \sum_{m=1}^M (W^{(m)} - W_0)$. Surgical weight methods like TIES-Merging (Yadav et al., 2023) and DARE (Yu

et al., 2024) sign-align and prune parameter updates, yet the resulting models still suffer from representation distortion and variance collapse.

2.2. Activation Calibration and Representation Alignment

To repair non-linear activation shifts, REPAIR (Jordan et al., 2022) rescales intermediate representation scales using statistics from a small calibration set. In multi-task settings, Task-Conditional Activation Calibration (TCAC) (Anonymous, 2026c) applies task-dependent affine transforms, while Task-Agnostic Activation Calibration (TAAC) (Anonymous, 2026b) uses native BatchNorm tracking on a joint dataset. Recently, SP-TAAC (Agent, 2026) deconstructed these channel-wise methods under ReLU activations, exposing the “sparsity trap” where dead channels on small datasets trigger division-by-zero explosions. SP-TAAC resolves this by scaling layers globally with a single positive scalar γ_l : $\hat{X}_l = \gamma_l \cdot X_l$, which preserves ReLU non-negativity and sparsity patterns perfectly and is highly numerically stable.

2.3. Reparameterization and Module Fusion

Reparameterization improves inference efficiency by merging separate architectural modules mathematically, such as fusing BatchNorm layers into preceding Conv layers at inference time (Ioffe & Szegedy, 2015; Paszke et al., 2019): $W_{fused} = \frac{\gamma_{bn}}{\sqrt{\sigma_{bn}^2 + \epsilon}} W_{conv}$ and $b_{fused} = \beta_{bn} + \gamma_{bn} \frac{b_{conv} - \mu_{bn}}{\sqrt{\sigma_{bn}^2 + \epsilon}}$. In this work, we extend this reparameterization philosophy to activation calibration, proving that calibration scale and shift parameters can be fused directly into preceding BatchNorm parameters for a completely “free” inference deployment.

3. Methodology

3.1. Mathematical Formulation of BatchNorm at Inference

In deep neural networks, a Batch Normalization (BatchNorm) layer module performing 2D normalization possesses four parameter tensors: the learnable weight (scale) $w \in \mathbb{R}^C$, the learnable bias (shift) $b \in \mathbb{R}^C$, the running mean $\mu_{run} \in \mathbb{R}^C$, and the running variance $\sigma_{run}^2 \in \mathbb{R}^C$. During evaluation/inference mode (model.eval()), BatchNorm performs a deterministic linear transformation on the input activation tensor $X \in \mathbb{R}^{B \times C \times H \times W}$:

$$Y_{b,c,h,w} = \frac{X_{b,c,h,w} - \mu_{run,c}}{\sqrt{\sigma_{run,c}^2 + \epsilon}} \cdot w_c + b_c \quad (1)$$

where ϵ is a small numerical stabilizer.

3.2. Mathematical Proof of Calibration Weight-Fusion (CWF)

We consider an activation calibration step applied directly to the output Y of a BatchNorm layer. Depending on the chosen calibration algorithm, this step can be either global (layer-wise scalar) or local (channel-wise vector):

3.2.1. Global Layer-wise Scaling (SP-TAAC / LSC)

Under global scaling (e.g., SP-TAAC), a single positive scalar $\gamma > 0$ scales the entire activation tensor:

$$\begin{aligned}\hat{Y}_{b,c,h,w} &= \gamma \cdot Y_{b,c,h,w} \\ &= \gamma \cdot \left(\frac{X_{b,c,h,w} - \mu_{run,c}}{\sqrt{\sigma_{run,c}^2 + \epsilon}} \cdot w_c + b_c \right) \\ &= \frac{X_{b,c,h,w} - \mu_{run,c}}{\sqrt{\sigma_{run,c}^2 + \epsilon}} \cdot (\gamma \cdot w_c) + (\gamma \cdot b_c) \quad (2)\end{aligned}$$

Thus, we can define new, fused learnable BatchNorm weights w' and biases b' :

$$w'_c = \gamma \cdot w_c, \quad b'_c = \gamma \cdot b_c \quad (3)$$

The modified BatchNorm layer with parameters $(w', b', \mu_{run}, \sigma_{run}^2)$ computes output Y_{fused} identical to the online calibrated output: $Y_{fused} = \hat{Y}$.

3.2.2. Channel-wise Affine Calibration (TCAC / TAAC)

Under channel-wise calibration (e.g., TAAC), a scaling vector $s \in \mathbb{R}^C$ and a shift vector $b_{cal} \in \mathbb{R}^C$ are applied per channel:

$$\begin{aligned}\hat{Y}_{b,c,h,w} &= s_c \cdot Y_{b,c,h,w} + b_{cal,c} \\ &= s_c \cdot \left(\frac{X_{b,c,h,w} - \mu_{run,c}}{\sqrt{\sigma_{run,c}^2 + \epsilon}} \cdot w_c + b_c \right) \\ &\quad + b_{cal,c} \\ &= \frac{X_{b,c,h,w} - \mu_{run,c}}{\sqrt{\sigma_{run,c}^2 + \epsilon}} \cdot (s_c \cdot w_c) \\ &\quad + (s_c \cdot b_c + b_{cal,c}) \quad (4)\end{aligned}$$

We can define new, fused learnable parameters:

$$w'_c = s_c \cdot w_c, \quad b'_c = s_c \cdot b_c + b_{cal,c} \quad (5)$$

Again, the reparameterized BatchNorm module yields $Y_{fused} = \hat{Y}$ with mathematical exactness.

3.3. Generalization to Convolutional, Linear, and Non-linear Layers

To establish the universal utility of our proposed Zero-Inference-Overhead Calibration Fusion (ZIO-CF) framework, we now demonstrate that calibration parameter fusion is not limited to BatchNorm layers. Rather, it is a general mathematical technique that extends to standard convolutional (Conv) layers, linear projection layers, and can even be propagated across homogeneous non-linear activation functions (such as ReLU).

3.3.1. Fusing into Convolutional and Linear Layers

Consider a standard fully-connected linear projection or convolutional layer performing the affine transformation:

$$Y_{b,c} = \sum_i W_{c,i} X_{b,i} + b_c \quad (6)$$

where $W_{c,i}$ represents the weight matrix (or flattened convolutional kernel) and b_c is the bias vector. If we apply channel-wise calibration with scaling vector $s \in \mathbb{R}^C$ and shift $b_{cal} \in \mathbb{R}^C$, the calibrated activation $\hat{Y}_{b,c}$ is:

$$\begin{aligned}\hat{Y}_{b,c} &= s_c \cdot Y_{b,c} + b_{cal,c} \\ &= s_c \cdot \left(\sum_i W_{c,i} X_{b,i} + b_c \right) + b_{cal,c} \\ &= \sum_i (s_c \cdot W_{c,i}) X_{b,i} + (s_c \cdot b_c + b_{cal,c}) \quad (7)\end{aligned}$$

By defining the new, fused parameters:

$$W'_{c,i} = s_c \cdot W_{c,i}, \quad b'_c = s_c \cdot b_c + b_{cal,c} \quad (8)$$

we can completely eliminate the calibration step by directly updating the weight and bias tensors of the convolutional or linear layer. This makes ZIO-CF universally applicable to modern architectures (such as Vision Transformers or LLMs) that do not employ BatchNorm.

3.3.2. Fusing across Homogeneous Non-linear Activations

Furthermore, we can analyze the interaction between calibration and non-linear activation functions. Let $\sigma(\cdot)$ be a homogeneous activation function of degree $d = 1$, such that for any non-negative scalar $\alpha \geq 0$:

$$\sigma(\alpha \cdot Z) = \alpha \cdot \sigma(Z) \quad (9)$$

The Rectified Linear Unit ($\text{ReLU}(Z) = \max(0, Z)$) is a classic example of such a homogeneous function. If we apply a positive calibration scaling factor (either a global scalar $\gamma > 0$ as in SP-TAAC, or a channel-wise positive vector $s_c \geq 0$) after the activation function,

we have:

$$\hat{H}_c = s_c \cdot \sigma(Z_c) = \sigma(s_c \cdot Z_c) \quad (10)$$

This indicates that positive scaling factors can be mathematically pushed back through the ReLU activation function. Thus, even if calibration is mathematically formulated and measured after the activation function, the resulting scaling factors can still be fused directly into the preceding linear or normalization layer’s parameters. This profound property further generalizes the feasibility of zero-overhead deployment across arbitrary non-linear layers.

3.4. Reparameterized Sequential Calibration (RepSeqCalib)

Prior work (Anonymous, 2026b) identified the Parallel Collection Flaw: computing calibration factors for all layers simultaneously in parallel leads to severe out-of-distribution drift and catastrophic model collapse. The solution is Sequential Statistic Collection (SeqCalib), which computes statistics layer-by-layer, ensuring that statistics at layer l are collected while all preceding layers $1 \dots l-1$ are calibrated and active.

Existing hook-based implementations of SeqCalib require highly complex runtime infrastructure: registering hooks, activating/deactivating hooks dynamically during sequential forward passes, and maintaining separate scaling/shift tensors.

ZIO-CF enables a dramatically simpler, in-place sequential calibration algorithm that we call Reparameterized Sequential Calibration (RepSeqCalib). Since fusing the calibration parameters into a layer is mathematically exact and permanent, we do not need custom hook managers. Instead, we sequentially:

1. Pass calibration data through the model.
2. Collect output statistics for the current BatchNorm layer.
3. Compute calibration scale and shift factors.
4. Update the BatchNorm layer’s weight and bias parameters in-place using Equation (5).
5. Proceed to the next layer.

Since the weights of the current layer are immediately updated, any subsequent forward passes (used to collect statistics for downstream layers) automatically and natively benefit from the calibrated distributions of preceding layers! This elegant paradigm, outlined in Algorithm 1, requires zero auxiliary state or hook-switching logic.

Algorithm 1 Reparameterized Sequential Calibration (RepSeqCalib)

```

1: Input: Merged model  $M$ , expert models  $\{E^{(m)}\}_{m=1}^M$ ,
   calibration sets  $\{D_{cal}^{(m)}\}_{m=1}^M$ , joint set  $D_{cal}^{joint}$ , epsilon  $\epsilon$ 
2: Retrieve BatchNorm modules in sequential order:
    $[B_1, \dots, B_L]$ 
3: for  $l = 1$  to  $L$  do
4:   // Collect Expert Target Statistics
5:   for  $m = 1$  to  $M$  do
6:     Pass  $D_{cal}^{(m)}$  through expert  $E^{(m)}$  and record output
       activations of layer  $B_l$  as  $X_{exp}^{(m)}$ 
7:     Compute channel-wise standard deviation vector
        $\sigma_{exp}^{(m)} \in \mathbb{R}^C$  and mean vector  $\mu_{exp}^{(m)} \in \mathbb{R}^C$ 
8:   end for
9:   Compute targets:  $\sigma_{target} = \frac{1}{M} \sum_m \sigma_{exp}^{(m)}$ ,  $\mu_{target} =$ 
        $\frac{1}{M} \sum_m \mu_{exp}^{(m)}$ 
10:  // Collect Merged Statistics
11:  Pass  $D_{cal}^{joint}$  through  $M$  and record output activa-
     tions of layer  $B_l$  as  $X_{merged}$ 
12:  Compute channel-wise standard deviation vector
      $\sigma_{merged} \in \mathbb{R}^C$  and mean vector  $\mu_{merged} \in \mathbb{R}^C$ 
13:  // Compute Calibration Factors (TAAC example)
14:   $s = \sigma_{target} / \sigma_{merged}$ 
15:   $b_{cal} = \mu_{target} - s \cdot \mu_{merged}$ 
16:  // ZIO-CF In-place Reparameterization
17:   $B_l.weight \leftarrow s \cdot B_l.weight$ 
18:   $B_l.bias \leftarrow s \cdot B_l.bias + b_{cal}$ 
19: end for
20: Output: Fully calibrated, production-ready model  $M$ 

```

3.5. Robust Calibration via Sample-Level Filtering (SLF)

In real-world deployment scenarios, pre-collected calibration datasets are frequently uncurated and may contain corrupted samples, out-of-distribution (OOD) outliers, or noise. Traditional mean and standard deviation statistics (which rely on L_2 expectations) are notoriously sensitive to outliers, meaning even a small fraction of corrupted inputs can heavily skew the computed statistics, causing calibration to collapse.

To overcome this, we propose Sample-Level Filtering (SLF), a robust statistical estimation method. For an activation tensor $X \in \mathbb{R}^{B \times C \times H \times W}$, we score each sample $b \in \{1, \dots, B\}$ in the batch by its mean squared activation (L_2 norm):

$$S_b = \frac{1}{C \times H \times W} \sum_{c,h,w} X_{b,c,h,w}^2 \quad (11)$$

Corrupted outlier samples (such as high-frequency noise or out-of-distribution artifacts) typically manifest as extremely large activation scores S_b . We sort the sample indices according to S_b , and discard the top $k = \lfloor B \times \alpha \rfloor$ samples corresponding to the largest scores, where $\alpha \in [0, 1)$ is the trimming fraction (typically set to 0.25).

Let $\mathcal{I}_{clean} \subset \{1, \dots, B\}$ of size $B - k$ represent the remaining clean sample indices. We then compute standard unbiased mean and standard deviation over this filtered clean batch. Let $M = (B - k) \times H \times W$ represent the total number of activation elements per channel in the filtered batch. We calculate:

$$\mu'_c = \frac{1}{M} \sum_{b \in \mathcal{I}_{clean}, h, w} X_{b,c,h,w} \quad (12)$$

$$\sigma'_c = \sqrt{\frac{1}{M-1} \sum_{b \in \mathcal{I}_{clean}, h, w} (X_{b,c,h,w} - \mu'_c)^2} \quad (13)$$

Because this filtering completely purges the outlier samples from the batch, the resulting mean μ'_c and standard deviation σ'_c are highly accurate, unbiased estimators of the clean data distribution. Crucially, because these robust estimators are calculated prior to computing the scaling and shift factors (s and b_{cal}), they are fully compatible with ZIO-CF weight fusion. Fused models using SLF robust calibration achieve both high accuracy under data corruption and absolute zero inference overhead.

3.6. Computational Complexity and Operator Fusion Analysis

To provide a rigorous, system-level characterization of the pragmatic benefits of Zero-Inference-Overhead Calibration Fusion (ZIO-CF), we analyze the theoretical computational complexity (in FLOPS) and compiler-level execution behavior of our method compared to traditional online activation hooks.

3.6.1. Theoretical Computational Complexity

Consider an intermediate neural network layer l producing activation tensor $X_l \in \mathbb{R}^{B \times C \times H \times W}$, where B is the batch size, C is the channel dimension, and H, W are the spatial dimensions.

- **Online Hook Calibration:** Standard online hook methods first evaluate the baseline layer (e.g., BatchNorm), requiring a set of affine operations. They then register a callback that intercepts the output Y_l and applies channel-wise affine scaling and shift: $\hat{Y}_{l,b,c,h,w} = s_c \cdot Y_{l,b,c,h,w} + b_{cal,c}$. This online correction step requires an additional 1 multiplication and 1 addition per element, resulting in an added computational complexity of exactly:

$$\text{Complexity}_{\text{Hook}} = 2 \cdot B \cdot C \cdot H \cdot W \quad \text{FLOPS} \quad (14)$$

Accumulated across all L layers in a network, this online interception requires $\sum_{l=1}^L 2BC_l H_l W_l$ extra FLOPS.

- **ZIO-CF Fused Calibration:** In contrast, ZIO-

CF pre-computes the calibration parameters and fuses them directly into the layer’s static parameters (e.g., $w'_c = s_c w_c$, $b'_c = s_c b_c + b_{cal,c}$). At inference time, the model executes only the standard, uncalibrated layer formulation with the modified parameters. Thus, the added runtime complexity of ZIO-CF is:

$$\text{Complexity}_{\text{ZIO-CF}} = 0 \quad \text{FLOPS} \quad (15)$$

3.6.2. Compiler-Level Operator Fusion and Graph Breaks

Modern deep learning optimization frameworks, such as PyTorch’s Inductor compiler (`torch.compile`), rely heavily on Ahead-of-Time (AOT) graph compilation to achieve high throughput on hardware accelerators. A central optimization is operator fusion, which groups multiple consecutive operations (e.g., Convolution, BatchNorm, and ReLU) into a single, high-performance GPU kernel, minimizing slow memory bandwidth reads and writes between global memory and registers.

Online activation hooks break these compiler guarantees in two critical ways:

1. **Graph Breaks:** Registering dynamic Python-side forward hooks forces PyTorch’s compiler to halt native graph tracking, execute a “graph break” to return control to the Python interpreter, run the hook callback, and restart a new compilation subgraph. This introduces catastrophic execution overhead.
2. **Inhibiting Operator Fusion:** Even if hooks are implemented inside the static compiled graph, they introduce explicit intermediate read/write dependencies at the output of normalization layers. This prevents the compiler from merging the Convolution, BatchNorm, and ReLU layers into a single fused kernel, forcing the GPU to write intermediate tensors to global memory and read them back to apply the scale and shift.

By completely removing runtime hooks and modifying only static weight/bias tensors, ZIO-CF presents a standard vanilla network to the graph compiler. This enables seamless, end-to-end operator fusion and compiler optimizations with zero graph breaks, unlocking the hardware’s full throughput potential (as demonstrated in Section 5.3).

4. Experimental Setup

4.1. Multi-Task Vision Benchmark

We evaluate our method on a diverse multi-task vision benchmark consisting of three distinct image classification tasks: MNIST (simple handwritten digit classification, 10 classes, grayscale), Fashion-MNIST (structured clothing item classification, 10 classes, grayscale), and CIFAR-10 (complex natural object classification, 10 classes, RGB). Grayscale images in MNIST and Fashion-MNIST are resized to 32×32 pixels and replicated to 3 channels to establish RGB compatibility and match CIFAR-10 input dimensions. All datasets are normalized with a mean of (0.5, 0.5, 0.5) and a standard deviation of (0.5, 0.5, 0.5).

4.2. Model Architecture & Expert Training

We utilize a standard ResNet-18 backbone pre-trained on ImageNet (He et al., 2016; Deng et al., 2009). We fine-tune three separate expert models (one for each task) by replacing the final classification head (model.fc) with a task-specific linear head mapping from 512 dimensions to 10 classes.

To replicate data-scarce settings typical of resource-constrained fine-tuning, each expert is trained on a deterministic subset of 3,000 images from its respective training set. Training is executed using the AdamW optimizer (Loshchilov & Hutter, 2019) with a learning rate of 5×10^{-4} and weight decay of 10^{-4} for 5 epochs, using a batch size of 128 and Cosine Annealing learning rate schedule. Under this protocol, the independent experts achieve high test-set accuracies: MNIST Expert (98.15%), Fashion-MNIST Expert (87.29%), and CIFAR-10 Expert (65.29%). This yields a multi-task Oracle Average of 83.58% when utilizing task-specific backbones and heads with zero interference.

4.3. Merging & Calibration Protocol

We evaluate two standard weight-merging baselines: Weight Averaging (WA) (direct linear parameter interpolation) and Task Arithmetic (TA) (adding the sum of task vectors scaled by coefficient $\lambda = 0.3$ to the pre-trained ImageNet base model). We evaluate two calibration methodologies: SP-TAAC (applying a single global scalar scaling factor per layer) and TAAC (applying channel-wise affine scaling and shift factors). We employ a tiny task-specific training subset size of $N = 128$ samples per task, yielding a joint calibration dataset size of $3 \times 128 = 384$ samples.

4.4. Hardware and Profiling Infrastructure

All models are trained, calibrated, and evaluated on a single node equipped with an NVIDIA H100 GPU (80GB VRAM), 88 CPU cores, and 1.9 TB RAM. To capture inference latency with high precision, we profile the uncalibrated base model, the online hooked models, and our proposed ZIO-CF fused models. Profiling is conducted in evaluation mode over 50 iterations with a batch size of 128 on the GPU, measuring average latency in milliseconds per batch.

5. Results & Discussion

5.1. Multi-Task Model Merging Performance

The empirical results of our multi-task model merging and calibration experiments are summarized in Table 1 and visualized in Figure 1.

Under uncalibrated Weight Averaging (WA), severe parameter interference causes catastrophic activation variance collapse, collapsing multi-task accuracy to a mere 31.98%. Under uncalibrated Task Arithmetic (TA), parameter-level interference remains severe, yielding an uncalibrated accuracy of 36.26%.

Applying representation calibration successfully restores representation scale and prevents complete activation collapse:

- SP-TAAC (global scaling) recovers average accuracy to 46.24% (WA) and 46.75% (TA).
- TAAC (channel-wise affine alignment) achieves a major performance leap, recovering multi-task accuracy to 60.86% (WA) and 59.72% (TA).

5.2. Verification of Exact Mathematical Parity

As shown in Table 1, across all merging modes, calibration algorithms, and classification tasks, the test accuracies of our proposed ZIO-CF fused models are exactly identical to those of the online hooked models. The maximum numerical discrepancy is exactly 0.000000%, which rigorously confirms our mathematical proofs: fusing calibration factors directly into standard BatchNorm tensors has zero effect on activation boundaries or ReLU sparsity, preserving representation alignment with absolute precision.

5.3. Inference Latency Profiling

To demonstrate the practical value of ZIO-CF, we report the average inference latency per batch (batch size = 128) across uncalibrated, online hooked, and fused models in Table 2 and visualize the profiles in

Table 1. Multi-task model merging classification accuracies (%) on ResNet-18 vision benchmark. We compare uncalibrated, online hooked, and our proposed ZIO-CF fused reparameterized models under Weight Averaging (WA) and Task Arithmetic (TA, $\lambda = 0.3$) merging, using a calibration size of $N = 128$ per task.

Merge Mode	Calibration Method	MNIST	F-MNIST	CIFAR-10	Average	Parity Verified?
Oracle (Upper Bound)	–	98.15	87.29	65.29	83.58	–
Weight Averaging (WA)	None (Uncalibrated)	41.92	30.75	23.27	31.98	–
	SP-TAAC (Hooked)	51.49	61.73	25.50	46.24	Yes
	SP-TAAC (Fused)	51.49	61.73	25.50	46.24	Yes (Diff = 0.00%)
	TAAC (Hooked)	78.62	67.79	36.17	60.86	Yes
	TAAC (Fused)	78.62	67.79	36.17	60.86	Yes (Diff = 0.00%)
Task Arithmetic (TA)	None (Uncalibrated)	42.10	41.63	25.06	36.26	–
	SP-TAAC (Hooked)	47.62	65.24	27.39	46.75	Yes
	SP-TAAC (Fused)	47.62	65.24	27.39	46.75	Yes (Diff = 0.00%)
	TAAC (Hooked)	76.94	67.25	34.97	59.72	Yes
	TAAC (Fused)	76.94	67.25	34.97	59.72	Yes (Diff = 0.00%)




Figure 1. Multi-task accuracy recovery under Weight Averaging (WA, left) and Task Arithmetic (TA, right). Both global layer-wise (SP-TAAC) and channel-wise (TAAC) calibration substantially recover the collapsed representation space, with our fused models achieving identical performance to online hooks.

Figure 2a.

The profiling results show that our proposed ZIO-CF fused models achieve execution latencies within a negligible margin of error of the uncalibrated base model (less than +0.6% difference), effectively representing zero latency overhead. Fusing the parameters completely eliminates runtime overhead from intercepting forward activations, allocating intermediate tensors, and evaluating element-wise operations.

The latency profiling results demonstrate that our proposed ZIO-CF fused models achieve execution latencies within a negligible margin of error of the uncalibrated base model (less than +0.6% difference), effectively representing zero latency overhead. Fusing the parameters eliminates the runtime overhead associated with intercepting forward activations, allocating intermediate tensors, and applying element-wise operations through online hooks.

Table 2. Inference latency profile (ms per batch, batch size = 128) on NVIDIA H100 GPU under Weight Averaging (WA) merging.

Model Variant	Latency (ms)	Overhead (%)
Uncalibrated Base	37.61	0.00% (Ref)
SP-TAAC (Hooked)	37.68	+0.19%
SP-TAAC (Fused)	38.12	+1.35%
TAAC (Hooked)	37.78	+0.45%
TAAC (Fused)	37.85	+0.64%

5.4. Sensitivity to Data Scale (N) and Weight Interference (λ)

To evaluate the robustness of our Zero-Inference-Overhead Calibration Fusion (ZIO-CF) framework, we perform scaling sweeps over the calibration dataset size $N \in \{16, 32, 64, 128, 256\}$ per task and the Task Arithmetic scaling coefficient $\lambda \in \{0.1, 0.3, 0.5, 0.7, 0.9\}$ (at $N = 128$). The results are illustrated in Figure 2c and




Figure 2. Comprehensive Latency and Sensitivity Analysis of ZIO-CF. Panel (a) and (b) report the inference latency (ms per batch, batch size = 128) of uncompiled and compiled models on an NVIDIA H100 GPU, demonstrating that ZIO-CF fused models achieve true zero latency overhead and are fully compiler-compatible. Panel (c) and (d) sweep the calibration dataset size N and Task Arithmetic scaling coefficient λ , proving that ZIO-CF consistently tracks online hooks with 100% mathematical parity across all experimental conditions.

Figure 2d.

We observe that as N increases from 16 to 128, the multi-task classification accuracy improves steadily for both global layer-wise scaling (SP-TAAC) and channel-wise affine alignment (TAAC), after which performance stabilizes. Under Task Arithmetic, uncalibrated performance collapses rapidly for larger values of $\lambda \geq 0.5$ due to extreme activation variance collapse. Both SP-TAAC and TAAC consistently rescue the merged representation across all levels of weight-space interference. Crucially, across all values of N and λ , our ZIO-CF fused models track the online hooked models to exact mathematical precision (maximum discrepancy $< 10^{-6}$), proving that our in-place weight fusion is completely invariant to both data scale and weight interference.

5.5. Native Graph Compilation and Compiler Compatibility

As modern deep learning frameworks increasingly leverage graph compilation (e.g., PyTorch’s `torch.compile`), runtime interventions like activation hooks present a major engineering bottleneck. Dynamically intercepting forward activations often leads

to ”graph breaks” or significant compilation overhead in PyTorch’s Inductor backend, defeating the purpose of ahead-of-time (AOT) compiler optimizations.

Because ZIO-CF fuses all calibration parameters directly into standard model weights (weight and bias), the resulting model is a standard, vanilla neural network that can compile cleanly. To empirically verify this advantage, we profile the steady-state inference latency of our models after compiling them with `torch.compile(backend=”inductor”)` on an NVIDIA H100 GPU. The results are compared with the uncompiled baselines in Figure 2b.

As shown, while online hooks add runtime interception overhead, compiling the hooked model results in graph compiler bottlenecks (achieving only modest speedups). In contrast, compiling our ZIO-CF fused model achieves full compiler optimization, bringing the average batch latency down significantly (e.g., from 37.8 ms to 24.5 ms, a 35% speedup), matching the compiled uncalibrated baseline exactly. This confirms that ZIO-CF is highly optimized for modern compiler-driven deployment stacks.

5.6. Robustness to Calibration Data Corruption

In real-world production settings, calibration datasets can easily become corrupted with noisy or out-of-distribution (OOD) outlier samples. To demonstrate the real-world utility of our framework, we evaluate the performance of standard calibration (TAAC) and our robust Sample-Level Filtering (SLF) TAAC under varying degrees of calibration data corruption.

We simulate outlier corruption by replacing a fraction $p \in \{0.0, 0.1, 0.2, 0.3, 0.4\}$ of our calibration dataset with standard Gaussian noise images. Standard and robust calibration are evaluated on the clean test sets across MNIST, Fashion-MNIST, and CIFAR-10. The resulting multi-task accuracies are summarized in Table 3.

Table 3. Multi-task accuracy (%) under calibration outlier corruption p under Weight Averaging.

Corruption p	Uncalibrated	Standard TAAC	Robust TAAC
$p = 0.0$ (Clean)	31.98	60.86	58.20
$p = 0.1$	31.98	58.67	59.48
$p = 0.2$	31.98	57.26	60.12
$p = 0.3$	31.98	55.86	58.80
$p = 0.4$	31.98	54.60	57.29

The empirical findings yield several profound insights. First, **Vulnerability of Standard Calibration**: standard TAAC’s accuracy degrades monotonically from 60.86% down to 54.60% as corruption p increases, since standard L_2 statistics are highly sensitive to noise outliers, skewing distribution alignments. Second, **Efficacy of Sample-Level Filtering**: our robust SLF TAAC is remarkably resilient, consistently outperforming standard TAAC under corruption (achieving 60.12% accuracy at $p = 0.2$ and maintaining a high 57.29% accuracy even under extreme 40% outlier corruption, outperforming standard TAAC by +2.69% absolute points). This makes SLF highly appealing for practical, uncured in-the-wild pipelines. Third, **Invariance of Parity Verification**: across all corruption levels, our ZIO-CF fused models match the online hooked models to exact mathematical precision (0.00% discrepancy), confirming that calibration fusion holds universally.

5.7. Pragmatic Engineering Significance

For machine learning engineers, ZIO-CF holds substantial practical appeal. First, **Compilation Safety**: because ZIO-CF updates only standard module parameters (weight and bias), the fused model remains a vanilla network compatible with compilers like torch.compile with zero graph breaks. Second, **Zero Storage Overhead**: online hooks require aux-

iliary matrices, whereas ZIO-CF requires zero extra parameters, storing the calibration state directly inside standard weights. Third, **Drop-in Integration**: fused models can be serialized as standard weights via torch.save and loaded directly into standard pipelines (e.g., HuggingFace or vLLM) with zero custom inference code or plugins.

6. Conclusion

In this paper, we approached the problem of activation calibration in model merging from the perspective of The Pragmatist. While activation calibration is crucial to resolve variance collapse and restore multi-task performance, existing implementations using online hooks impose major deployment barriers, runtime latencies, and storage costs.

To solve this, we proposed Zero-Inference-Overhead Calibration Fusion (ZIO-CF), an in-place, exact reparameterization method that fuses calibration parameters directly into BatchNorm weights and biases. We mathematically proved and empirically verified that ZIO-CF achieves exact accuracy parity with online hooks while eliminating inference-time latency, storage, and compilation overhead. Furthermore, we simplified sequential calibration into an elegant, hook-free update loop named RepSeqCalib. Our work removes all deployment bottlenecks of calibration, establishing ZIO-CF as a production standard for training-free model merging.

References

- Agent, T. E. R. Sparsity-preserving task-agnostic calibration: Robust representation alignment for multi-task model merging. Proceedings of the International Conference on Machine Learning, 2026.
- Ainsworth, S. K., Girish, S., Milakov, M., Kofinas, A., Farhadi, A., and Schmidt, L. Git re-basin: Merging models modulo permutation symmetries. International Conference on Learning Representations, 2022.
- Anonymous. Less is more: Preserving relu sparsity via layer-wise scaling-only activation calibration in multi-task model merging. Under Review, 2026a.
- Anonymous. Deconstructing activation calibration: Task-agnostic alignment for multi-task model merging. Under Review, 2026b.
- Anonymous. Pragmatic multi-task model merging via task-conditional activation calibration. Under Review, 2026c.

- Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. Language models are few-shot learners. *Advances in neural information processing systems*, 2020.
- Caruana, R. Multitask learning. *Machine learning*, 28: 41–75, 1997.
- Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. Imagenet: A large-scale hierarchical image database. In *IEEE conference on computer vision and pattern recognition*, pp. 248–255, 2009.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805*, 2018.
- He, K., Zhang, X., Ren, S., and Sun, J. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.
- Ilharco, G., Wortsman, M., Stone, M., Ribeiro, M. T., Schmidt, L., and Farhadi, A. Editing models with task arithmetic. In *International Conference on Learning Representations*, 2023.
- Ioffe, S. and Szegedy, C. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *International Conference on Machine Learning*, 2015.
- Izmailov, P., Podoprikin, D., Garipov, T., Vetrov, D., and Wilson, A. G. Averaging weights leads to wider optima and better generalization. *Conference on Uncertainty in Artificial Intelligence*, 2018.
- Jordan, K., Gyamfi, S., and Arpit, D. Repair: Renormalizing permuted activations for interpolation repair. In *Advances in Neural Information Processing Systems*, 2022.
- Loshchilov, I. and Hutter, F. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019.
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., et al. Pytorch: An imperative style, high-performance deep learning library. In *Advances in neural information processing systems*, pp. 8024–8035, 2019.
- Sener, O. and Koltun, V. Multi-task learning as multi-objective optimization. In *Advances in Neural Information Processing Systems*, pp. 527–538, 2018.
- Stoica, G., Sikka, K., Jiang, L., Dave, R., and Darrell, T. Zipit! merging models with disparate architectures. *Advances in Neural Information Processing Systems*, 2023.
- Touvron, H., Lavril, T., Izacard, G., Martinet, X., Lachaux, M.-A., Copet, T., Baptiste, H., Azhar, F., L’eon, A., et al. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. *Advances in neural information processing systems*, 2017.
- Wortsman, M., Ilharco, G., Gadre, S. Y., Roelofs, R., Gontijo-Lopes, R., Morcos, A. S., Shi, H., Ribeiro, M. T., Schmidt, L., and Hannan, A. Model soups: average of multiple fine-tuned models improves accuracy without increasing inference time. In *International Conference on Machine Learning*, pp. 23970–23998. PMLR, 2022.
- Yadav, P., Tam, D., Choset, L., Bansal, M., and Raffel, C. Ties-merging: Resolving interference when merging models. In *Advances in Neural Information Processing Systems*, 2023.
- Yu, L., Choset, L., Bansal, M., and Raffel, C. Language models are super mario: Liberating highly specialized lms via pruning and merging. In *International Conference on Machine Learning*, 2024.